

1_Baseline_Models

February 13, 2026

1 King Country House Sales

[Dataset](#)

1.1 Data Description

Column Name	Description	Data Type
id	Unique numeric number assigned to each house being sold	<code>int64</code>
date	Date on which the house was sold out	<code>str</code>
price	Price of house (Target variable)	<code>float64</code>
bedrooms	Number of bedrooms in a house	<code>int64</code>
bathrooms	Number of bathrooms in a bedroom of a house	<code>float64</code>
sqft_living	Measurement of house in square foot	<code>int64</code>
sqft_lot	Square foot of the lot	<code>int64</code>
floors	Total floors (levels) of house	<code>float64</code>
waterfront	Waterfront view (0=No, 1=Yes)	<code>int64</code>
view	House viewed or not (0=No, 1=Yes)	<code>int64</code>
condition	Overall condition of a house (Scale 1-5)	<code>int64</code>
grade	Overall grade given to the housing unit (Scale 1-11)	<code>int64</code>
sqft_above	Square footage of house apart from basement	<code>int64</code>
sqft_basement	Square footage of the basement of the house	<code>int64</code>
yr_built	Date of building of the house	<code>int64</code>
yr_renovated	Year of renovation of house	<code>int64</code>

Column Name	Description	Data Type
zipcode	Zipcode of the location of the house	int64
lat	Latitude of the location of the house	float64
long	Longitude of the location of the house	float64
sqft_living15	Living room area in 2015 (implies some renovations)	int64
sqft_lot15	LotSize area in 2015 (implies some renovations)	int64

1.2 Imports and config

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Models & Normalization
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression, ridge_regression, Lasso

# Evaluation
import statsmodels.api as sm

# Extra
from utils import *
```

1.3 Data Loading

```
[2]: file_path = "Data/king_ country_ houses_aa.csv"

df = pd.read_csv(file_path)
df.head()
```

```
[2]:      id        date   price  bedrooms  bathrooms  sqft_living \
0  7129300520  20141013T000000  221900.0          3       1.00      1180
1  6414100192  20141209T000000  538000.0          3       2.25      2570
2  5631500400  20150225T000000  180000.0          2       1.00       770
3  2487200875  20141209T000000  604000.0          4       3.00      1960
4  1954400510  20150218T000000  510000.0          3       2.00      1680

   sqft_lot  floors  waterfront  view ...  grade  sqft_above  sqft_basement \

```

```

0      5650    1.0          0      0 ...      7      1180          0
1      7242    2.0          0      0 ...      7      2170        400
2     10000    1.0          0      0 ...      6       770          0
3      5000    1.0          0      0 ...      7      1050        910
4      8080    1.0          0      0 ...      8      1680          0

    yr_built  yr_renovated  zipcode      lat      long  sqft_living15 \
0      1955                  0      98178  47.5112 -122.257           1340
1      1951                 1991     98125  47.7210 -122.319           1690
2      1933                  0      98028  47.7379 -122.233           2720
3      1965                  0      98136  47.5208 -122.393           1360
4      1987                  0      98074  47.6168 -122.045           1800

  sqft_lot15
0      5650
1      7639
2      8062
3      5000
4      7503

```

[5 rows x 21 columns]

[3]: df.dtypes

```

[3]: id            int64
date           str
price          float64
bedrooms       int64
bathrooms      float64
sqft_living    int64
sqft_lot       int64
floors         float64
waterfront     int64
view           int64
condition      int64
grade           int64
sqft_above     int64
sqft_basement  int64
yr_built       int64
yr_renovated   int64
zipcode        int64
lat            float64
long           float64
sqft_living15  int64
sqft_lot15     int64
dtype: object

```

```
[4]: # Unique values for each column
df.nunique()
```

```
[4]: id           21436
date          372
price         4028
bedrooms      13
bathrooms     30
sqft_living   1038
sqft_lot      9782
floors        6
waterfront    2
view          5
condition     5
grade          12
sqft_above    946
sqft_basement 306
yr_built      116
yr_renovated  70
zipcode       70
lat            5034
long           752
sqft_living15 777
sqft_lot15    8689
dtype: int64
```

1.4 Minimal Data Cleaning and Transformation

To perform the fitting to the baseline model, we would like to perform the minimum data cleaning and transformation possible.

We are transforming the date column into 3 different columns containing the day, the month and the year. This transformation is important to keep the information about the date which was previously in a format that cannot be fed to the regression models, while verifying which part of the date has relevance to the prediction.

```
[5]: df.date = pd.to_datetime(df.date)
df["year_sold"] = df.date.dt.year
df["month_sold"] = df.date.dt.month
df["day_sold"] = df.date.dt.day
```

```
[6]: df.describe()
```

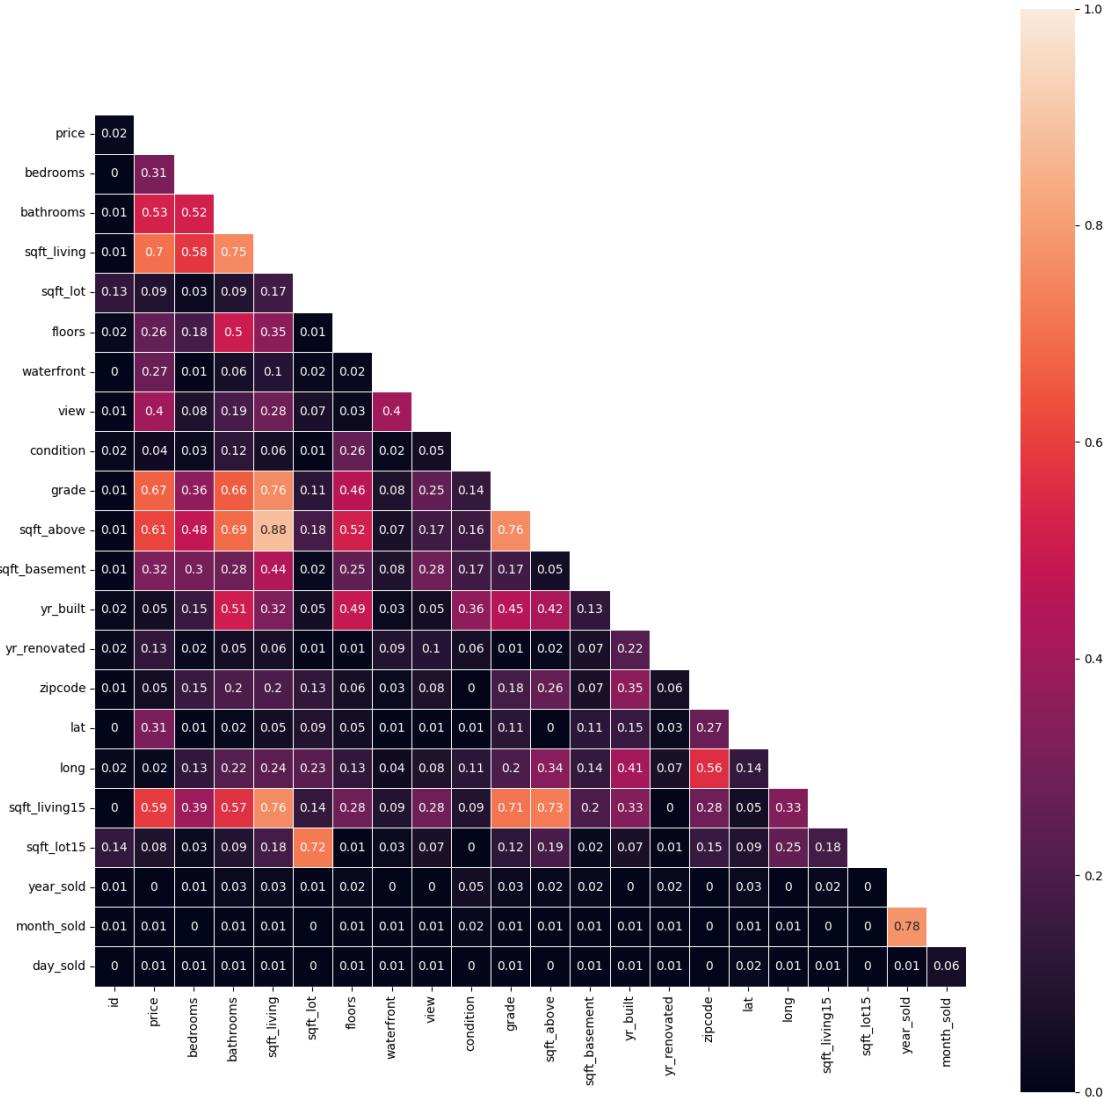
```
[6]:          id                  date        price  bedrooms \
count  2.161300e+04                21613  2.161300e+04  21613.000000
mean   4.580302e+09  2014-10-29 04:38:01.959931  5.400881e+05   3.370842
min    1.000102e+06                2014-05-02 00:00:00  7.500000e+04   0.000000
25%   2.123049e+09                2014-07-22 00:00:00  3.219500e+05   3.000000
```

50%	3.904930e+09		2014-10-16 00:00:00	4.500000e+05	3.000000
75%	7.308900e+09		2015-02-17 00:00:00	6.450000e+05	4.000000
max	9.900000e+09		2015-05-27 00:00:00	7.700000e+06	33.000000
std	2.876566e+09		NaN	3.671272e+05	0.930062
	bathrooms	sqft_living	sqft_lot	floors	waterfront \
count	21613.000000	21613.000000	2.161300e+04	21613.000000	21613.000000
mean	2.114757	2079.899736	1.510697e+04	1.494309	0.007542
min	0.000000	290.000000	5.200000e+02	1.000000	0.000000
25%	1.750000	1427.000000	5.040000e+03	1.000000	0.000000
50%	2.250000	1910.000000	7.618000e+03	1.500000	0.000000
75%	2.500000	2550.000000	1.068800e+04	2.000000	0.000000
max	8.000000	13540.000000	1.651359e+06	3.500000	1.000000
std	0.770163	918.440897	4.142051e+04	0.539989	0.086517
	view	...	yr_built	yr_renovated	zipcode \
count	21613.000000	...	21613.000000	21613.000000	21613.000000
mean	0.234303	...	1971.005136	84.402258	98077.939805
min	0.000000	...	1900.000000	0.000000	98001.000000
25%	0.000000	...	1951.000000	0.000000	98033.000000
50%	0.000000	...	1975.000000	0.000000	98065.000000
75%	0.000000	...	1997.000000	0.000000	98118.000000
max	4.000000	...	2015.000000	2015.000000	98199.000000
std	0.766318	...	29.373411	401.679240	53.505026
	lat	long	sqft_living15	sqft_lot15	year_sold \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	47.560053	-122.213896	1986.552492	12768.455652	2014.322954
min	47.155900	-122.519000	399.000000	651.000000	2014.000000
25%	47.471000	-122.328000	1490.000000	5100.000000	2014.000000
50%	47.571800	-122.230000	1840.000000	7620.000000	2014.000000
75%	47.678000	-122.125000	2360.000000	10083.000000	2015.000000
max	47.777600	-121.315000	6210.000000	871200.000000	2015.000000
std	0.138564	0.140828	685.391304	27304.179631	0.467616
	month_sold	day_sold			
count	21613.000000	21613.000000			
mean	6.574423	15.688197			
min	1.000000	1.000000			
25%	4.000000	8.000000			
50%	6.000000	16.000000			
75%	9.000000	23.000000			
max	12.000000	31.000000			
std	3.115308	8.635063			

[8 rows x 24 columns]

1.4.1 Correlation between the features

```
[7]: generate_heatmap(df.select_dtypes(include="number"))
```



```
[8]: # remove the id since its an identifier and has no relevance in the housepricing.
# remove the date after its transformation
df_clean = df.drop(["id", "date"], axis=1)
```

```
[9]: # Export cleaned and transformed data
# filename = "Data/cleaned_house_sales.csv"
# df_clean.to_csv(filename, index=False)
```

```
[10]: # The price is the target variable
y = df_clean["price"]

# All other variables are the features for the baseline model
X = df_clean.drop("price", axis=1)
```

1.5 Baseline Model

```
[11]: # Split into train and test
seed = 13
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=seed)
```

```
[12]: X_train
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	\
1571	4	1.50	2000	6778	1.0	0	0	
16330	4	2.50	2630	48706	2.0	0	0	
12786	4	2.50	2620	9525	2.5	0	0	
12524	3	2.50	1610	6000	2.0	0	0	
16179	3	1.00	880	18205	1.0	0	0	
...	
153	4	3.25	5180	19850	2.0	0	3	
866	3	2.50	3460	6590	2.0	0	0	
74	3	1.75	1790	50529	1.0	0	0	
14512	2	1.00	820	5040	1.0	0	0	
338	3	1.75	1420	8250	1.0	0	0	
	condition	grade	sqft_above	...	yr_built	yr_renovated	zipcode	\
1571	4	7	1170	...	1962	0	98198	
16330	3	8	2630	...	1986	0	98072	
12786	4	9	2620	...	1974	0	98040	
12524	4	7	1610	...	1993	0	98038	
16179	4	6	880	...	1945	0	98178	
...	
153	3	12	3540	...	2006	0	98006	
866	3	7	3460	...	2001	0	98056	
74	5	7	1090	...	1965	0	98042	
14512	3	7	820	...	1953	0	98199	
338	3	7	1420	...	1954	0	98133	
	lat	long	sqft_living15	sqft_lot15	year_sold	month_sold	\	
1571	47.3708	-122.311	1940	7531	2015	3		
16330	47.7750	-122.125	2680	48706	2014	5		
12786	47.5631	-122.219	2580	9525	2014	8		
12524	47.3490	-122.036	1570	6000	2014	8		
16179	47.5013	-122.244	1110	16115	2014	6		

```

...
153    47.5620 -122.162      ...
866    47.4802 -122.188      ...
74     47.3511 -122.073      ...
14512   47.6498 -122.388      ...
338    47.7535 -122.354      ...

    day_sold
1571      23
16330     21
12786      5
12524     26
16179     24

...
153       1
866      27
74       16
14512    20
338      26

```

[17290 rows x 21 columns]

[13]: y_train

```

[13]: 1571      284950.0
16330     625000.0
12786     838400.0
12524     282000.0
16179     218000.0
...
153       2250000.0
866      467000.0
74       349000.0
14512    508000.0
338      265000.0
Name: price, Length: 17290, dtype: float64

```

1.5.1 Metrics Dataframe

Column	Description
Model	Name of the algorithm used (e.g., Linear Regression, Random Forest).
Split	Indicates the dataset partition (Train or Test) to check for overfitting/underfitting.
R2	R-Squared: The proportion of variance in the target variable explained by the model (0 to 1).
Adjusted_R2	corrected for the number of predictors. Use this to compare models with different numbers of features.

Column	Description
MAE	Mean Absolute Error: The average absolute difference between predicted and actual values. Robust to outliers.
RMSE	Root Mean Squared Error: The square root of average squared errors. Penalizes large errors more heavily than MAE.
MAPE	Mean Absolute Percentage Error: The average error expressed as a percentage. Easy to interpret for business stakeholders.
Comments	Custom notes on model configuration, hyperparameters, or specific observations.

```
[14]: # Create dataframe to save all metrics
metrics_df = create_metrics_df()
metrics_df
```

[14]: Empty DataFrame
Columns: [Model, Split, R2, Adjusted_R2, MAE, RMSE, MAPE, Comments]
Index: []

1.5.2 Linear Model: Not normalized

```
[15]: # Not normalized:
lr_not_norm = LinearRegression()
lr_not_norm.fit(X_train, y_train)

y_train_pred_nn = lr_not_norm.predict(X_train)
y_test_pred_nn = lr_not_norm.predict(X_test)

get_r_squared(y_train=y_train, y_pred_train=y_train_pred_nn, y_test = y_test, y_pred_test=y_test_pred_nn)
```

R2 score:
train | 0.6977440607209231
test | 0.7161824379295636

[15]: (0.6977440607209231, 0.7161824379295636)

```
[16]: metrics_df = add_new_metrics(metrics_df,
                                 lr_not_norm,
                                 X_train,
                                 y_train,
                                 split="train",
                                 comments="Baseline model")
```

```
[17]: metrics_df = add_new_metrics(metrics_df,
                                 lr_not_norm,
```

```
X_test,  
y_test,  
split="test",  
comments="Baseline model")
```

```
[18]: metrics_df
```

```
[18]:
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	0.2561	
1	LinearRegression	test	0.7162	0.7148	125985.6747	0.2596	

	RMSE	Comments
0	202864.5703	Baseline model
1	191531.3335	Baseline model

```
[19]: # OLS: not normalized
```

```
import statsmodels.api as sm  
estimator_nn = sm.OLS(y_train, X_train) # Creates an object OLS estimator  
estimator_nn = estimator_nn.fit()  
estimator_nn.summary()
```

```
[19]:
```

Dep. Variable:	price	R-squared (uncentered):	0.904
Model:	OLS	Adj. R-squared (uncentered):	0.903
Method:	Least Squares	F-statistic:	8088.
Date:	Fri, 13 Feb 2026	Prob (F-statistic):	0.00
Time:	18:23:22	Log-Likelihood:	-2.3584e+05
No. Observations:	17290	AIC:	4.717e+05
Df Residuals:	17270	BIC:	4.719e+05
Df Model:	20		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
bedrooms	-3.637e+04	2114.370	-17.202	0.000	-4.05e+04	-3.22e+04
bathrooms	4.161e+04	3683.098	11.297	0.000	3.44e+04	4.88e+04
sqft_living	111.0403	2.552	43.509	0.000	106.038	116.043
sqft_lot	0.1139	0.054	2.123	0.034	0.009	0.219
floors	5371.0191	4071.934	1.319	0.187	-2610.383	1.34e+04
waterfront	5.778e+05	1.99e+04	29.076	0.000	5.39e+05	6.17e+05
view	5.206e+04	2426.290	21.458	0.000	4.73e+04	5.68e+04
condition	2.551e+04	2646.231	9.641	0.000	2.03e+04	3.07e+04
grade	9.669e+04	2423.262	39.902	0.000	9.19e+04	1.01e+05
sqft_above	73.1607	2.532	28.894	0.000	68.198	78.124
sqft_basement	37.8796	2.990	12.671	0.000	32.020	43.739
yr_built	-2653.8504	81.477	-32.572	0.000	-2813.554	-2494.146
yr_renovated	20.1984	4.139	4.881	0.000	12.086	28.310
zipcode	-644.8994	36.078	-17.875	0.000	-715.616	-574.183
lat	6.072e+05	1.21e+04	50.121	0.000	5.83e+05	6.31e+05
long	-2.257e+05	1.48e+04	-15.286	0.000	-2.55e+05	-1.97e+05
sqft_living15	19.6702	3.876	5.074	0.000	12.072	27.268
sqft_lot15	-0.3846	0.083	-4.644	0.000	-0.547	-0.222
year_sold	5635.3040	1568.428	3.593	0.000	2561.027	8709.581
month_sold	-2811.1927	529.712	-5.307	0.000	-3849.481	-1772.904
day_sold	-424.6159	179.656	-2.363	0.018	-776.760	-72.472
Omnibus:		15223.220	Durbin-Watson:		2.018	
Prob(Omnibus):		0.000	Jarque-Bera (JB):		1731607.437	
Skew:		3.738	Prob(JB):		0.00	
Kurtosis:		51.453	Cond. No.		5.85e+17	

Notes:

- [1] R² is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The smallest eigenvalue is 5.12e-22. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

1.5.3 Linear Model: Normalized

[20]: # normalization

```
normalizer = MinMaxScaler()

X_train_norm = normalizer.fit_transform(X_train)
X_test_norm = normalizer.fit_transform(X_test)
```

[21]: # Linear Regression: normalized

```
lr = LinearRegression()
lr.fit(X_train_norm, y_train)
```

```

y_train_pred = lr.predict(X_train_norm)
y_test_pred = lr.predict(X_test_norm)

get_r_squared(y_train=y_train, y_pred_train=y_train_pred, y_test = y_test, ↴
    ↪y_pred_test=y_test_pred)

```

R2 score:
train | 0.6977440607209233
test | 0.6834108053685031

[21]: (0.6977440607209233, 0.6834108053685031)

```

[22]: # train
metrics_df = add_new_metrics(metrics_df,
                             lr,
                             X_train_norm,
                             y_train,
                             split = "train",
                             comments="Normalized version of the baseline"
)

# test
metrics_df = add_new_metrics(metrics_df,
                             lr,
                             X_test_norm,
                             y_test,
                             split = "test",
                             comments="Normalized version of the baseline"
)

```

[23]: metrics_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	0.2561	
1	LinearRegression	test	0.7162	0.7148	125985.6747	0.2596	
2	LinearRegression	train	0.6977	0.6974	125948.1118	0.2561	
3	LinearRegression	test	0.6834	0.6819	137988.4158	0.2836	
			RMSE		Comments		
0	202864.5703				Baseline model		
1	191531.3335				Baseline model		
2	202864.5703				Normalized version of the baseline		
3	202287.1258				Normalized version of the baseline		

```
[24]: estimator = sm.OLS(y_train, X_train_norm) # Creates an object OLS estimator
estimator = estimator.fit()
estimator.summary()
```

Dep. Variable:	price	R-squared (uncentered):	0.902			
Model:	OLS	Adj. R-squared (uncentered):	0.902			
Method:	Least Squares	F-statistic:	7984.			
Date:	Fri, 13 Feb 2026	Prob (F-statistic):	0.00			
Time:	18:23:22	Log-Likelihood:	-2.3594e+05			
No. Observations:	17290	AIC:	4.719e+05			
Df Residuals:	17270	BIC:	4.721e+05			
Df Model:	20					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
x1	-1.564e+06	6.58e+04	-23.766	0.000	-1.69e+06	-1.44e+06
x2	3.188e+05	2.96e+04	10.764	0.000	2.61e+05	3.77e+05
x3	9.373e+05	1.75e+04	53.441	0.000	9.03e+05	9.72e+05
x4	1.685e+05	8.9e+04	1.893	0.058	-5999.748	3.43e+05
x5	2.793e+04	1.02e+04	2.742	0.006	7964.623	4.79e+04
x6	5.677e+05	2e+04	28.414	0.000	5.29e+05	6.07e+05
x7	2.115e+05	9762.143	21.669	0.000	1.92e+05	2.31e+05
x8	2.435e+04	9176.473	2.654	0.008	6365.955	4.23e+04
x9	8.356e+05	2.28e+04	36.611	0.000	7.91e+05	8.8e+05
x10	1.188e+06	2.41e+04	49.291	0.000	1.14e+06	1.23e+06
x11	3.9e+05	1.75e+04	22.316	0.000	3.56e+05	4.24e+05
x12	-3.377e+05	9139.698	-36.953	0.000	-3.56e+05	-3.2e+05
x13	2.434e+04	8307.838	2.930	0.003	8060.319	4.06e+04
x14	-1.606e+05	6859.606	-23.406	0.000	-1.74e+05	-1.47e+05
x15	3.519e+05	7379.332	47.682	0.000	3.37e+05	3.66e+05
x16	-3.519e+05	1.71e+04	-20.547	0.000	-3.85e+05	-3.18e+05
x17	1.099e+05	2.27e+04	4.844	0.000	6.54e+04	1.54e+05
x18	-3.397e+05	7.25e+04	-4.684	0.000	-4.82e+05	-1.98e+05
x19	2726.7556	4961.512	0.550	0.583	-6998.311	1.25e+04
x20	-4.898e+04	8118.907	-6.033	0.000	-6.49e+04	-3.31e+04
x21	-2.83e+04	5308.269	-5.331	0.000	-3.87e+04	-1.79e+04
Omnibus:	14877.120	Durbin-Watson:	2.019			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1551903.144			
Skew:	3.624	Prob(JB):	0.00			
Kurtosis:	48.844	Cond. No.	2.78e+16			

Notes:

- [1] R^2 is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The smallest eigenvalue is 5.56e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[25]: i = 1

for item in X_train.columns:
    print(i, item)
    i+=1
```

```
1 bedrooms
2 bathrooms
3 sqft_living
4 sqft_lot
5 floors
6 waterfront
7 view
8 condition
9 grade
10 sqft_above
11 sqft_basement
12 yr_built
13 yr_renovated
14 zipcode
15 lat
16 long
17 sqft_living15
18 sqft_lot15
19 year_sold
20 month_sold
21 day_sold
```

1.6 Random Forest

```
[26]: # most common hyperparameters or the default ones
from sklearn.ensemble import RandomForestRegressor

rf_regressor = RandomForestRegressor(random_state=seed)#default values +
random_state = 13
rf_regressor.fit(X_train, y_train)
```

```
[26]: RandomForestRegressor(random_state=13)
```

```
[27]: metrics_df = add_new_metrics(metrics_df,
                                  rf_regressor,
                                  X_train,
                                  y_train,
                                  split = "train",
                                  comments="Baseline, no normalization, random_state=
13, default values.")
```

```
[28]: metrics_df = add_new_metrics(metrics_df,
                                    rf_regressor,
                                    X_test,
                                    y_test,
                                    split = "test",
                                    comments="Baseline, no normalization, random_state=13, default values.")
```

```
[29]: metrics_df
```

```
[29]:
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	0.2561	
1	LinearRegression	test	0.7162	0.7148	125985.6747	0.2596	
2	LinearRegression	train	0.6977	0.6974	125948.1118	0.2561	
3	LinearRegression	test	0.6834	0.6819	137988.4158	0.2836	
4	RandomForestRegressor	train	0.9820	0.9819	25948.4444	0.0486	
5	RandomForestRegressor	test	0.8959	0.8954	67269.8770	0.1271	

	RMSE	Comments
0	202864.5703	Baseline model
1	191531.3335	Baseline model
2	202864.5703	Normalized version of the baseline
3	202287.1258	Normalized version of the baseline
4	49547.1197	Baseline, no normalization, random_state 13, d...
5	115994.2051	Baseline, no normalization, random_state 13, d...

1.7 XGBoost

```
[30]: import xgboost as xgb
```

```
xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[30]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

```
[31]: metrics_df = add_new_metrics(metrics_df,
                                    xgb_clf,
```

```

        X_train,
        y_train,
        split = "train",
        comments="Baseline, no normalization, default\u
        ↵values.")

```

```
[32]: metrics_df = add_new_metrics(metrics_df,
                                    xgb_clf,
                                    X_test,
                                    y_test,
                                    split = "test",
                                    comments="Baseline, no normalization, default\u
                                    ↵values.")
```

```
[33]: metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	0.2561	
1	LinearRegression	test	0.7162	0.7148	125985.6747	0.2596	
2	LinearRegression	train	0.6977	0.6974	125948.1118	0.2561	
3	LinearRegression	test	0.6834	0.6819	137988.4158	0.2836	
4	RandomForestRegressor	train	0.9820	0.9819	25948.4444	0.0486	
5	RandomForestRegressor	test	0.8959	0.8954	67269.8770	0.1271	
6	XGBRegressor	train	0.9780	0.9780	39126.7558	0.0872	
7	XGBRegressor	test	0.9015	0.9010	65712.7448	0.1246	

	RMSE	Comments
0	202864.5703	Baseline model
1	191531.3335	Baseline model
2	202864.5703	Normalized version of the baseline
3	202287.1258	Normalized version of the baseline
4	49547.1197	Baseline, no normalization, random_state 13, d...
5	115994.2051	Baseline, no normalization, random_state 13, d...
6	54676.7727	Baseline, no normalization, default values.
7	112860.4995	Baseline, no normalization, default values.

1.7.1 Insights

1.8 Conclusions & Export

```
[34]: # metrics_df.to_csv("Metrics/baseline_metrics.csv", index=False)
```

```
[ ]:
```