

# 2\_Feature\_Importance

February 13, 2026

## 1 Feature Importance

### 1.1 Imports

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Models & Normalization
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso

# Evaluation
import statsmodels.api as sm

# Extra
from utils import *
```

### 1.2 Open the Clean data

```
[2]: file_path = "Data/cleaned_house_sales.csv"

df_clean = pd.read_csv(file_path)
df_clean.head()
```

```
[2]:    price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront \
0  221900.0        3       1.00      1180      5650     1.0          0
1  538000.0        3       2.25      2570      7242     2.0          0
2  180000.0        2       1.00       770     10000     1.0          0
3  604000.0        4       3.00      1960      5000     1.0          0
4  510000.0        3       2.00      1680      8080     1.0          0

  view  condition  grade  ...  yr_built  yr_renovated  zipcode      lat \
0     0           3     7  ...    1955                  0    98178  47.5112
1     0           3     7  ...    1951                1991    98125  47.7210
```

```

2      0          3      6 ...      1933          0    98028  47.7379
3      0          5      7 ...      1965          0    98136  47.5208
4      0          3      8 ...      1987          0    98074  47.6168

   long  sqft_living15  sqft_lot15  year_sold  month_sold  day_sold
0 -122.257           1340        5650    2014          10         13
1 -122.319           1690        7639    2014          12          9
2 -122.233           2720        8062    2015           2         25
3 -122.393           1360        5000    2014          12          9
4 -122.045           1800        7503    2015           2         18

[5 rows x 22 columns]

```

```
[3]: # Same seed for all random states
seed = 13

# The price is the target variable
y = df_clean["price"]

# All other variables are the features for the baseline model
X = df_clean.drop("price", axis=1)

# Split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=seed)
```

### 1.3 Lasso (L1 regularization)

Should reduce less important features to 0. Applying it to find features we can drop.

```
[4]: # standardization
from sklearn.preprocessing import StandardScaler

std_scaler = StandardScaler()

X_train_standard = std_scaler.fit_transform(X_train)
X_test_standard = std_scaler.fit_transform(X_test)
```

```
[5]: print(X_train_standard.min(), X_train_standard.max())
```

-3.956684808851382 39.35388067380475

The max value being so extreme even after standardization points to a distribution with many outliers.

```
[6]: features = X_train.columns
features
```

```
[6]: Index(['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
       'waterfront', 'view', 'condition', 'grade', 'sqft_above',
       'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode', 'lat', 'long',
       'sqft_living15', 'sqft_lot15', 'year_sold', 'month_sold', 'day_sold'],
      dtype='str')
```

```
[7]: from IPython.display import clear_output

lasso_regressor = Lasso(random_state=seed)

lasso_regressor.fit(X_train_standard, y_train)

# Remove warning
clear_output()
```

```
[8]: # Regression coefficients


coefs_lasso = pd.Series(np.abs(lasso_regressor.coef_), features).
    ↪sort_values(ascending=False)

coefs_lasso
```

```
[8]: sqft_living      260890.945338
grade           113621.446101
lat             84316.523244
sqft_above      81991.026779
yr_built        76475.889246
sqft_basement   59472.818333
waterfront      49274.007710
view            39200.831428
bedrooms        33789.239856
long            32117.139742
bathrooms       32082.314821
zipcode         31555.515044
condition       17557.032206
year_sold       16266.798818
sqft_living15   14460.312554
sqft_lot15      10391.559985
yr_renovated    8500.402589
sqft_lot        4863.151159
day_sold        2879.776176
floors          2293.208022
month_sold      1979.637736
dtype: float64
```

```
[9]: sns.set_theme()
coefs_lasso.plot(kind='bar', title='Lasso Coefficients')
```

```
[9]: <Axes: title={'center': 'Lasso Coefficients'}>
```



The lasso regressor did not identify any irrelevant parameters that we could currently drop.

## 1.4 Ridge (L2 regularization)

Reduces coefficients magnitudes for correlated features. Should give us an idea of multicollinearity.


```
[10]: ridge_regressor = Ridge(random_state=seed)
ridge_regressor.fit(X_train_standard, y_train)

# Regression coefficients

coefs_ridge = pd.Series(np.abs(ridge_regressor.coef_), features).
    sort_values(ascending=False)
```

```
[11]: sns.set_theme()  
coefs_ridge.plot(kind='bar', title='Ridge Coefficients')
```

```
[11]: <Axes: title={'center': 'Ridge Coefficients'}>
```



#### 1.4.1 Lasso vs. Ridge

```
[12]: # combined  
  
feature_importance = pd.DataFrame({"lasso": coefs_lasso, "ridge": coefs_ridge})
```

```
[13]: feature_importance = feature_importance.reset_index()  
feature_importance
```

```
[13]:      index          lasso          ridge  
0    bathrooms  32082.314821  32087.628497
```

1	bedrooms	33789.239856	33783.663620
2	condition	17557.032206	17559.541392
3	day_sold	2879.776176	2880.548680
4	floors	2293.208022	2295.681048
5	grade	113621.446101	113606.366537
6	lat	84316.523244	84314.151501
7	long	32117.139742	32118.058335
8	month_sold	1979.637736	1982.151021
9	sqft_above	81991.026779	78520.819072
10	sqft_basement	59472.818333	25543.707561
11	sqft_living	260890.945338	83124.018496
12	sqft_living15	14460.312554	14478.214453
13	sqft_lot	4863.151159	4866.064328
14	sqft_lot15	10391.559985	10392.616045
15	view	39200.831428	39200.967653
16	waterfront	49274.007710	49272.716356
17	year_sold	16266.798818	16268.574981
18	yr_built	76475.889246	76466.727470
19	yr_renovated	8500.402589	8503.739225
20	zipcode	31555.515044	31551.106899

```
[14]: import pandas as pd
import plotly.express as px

# Fix formating of the dataframe for plotting
df_melted = feature_importance.melt(id_vars='index',
                                      value_vars=['lasso', 'ridge'],
                                      var_name='Model',
                                      value_name='Coefficient')

# Create a interactive bar plot
fig = px.bar(df_melted.sort_values("Coefficient", ascending=False),
              x='index',
              y='Coefficient',
              color='Model',
              barmode='group',           # side-by-side, not stacked
              color_discrete_map={
                  'lasso': '#4c72b0',
                  'ridge': '#dd8452'},
              title='Feature Importance: Lasso vs Ridge')

# 3. Rotate x-axis labels
fig.update_layout(
    xaxis_tickangle=-45,
    xaxis_title='Features',
    yaxis_title='Absolute Coefficient Value'
)
```

```
fig.show()
```

## 1.5 XGBoost

```
[15]: import xgboost as xgb
```

```
xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[15]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

```
[16]: features = X_train.columns
coefs_xgb = pd.Series(np.abs(xgb_clf.feature_importances_), features).
    ↪sort_values(ascending=False)
```


```
coefs_xgb
```

```
[16]: grade          0.368874
waterfront       0.170580
sqft_living      0.167270
lat              0.075877
long             0.041144
view             0.037144
sqft_living15   0.020389
yr_built         0.018146
zipcode          0.016965
bathrooms        0.012077
sqft_above       0.011518
year_sold        0.011033
condition        0.009306
yr_renovated     0.008090
sqft_lot         0.007944
sqft_basement    0.005541
sqft_lot15       0.005373
month_sold       0.004566
day_sold         0.003449
```

```
floors          0.003055  
bedrooms        0.001659  
dtype: float32
```

```
[17]: coefs_xgb.plot(kind='bar', title='XGBoost Coefficients')
```

```
[17]: <Axes: title={'center': 'XGBoost Coefficients'}>
```



## 1.6 Random Forest

```
[18]: # most common hyperparameters or the default ones  
from sklearn.ensemble import RandomForestRegressor  
  
rf_regressor = RandomForestRegressor(random_state=seed)  
#default values +  
#random_state = 13  
rf_regressor.fit(X_train, y_train)
```

```
[18]: RandomForestRegressor(random_state=13)
```


```
[19]: features = X_train.columns
coefs_rf = pd.Series(np.abs(rf_regressor.feature_importances_), features).
    sort_values(ascending=False)

coefs_rf
```

```
[19]: grade          0.313353
sqft_living      0.274236
lat              0.159146
long             0.069911
sqft_living15    0.030356
waterfront       0.025279
yr_built         0.023713
sqft_above        0.018558
zipcode          0.012953
sqft_lot          0.012211
view              0.011875
sqft_lot15        0.011449
bathrooms         0.007467
day_sold          0.006681
month_sold        0.006077
sqft_basement     0.005584
bedrooms          0.002960
condition         0.002917
yr_renovated      0.002241
floors            0.001909
year_sold         0.001126
dtype: float64
```

```
[20]: coefs_rf.plot(kind='bar', title='Random Forest Coefficients')
```

```
[20]: <Axes: title={'center': 'Random Forest Coefficients'}>
```



### 1.6.1 Random Forest vs XGBoost

```
[21]: feature_importance = pd.DataFrame({"rf": coefs_rf, "xgb": coefs_xgb})
feature_importance = feature_importance.reset_index()
feature_importance
```

```
[21]:      index      rf      xgb
0    bathrooms  0.007467  0.012077
1    bedrooms  0.002960  0.001659
2  condition  0.002917  0.009306
3   day_sold  0.006681  0.003449
4     floors  0.001909  0.003055
5     grade  0.313353  0.368874
6       lat  0.159146  0.075877
7     long  0.069911  0.041144
8  month_sold  0.006077  0.004566
```

```

9      sqft_above  0.018558  0.011518
10     sqft_basement  0.005584  0.005541
11      sqft_living  0.274236  0.167270
12    sqft_living15  0.030356  0.020389
13      sqft_lot   0.012211  0.007944
14    sqft_lot15   0.011449  0.005373
15      view    0.011875  0.037144
16    waterfront   0.025279  0.170580
17    year_sold   0.001126  0.011033
18    yr_built    0.023713  0.018146
19  yr_renovated  0.002241  0.008090
20      zipcode   0.012953  0.016965

```

```
[22]: import pandas as pd
import plotly.express as px

# Fix formating of the dataframe for plotting
df_melted = feature_importance.melt(id_vars='index',
                                      value_vars=['rf', 'xgb'],
                                      var_name='Model',
                                      value_name='Coefficient')

# Create a interactive bar plot
fig = px.bar(df_melted.sort_values("Coefficient", ascending=False),
              x='index',
              y='Coefficient',
              color='Model',
              barmode='group',           # side-by-side, not stacked
              color_discrete_map={
                  'rf': '#4c72b0',
                  'xgb': '#dd8452'},
              title='Feature Importance: Random Forest vs XGBoost')

# 3. Rotate x-axis labels
fig.update_layout(
    xaxis_tickangle=-45,
    xaxis_title='Features',
    yaxis_title='Absolute Coefficient Value'
)

fig.show()
```

### 1.6.2 Insights

The two ensemble algorithms, Random Forest and XGBoost, generally agree about their top features, even though they have varying degrees of importance.

[ ]: