

2_Feature_Importance

February 14, 2026

1 Feature Importance

1.1 Imports

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Models & Normalization
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb

# Evaluation
import statsmodels.api as sm

# Extra
from utils import *
```

```
[2]: # Exporting plotly plots to pdf
import plotly.io as pio

# Configuration:

# Set to False for interactive zooming/hovering (Exploration)
# Set to True for static images (PDF Export)
EXPORT_MODE = True

if EXPORT_MODE:
    # Forces all charts to be static images (requires 'pip install -U kaleido')
    pio.renderers.default = "png" # Use "png" if svg gives you trouble
    print(" EXPORT_MODE is ON. Charts will be static images.")
else:
    # Default interactive plotly
```

```

pio.renderers.default = "notebook_connected"
print(" Interactive mode. Charts will include zoom/hover.")

```

EXPORT_MODE is ON. Charts will be static images.

1.2 Open the Clean data

```
[3]: file_path = "Data/v0_cleaned_house_sales.csv"

df_clean = pd.read_csv(file_path)
df_clean.head()
```

```
[3]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront \
0  221900.0         3       1.00      1180      5650     1.0          0
1  538000.0         3       2.25      2570      7242     2.0          0
2  180000.0         2       1.00       770     10000     1.0          0
3  604000.0         4       3.00      1960      5000     1.0          0
4  510000.0         3       2.00      1680      8080     1.0          0

      view  condition  grade  ...  yr_built  yr_renovated  zipcode        lat \
0      0          3     7  ...    1955                  0   98178  47.5112
1      0          3     7  ...    1951                1991   98125  47.7210
2      0          3     6  ...    1933                  0   98028  47.7379
3      0          5     7  ...    1965                  0   98136  47.5208
4      0          3     8  ...    1987                  0   98074  47.6168

      long  sqft_living15  sqft_lot15  year_sold  month_sold  day_sold
0 -122.257        1340        5650    2014        10        13
1 -122.319        1690        7639    2014        12         9
2 -122.233        2720        8062    2015         2        25
3 -122.393        1360        5000    2014        12         9
4 -122.045        1800        7503    2015         2        18
```

[5 rows x 22 columns]

```
[4]: # Same seed for all random states
seed = 13

# The price is the target variable
y = df_clean["price"]

# All other variables are the features for the baseline model
X = df_clean.drop("price", axis=1)

# Split into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=seed)
```

1.3 Lasso (L1 regularization)

Should reduce less important features to 0. Applying it to find features we can drop.

```
[5]: # standardization  
  
std_scaler = StandardScaler()  
  
X_train_standard = std_scaler.fit_transform(X_train)  
X_test_standard = std_scaler.fit_transform(X_test)
```

```
[6]: print(X_train_standard.min(), X_train_standard.max())
```

-3.956684808851382 39.35388067380475

The max value being so extreme even after standardization points to a distribution with many outliers.

```
[7]: features = X_train.columns  
features
```

```
[7]: Index(['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',  
         'waterfront', 'view', 'condition', 'grade', 'sqft_above',  
         'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode', 'lat', 'long',  
         'sqft_living15', 'sqft_lot15', 'year_sold', 'month_sold', 'day_sold'],  
         dtype='str')
```

```
[8]: from IPython.display import clear_output  
  
lasso_regressor = Lasso(random_state=seed)  
  
lasso_regressor.fit(X_train_standard, y_train)  
  
# Remove warning  
clear_output()
```


```
[9]: # Regression coefficients  
  
coefs_lasso = pd.Series(np.abs(lasso_regressor.coef_), features).  
    ↪sort_values(ascending=False)  
  
coefs_lasso
```

```
[9]: sqft_living      260890.945338  
grade            113621.446101  
lat              84316.523244  
sqft_above       81991.026779  
yr_built        76475.889246  
sqft_basement   59472.818333
```

```
waterfront      49274.007710
view            39200.831428
bedrooms        33789.239856
long             32117.139742
bathrooms       32082.314821
zipcode          31555.515044
condition        17557.032206
year_sold        16266.798818
sqft_living15   14460.312554
sqft_lot15      10391.559985
yr_renovated     8500.402589
sqft_lot         4863.151159
day_sold         2879.776176
floors           2293.208022
month_sold       1979.637736
dtype: float64
```

```
[10]: sns.set_theme()
coefs_lasso.plot(kind='bar', title='Lasso Coefficients')
```

```
[10]: <Axes: title={'center': 'Lasso Coefficients'}>
```



The lasso regressor did not identify any irrelevant parameters that we could currently drop.

1.4 Ridge (L2 regularization)

Reduces coefficients magnitudes for correlated features. Should give us an idea of multicollinearity.


```
[11]: ridge_regressor = Ridge(random_state=seed)
ridge_regressor.fit(X_train_standard, y_train)

# Regression coefficients

coefs_ridge = pd.Series(np.abs(ridge_regressor.coef_), features).
    ↪sort_values(ascending=False)
```

```
[12]: sns.set_theme()
coefs_ridge.plot(kind='bar', title='Ridge Coefficients')
```

```
[12]: <Axes: title={'center': 'Ridge Coefficients'}>
```



1.4.1 Lasso vs. Ridge

```
[13]: # combined
```

```
feature_importance = pd.DataFrame({'lasso': coefs_lasso, "ridge": coefs_ridge})
```

```
[14]: feature_importance = feature_importance.reset_index()
feature_importance
```

```
[14]:      index      lasso      ridge
0    bathrooms  32082.314821  32087.628497
1    bedrooms   33789.239856  33783.663620
2  condition   17557.032206  17559.541392
3   day_sold   2879.776176  2880.548680
```

```

4      floors    2293.208022    2295.681048
5      grade    113621.446101   113606.366537
6      lat     84316.523244    84314.151501
7      long    32117.139742    32118.058335
8      month_sold 1979.637736    1982.151021
9      sqft_above 81991.026779   78520.819072
10     sqft_basement 59472.818333   25543.707561
11     sqft_living 260890.945338   83124.018496
12     sqft_living15 14460.312554   14478.214453
13     sqft_lot    4863.151159    4866.064328
14     sqft_lot15 10391.559985    10392.616045
15     view     39200.831428    39200.967653
16     waterfront 49274.007710    49272.716356
17     year_sold 16266.798818    16268.574981
18     yr_built  76475.889246    76466.727470
19     yr_renovated 8500.402589   8503.739225
20     zipcode   31555.515044    31551.106899

```


```
[15]: # Fix formating of the dataframe for plotting
df_melted = feature_importance.melt(id_vars='index',
                                      value_vars=['lasso', 'ridge'],
                                      var_name='Model',
                                      value_name='Coefficient')

# Create a interactive bar plot
fig = px.bar(df_melted.sort_values("Coefficient", ascending=False),
              x='index',
              y='Coefficient',
              color='Model',
              barmode='group',           # side-by-side, not stacked
              color_discrete_map={
                  'lasso': '#4c72b0',
                  'ridge': '#dd8452'},
              title='Feature Importance: Lasso vs Ridge')

# 3. Rotate x-axis labels
fig.update_layout(
    xaxis_tickangle=-45,
    xaxis_title='Features',
    yaxis_title='Absolute Coefficient Value'
)

fig.show()
```

Feature Importance: Lasso vs Ridge



1.5 XGBoost

```
[16]: xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[16]: XGBRegressor(base_score=None, booster=None, callbacks=None,
       colsample_bylevel=None, colsample_bynode=None,
       colsample_bytree=None, device=None, early_stopping_rounds=None,
       enable_categorical=False, eval_metric=None, feature_types=None,
       feature_weights=None, gamma=None, grow_policy=None,
       importance_type=None, interaction_constraints=None,
       learning_rate=None, max_bin=None, max_cat_threshold=None,
       max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
       max_leaves=None, min_child_weight=None, missing=nan,
       monotone_constraints=None, multi_strategy=None, n_estimators=None,
       n_jobs=None, num_parallel_tree=None, ...)
```


```
[17]: features = X_train.columns
coefs_xgb = pd.Series(np.abs(xgb_clf.feature_importances_), features).
    sort_values(ascending=False)

coefs_xgb
```

```
[17]: grade          0.368874  
waterfront      0.170580  
sqft_living     0.167270  
lat             0.075877  
long            0.041144  
view            0.037144  
sqft_living15   0.020389  
yr_built        0.018146  
zipcode         0.016965  
bathrooms       0.012077  
sqft_above       0.011518  
year_sold        0.011033  
condition        0.009306  
yr_renovated    0.008090  
sqft_lot         0.007944  
sqft_basement    0.005541  
sqft_lot15       0.005373  
month_sold       0.004566  
day_sold         0.003449  
floors           0.003055  
bedrooms         0.001659  
dtype: float32
```

```
[18]: coefs_xgb.plot(kind='bar', title='XGBoost Coefficients')
```

```
[18]: <Axes: title={'center': 'XGBoost Coefficients'}>
```



1.6 Random Forest

```
[19]: # most common hyperparameters or the default ones

rf_regressor = RandomForestRegressor(random_state=seed)#default values +
random_state = 13
rf_regressor.fit(X_train, y_train)
```

```
[19]: RandomForestRegressor(random_state=13)
```


```
[20]: features = X_train.columns
coefs_rf = pd.Series(np.abs(rf_regressor.feature_importances_), features).
sort_values(ascending=False)

coefs_rf
```

```
[20]: grade          0.313353
      sqft_living     0.274236
      lat             0.159146
      long            0.069911
      sqft_living15   0.030356
      waterfront       0.025279
      yr_built        0.023713
      sqft_above       0.018558
      zipcode          0.012953
      sqft_lot         0.012211
      view             0.011875
      sqft_lot15       0.011449
      bathrooms         0.007467
      day_sold          0.006681
      month_sold        0.006077
      sqft_basement     0.005584
      bedrooms          0.002960
      condition         0.002917
      yr_renovated      0.002241
      floors             0.001909
      year_sold          0.001126
      dtype: float64
```

```
[21]: coefs_rf.plot(kind='bar', title='Random Forest Coefficients')
```

```
[21]: <Axes: title={'center': 'Random Forest Coefficients'}>
```



1.6.1 Random Forest vs XGBoost

```
[22]: feature_importance = pd.DataFrame({"rf": coefs_rf, "xgb": coefs_xgb})
feature_importance = feature_importance.reset_index()
feature_importance
```

```
[22]:      index      rf      xgb
0    bathrooms  0.007467  0.012077
1    bedrooms  0.002960  0.001659
2  condition  0.002917  0.009306
3   day_sold  0.006681  0.003449
4     floors  0.001909  0.003055
5     grade  0.313353  0.368874
6       lat  0.159146  0.075877
7     long  0.069911  0.041144
8  month_sold  0.006077  0.004566
```

```

9      sqft_above  0.018558  0.011518
10     sqft_basement  0.005584  0.005541
11      sqft_living  0.274236  0.167270
12    sqft_living15  0.030356  0.020389
13      sqft_lot   0.012211  0.007944
14    sqft_lot15   0.011449  0.005373
15      view     0.011875  0.037144
16    waterfront   0.025279  0.170580
17    year_sold   0.001126  0.011033
18    yr_built    0.023713  0.018146
19  yr_renovated  0.002241  0.008090
20      zipcode   0.012953  0.016965

```


```
[23]: # Fix formating of the dataframe for plotting
df_melted = feature_importance.melt(id_vars='index',
                                      value_vars=['rf', 'xgb'],
                                      var_name='Model',
                                      value_name='Coefficient')

# Create a interactive bar plot
fig = px.bar(df_melted.sort_values("Coefficient", ascending=False),
              x='index',
              y='Coefficient',
              color='Model',
              barmode='group',           # side-by-side, not stacked
              color_discrete_map={
                  'rf': '#4c72b0',
                  'xgb': '#dd8452'},
              title='Feature Importance: Random Forest vs XGBoost')

# 3. Rotate x-axis labels
fig.update_layout(
    xaxis_tickangle=-45,
    xaxis_title='Features',
    yaxis_title='Absolute Coefficient Value'
)

fig.show()
```

Feature Importance: Random Forest vs XGBoost



1.6.2 Insights

The two ensemble algorithms, Random Forest and XGBoost, generally agree about their top features, even though they have varying degrees of importance.

[]: