

# 3\_Outliers\_Analysis

February 14, 2026

## 1 Analyse the Influence of Outliers

### 1.1 Imports

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Models & Normalization
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb

# Hypothesis testing
from sklearn.model_selection import cross_val_score
from scipy import stats
import numpy as np

from utils import *

[2]: # Exporting plotly plots to pdf
import plotly.io as pio

# Configuration:

# Set to False for interactive zooming/hovering (Exploration)
# Set to True for static images (PDF Export)
EXPORT_MODE = True

if EXPORT_MODE:
    # Forces all charts to be static images (requires 'pip install -U kaleido')
    pio.renderers.default = "png" # Use "png" if svg gives you trouble
    print(" EXPORT_MODE is ON. Charts will be static images.")
else:
    # Default interactive plotly
    pio.renderers.default = "notebook_connected"
```

```
print(" Interactive mode. Charts will include zoom/hover.")
```

EXPORT\_MODE is ON. Charts will be static images.

## 1.2 Open the clean data

```
[3]: file_path = "Data/v0_cleaned_house_sales.csv"

df_clean = pd.read_csv(file_path)
df_clean.head()
```

```
[3]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront \
0  221900.0         3       1.00      1180      5650     1.0          0
1  538000.0         3       2.25      2570      7242     2.0          0
2  180000.0         2       1.00       770     10000     1.0          0
3  604000.0         4       3.00      1960      5000     1.0          0
4  510000.0         3       2.00      1680      8080     1.0          0

      view  condition  grade  ...  yr_built  yr_renovated  zipcode        lat \
0      0           3     7  ...    1955                  0   98178  47.5112
1      0           3     7  ...    1951                 1991  98125  47.7210
2      0           3     6  ...    1933                  0   98028  47.7379
3      0           5     7  ...    1965                  0   98136  47.5208
4      0           3     8  ...    1987                  0   98074  47.6168

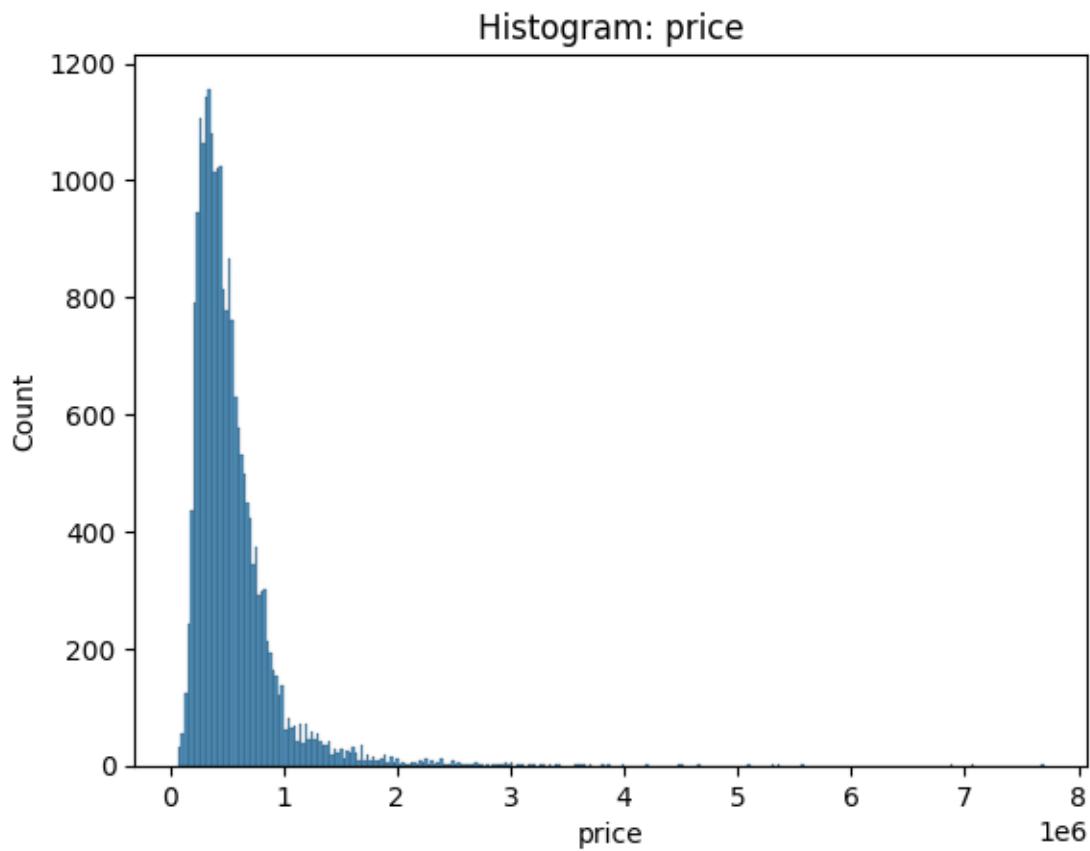
      long  sqft_living15  sqft_lot15  year_sold  month_sold  day_sold
0 -122.257         1340       5650    2014        10        13
1 -122.319         1690       7639    2014        12         9
2 -122.233         2720       8062    2015         2        25
3 -122.393         1360       5000    2014        12         9
4 -122.045         1800       7503    2015         2        18

[5 rows x 22 columns]
```

## 1.3 Analysing Prices

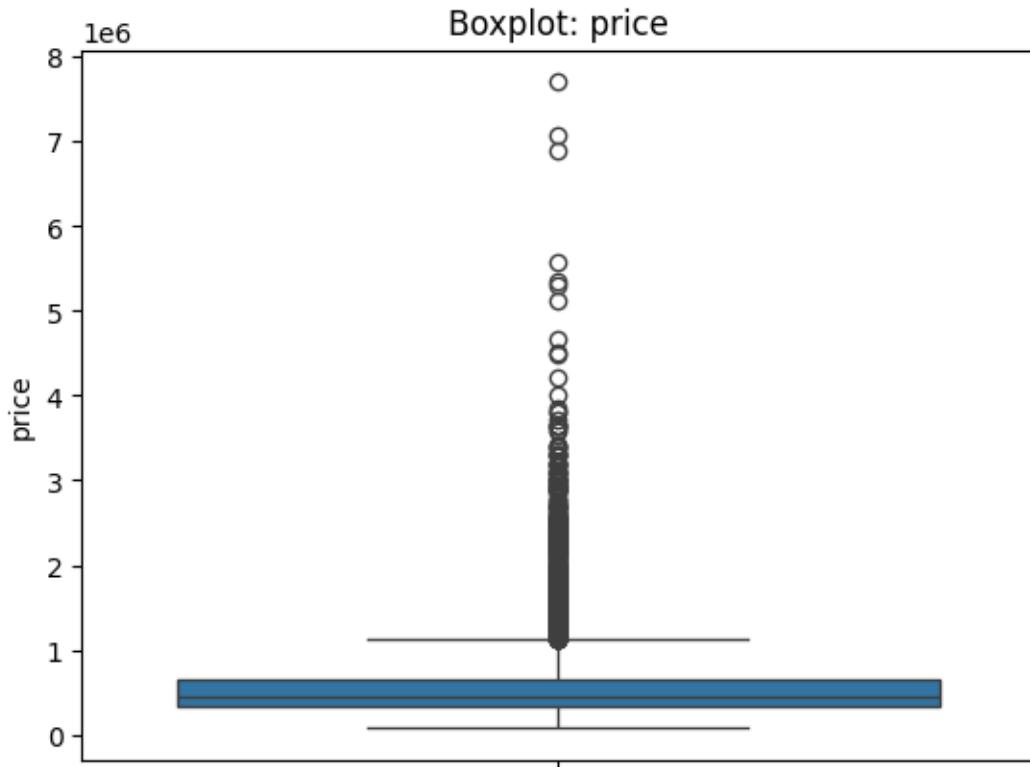
```
[4]: sns.histplot(df_clean.price)
plt.title("Histogram: price")
```

```
[4]: Text(0.5, 1.0, 'Histogram: price')
```



```
[5]: sns.boxplot(df_clean.price)
plt.title("Boxplot: price")
```

```
[5]: Text(0.5, 1.0, 'Boxplot: price')
```



```
[6]: df_clean["price"].describe()
```

```
[6]: count      2.161300e+04
mean       5.400881e+05
std        3.671272e+05
min        7.500000e+04
25%        3.219500e+05
50%        4.500000e+05
75%        6.450000e+05
max        7.700000e+06
Name: price, dtype: float64
```

```
[7]: # 11 houses over 4 million
df_clean[df_clean.price > 4e6]
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
1164	5110800.0	5	5.25	8010	45517	2.0	
1315	5300000.0	6	6.00	7390	24829	2.0	
1448	5350000.0	5	5.00	8000	23985	2.0	
2626	4500000.0	5	5.50	6640	40014	2.0	
3914	7062500.0	5	4.50	10040	37325	2.0	
4411	5570000.0	5	5.75	9200	35069	2.0	

7252	7700000.0	6	8.00	12050	27600	2.5		
8092	4668000.0	5	6.75	9640	13068	1.0		
8638	4489000.0	4	3.00	6430	27517	2.0		
9254	6885000.0	6	7.75	9890	31374	2.0		
12370	4208000.0	5	6.00	7440	21540	2.0		
	waterfront	view	condition	grade	...	yr_built	yr_renovated	\
1164	1	4	3	12	...	1999	0	
1315	1	4	4	12	...	1991	0	
1448	0	4	3	12	...	2009	0	
2626	1	4	3	12	...	2004	0	
3914	1	2	3	11	...	1940	2001	
4411	0	0	3	13	...	2001	0	
7252	0	3	4	13	...	1910	1987	
8092	1	4	3	12	...	1983	2009	
8638	0	0	3	12	...	2001	0	
9254	0	4	3	13	...	2001	0	
12370	0	0	3	12	...	2003	0	
	zipcode	lat	long	sqft_living15	sqft_lot15	year_sold	\	
1164	98033	47.6767	-122.211	3430	26788	2014		
1315	98040	47.5631	-122.210	4320	24619	2015		
1448	98004	47.6232	-122.220	4600	21750	2015		
2626	98155	47.7493	-122.280	3030	23408	2014		
3914	98004	47.6500	-122.214	3930	25449	2014		
4411	98039	47.6289	-122.233	3560	24345	2014		
7252	98102	47.6298	-122.323	3940	8800	2014		
8092	98040	47.5570	-122.210	3270	10454	2014		
8638	98004	47.6208	-122.219	3720	14592	2014		
9254	98039	47.6305	-122.240	4540	42730	2014		
12370	98006	47.5692	-122.189	4740	19329	2015		
	month_sold	day_sold						
1164	10	20						
1315	4	13						
1448	4	13						
2626	8	15						
3914	6	11						
4411	8	4						
7252	10	13						
8092	6	17						
8638	6	18						
9254	9	19						
12370	5	6						

[11 rows x 22 columns]

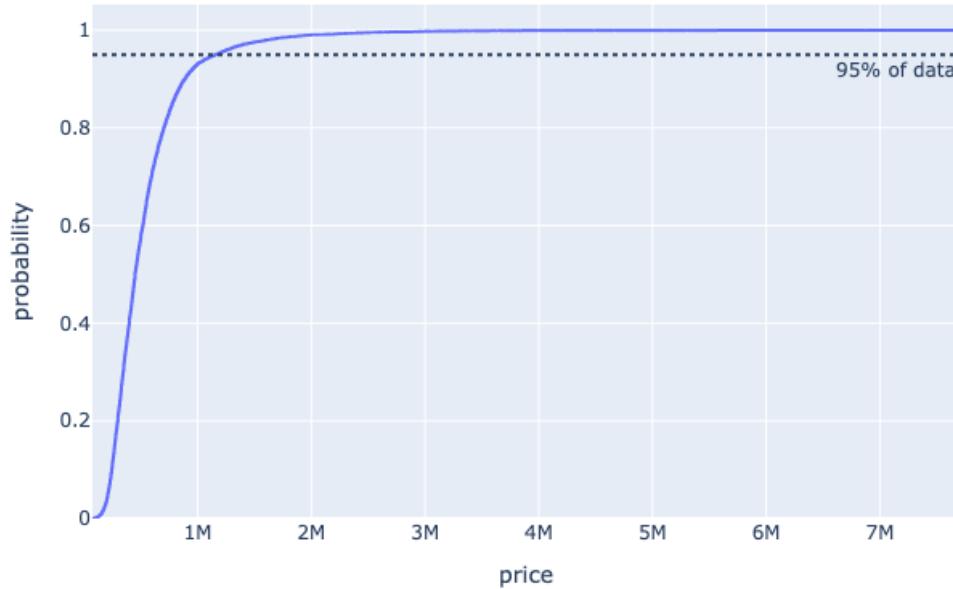
## 1.4 How much of the data is represented by outliers?

```
[8]: df_sorted = df_clean["price"].sort_values()

# Cumulative Quatiles
fig = px.ecdf(df_sorted, x="price", title="Cumulative Distribution of House Prices")

# Add a marker line at the 95th and 99th percentiles
fig.add_hline(y=0.95, line_dash="dot", annotation_text="95% of data", annotation_position="bottom right")
fig.show()
```

Cumulative Distribution of House Prices



```
[9]: # quantiles
quantile_99 = df_clean.price.quantile(0.99)
quantile_95 = df_clean.price.quantile(0.95)

print("Quantile 99: ", quantile_99)
print("Quantile 95: ", quantile_95)
```

Quantile 99: 1964400.0000000051

Quantile 95: 1156479.999999974

```
[10]: # Flag the datapoints inside each quantiles 99 and 95
df_clean["q_99"] = (df_clean.price < quantile_99).astype(int)
df_clean["q_95"] = (df_clean.price < quantile_95).astype(int)

df_clean.head()
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	\
0	221900.0	3	1.00	1180	5650	1.0	0	
1	538000.0	3	2.25	2570	7242	2.0	0	
2	180000.0	2	1.00	770	10000	1.0	0	
3	604000.0	4	3.00	1960	5000	1.0	0	
4	510000.0	3	2.00	1680	8080	1.0	0	

	view	condition	grade	...	zipcode	lat	long	sqft_living15	\
0	0	3	7	...	98178	47.5112	-122.257	1340	
1	0	3	7	...	98125	47.7210	-122.319	1690	
2	0	3	6	...	98028	47.7379	-122.233	2720	
3	0	5	7	...	98136	47.5208	-122.393	1360	
4	0	3	8	...	98074	47.6168	-122.045	1800	

	sqft_lot15	year_sold	month_sold	day_sold	q_99	q_95		
0	5650	2014		10	13	1	1	
1	7639	2014		12	9	1	1	
2	8062	2015		2	25	1	1	
3	5000	2014		12	9	1	1	
4	7503	2015		2	18	1	1	

[5 rows x 24 columns]

```
[11]: print(f"The 99th quantile exclude {df_clean.shape[0] - df_clean.q_99.sum()} datapoints")
print(f"The 95th quantile exclude {df_clean.shape[0] - df_clean.q_95.sum()} datapoints")
```

The 99th quantile exclude 217 datapoints  
The 95th quantile exclude 1081 datapoints

## 1.5 Sensitivity to outliers (metrics)

```
[12]: metrics_df = create_metrics_df()
```

### 1.5.1 Random Forest flagging the outliers

```
[13]: # Split into train and test
seed = 13
# The price is the target variable
y = df_clean["price"]
```

```
# All other variables are the features for the baseline model
X = df_clean.drop(["price"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↴random_state=seed)
```

[14]: X\_train

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	\
1571	4	1.50	2000	6778	1.0	0	0	
16330	4	2.50	2630	48706	2.0	0	0	
12786	4	2.50	2620	9525	2.5	0	0	
12524	3	2.50	1610	6000	2.0	0	0	
16179	3	1.00	880	18205	1.0	0	0	
...	...	...	...	...	...	...	...	
153	4	3.25	5180	19850	2.0	0	3	
866	3	2.50	3460	6590	2.0	0	0	
74	3	1.75	1790	50529	1.0	0	0	
14512	2	1.00	820	5040	1.0	0	0	
338	3	1.75	1420	8250	1.0	0	0	
	condition	grade	sqft_above	...	zipcode	lat	long	\
1571	4	7	1170	...	98198	47.3708	-122.311	
16330	3	8	2630	...	98072	47.7750	-122.125	
12786	4	9	2620	...	98040	47.5631	-122.219	
12524	4	7	1610	...	98038	47.3490	-122.036	
16179	4	6	880	...	98178	47.5013	-122.244	
...	...	...	...	...	...	...	...	
153	3	12	3540	...	98006	47.5620	-122.162	
866	3	7	3460	...	98056	47.4802	-122.188	
74	5	7	1090	...	98042	47.3511	-122.073	
14512	3	7	820	...	98199	47.6498	-122.388	
338	3	7	1420	...	98133	47.7535	-122.354	
	sqft_living15	sqft_lot15	year_sold	month_sold	day_sold	q_99	q_95	
1571	1940	7531	2015	3	23	1	1	
16330	2680	48706	2014	5	21	1	1	
12786	2580	9525	2014	8	5	1	1	
12524	1570	6000	2014	8	26	1	1	
16179	1110	16115	2014	6	24	1	1	
...	...	...	...	...	...	...	...	
153	3160	9750	2015	4	1	0	0	
866	2490	6312	2015	4	27	1	1	
74	1940	50529	2015	3	16	1	1	
14512	1730	5760	2014	8	20	1	1	

```
338          1740        8000       2014         8          26         1         1
```

[17290 rows x 23 columns]

```
[15]: # most common hyperparameters or the default ones
```

```
rf_regressor = RandomForestRegressor(random_state=seed)#default values +
↳random_state = 13
rf_regressor.fit(X_train, y_train)

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_train,
                             y_train,
                             split = "train",
                             comments="Outliers flagging, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Outliers flagging, no normalization.")
```

```
[16]: metrics_df
```

```
Model   Split      R2  Adjusted_R2      MAE    MAPE \
0  RandomForestRegressor  train  0.9879      0.9879  22769.7215  0.0456
1  RandomForestRegressor  test   0.9322      0.9319  59863.8409  0.1203

RMSE                               Comments
0  40566.2303  Outliers flagging, no normalization.
1  93599.4998  Outliers flagging, no normalization.
```

## 1.5.2 Random Forest removing top 1% outliers

```
# Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_99 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_99 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,↳
↳random_state=seed)
```

[18]: X\_train

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	\
16385	3	1.50	1290	8175	1.0	0	0	
16816	3	2.00	1830	10873	1.0	0	0	
6038	3	1.75	1840	8086	1.0	0	0	
16785	3	1.75	1540	10545	2.0	0	0	
15744	3	2.00	2090	15790	1.0	0	0	
...	...	...	...	...	...	...	...	
155	3	1.00	1180	7669	1.0	0	0	
878	3	1.00	1020	8100	1.0	0	0	
75	4	4.00	3430	35102	2.0	0	0	
14650	4	3.00	2530	5625	1.0	0	0	
345	4	1.00	1000	7134	1.0	0	0	
	condition	grade	sqft_above	...	yr_built	yr_renovated	zipcode	\
16385	4	7	820	...	1952	0	98004	
16816	3	8	1830	...	1989	0	98023	
6038	4	8	1840	...	1964	0	98052	
16785	4	6	1540	...	1978	0	98045	
15744	3	9	2090	...	1992	0	98034	
...	...	...	...	...	...	...	...	
155	4	7	1180	...	1967	0	98058	
878	3	7	1020	...	1954	0	98198	
75	4	10	2390	...	1986	0	98075	
14650	3	8	1470	...	1976	0	98034	
345	3	6	1000	...	1943	0	98178	
	lat	long	sqft_living15	sqft_lot15	year_sold	month_sold	\	
16385	47.6296	-122.205	2130	8577	2014	8		
16816	47.3066	-122.394	2490	8976	2014	7		
6038	47.6700	-122.155	1840	8060	2014	6		
16785	47.4451	-121.763	1540	10000	2015	3		
15744	47.7296	-122.199	1820	8770	2014	11		
...	...	...	...	...	...	...	...	
155	47.4479	-122.176	1190	7669	2014	7		
878	47.3586	-122.314	1020	8100	2014	12		
75	47.5822	-121.987	3240	35020	2014	11		
14650	47.7094	-122.233	1840	7070	2014	7		
345	47.4897	-122.240	1020	7138	2014	7		
	day_sold							
16385		20						
16816		31						
6038		23						
16785		11						
15744		20						

```

...
155      28
878      19
75       5
14650    11
345      23

```

[17116 rows x 21 columns]

[19]: # most common hyperparameters or the default ones

```

rf_regressor = RandomForestRegressor(random_state=seed)#default values +
random_state = 13
rf_regressor.fit(X_train, y_train)

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_train,
                             y_train,
                             split = "train",
                             comments="Removing top 1%, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 1%, no normalization.")

```

[20]: metrics\_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
	RMSE			Comments			
0	40566.2303			Outliers flagging, no normalization.			
1	93599.4998			Outliers flagging, no normalization.			
2	37273.0054			Removing top 1%, no normalization.			
3	104088.2866			Removing top 1%, no normalization.			

### 1.5.3 Random Forest removing top 5% outliers

```
[21]: # Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_95 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_95 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=seed)
```

```
[22]: X_train
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	\
7523	3	1.00	1200	7810	1.0	0	0	
20438	4	2.50	3250	4500	2.0	0	0	
8571	4	2.50	2180	3893	2.0	0	0	
15677	3	1.75	1990	5600	1.0	0	1	
8032	3	1.75	2360	4063	1.0	0	0	
...	...	...	...	...	...	...	...	
159	4	1.75	1760	11180	1.0	0	0	
907	3	1.75	1410	9315	1.0	0	0	
78	3	1.00	1410	5060	1.0	0	0	
15224	4	2.50	2570	8178	1.0	0	2	
353	2	1.00	990	3120	1.0	0	2	
	condition	grade	sqft_above	...	yr_built	yr_renovated	zipcode	\
7523	4	7	1200	...	1967	0	98038	
20438	3	8	3250	...	2008	0	98059	
8571	3	8	2180	...	1999	0	98117	
15677	3	8	1330	...	1941	0	98199	
8032	5	7	1180	...	1940	0	98117	
...	...	...	...	...	...	...	...	
159	4	8	1760	...	1968	0	98059	
907	5	7	1410	...	1960	0	98031	
78	4	7	910	...	1956	0	98133	
15224	3	8	1710	...	1961	0	98118	
353	5	7	790	...	1907	0	98103	
	lat	long	sqft_living15	sqft_lot15	year_sold	month_sold	\	
7523	47.3631	-122.050	1590	7800	2015	4		
20438	47.4944	-122.150	3030	4598	2014	11		
8571	47.6886	-122.388	1710	4550	2014	12		
15677	47.6500	-122.415	2630	6780	2014	9		
8032	47.6902	-122.382	1660	4063	2014	8		

```

...
159    47.4715 -122.118      ...    1730    ...    11180    ...    2014    ...
907    47.3969 -122.198      ...    1630    ...    8250    ...    2014    ...
78     47.7073 -122.340      ...    1130    ...    5693    ...    2014    ...
15224   47.5483 -122.261      ...    2050    ...    7500    ...    2015    ...
353    47.6800 -122.353      ...    1930    ...    3120    ...    2014    ...
day_sold
7523      25
20438     4
8571      8
15677     3
8032      18
...
159      3
907      7
78       9
15224    9
353      3

```

[16425 rows x 21 columns]

[23]: # most common hyperparameters or the default ones

```

rf_regressor = RandomForestRegressor(random_state=seed)#default values +
↳random_state = 13
rf_regressor.fit(X_train, y_train)

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_train,
                             y_train,
                             split = "train",
                             comments="Removing top 5%, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 5%, no normalization.")

```

[24]: metrics\_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	

1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458
5	RandomForestRegressor	test	0.8717	0.8710	53038.5069	0.1212

	RMSE	Comments
0	40566.2303	Outliers flagging, no normalization.
1	93599.4998	Outliers flagging, no normalization.
2	37273.0054	Removing top 1%, no normalization.
3	104088.2866	Removing top 1%, no normalization.
4	29386.2798	Removing top 5%, no normalization.
5	76308.9748	Removing top 5%, no normalization.

### 1.5.4 XGBoost flagging the outliers

```
[25]: # Split into train and test
seed = 13
# The price is the target variable
y = df_clean["price"]

# All other variables are the features for the baseline model
X = df_clean.drop(["price"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=seed)
```

```
[26]: xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[26]: XGBRegressor(base_score=None, booster=None, callbacks=None,
       colsample_bylevel=None, colsample_bynode=None,
       colsample_bytree=None, device=None, early_stopping_rounds=None,
       enable_categorical=False, eval_metric=None, feature_types=None,
       feature_weights=None, gamma=None, grow_policy=None,
       importance_type=None, interaction_constraints=None,
       learning_rate=None, max_bin=None, max_cat_threshold=None,
       max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
       max_leaves=None, min_child_weight=None, missing=nan,
       monotone_constraints=None, multi_strategy=None, n_estimators=None,
       n_jobs=None, num_parallel_tree=None, ...)
```

```
[27]: metrics_df = add_new_metrics(metrics_df,
                                  xgb_clf,
                                  X_train,
                                  y_train,
```

```

        split = "train",
        comments="Outliers flagging, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             xgb_clf,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Outliers flagging, no normalization.")

```

[28]: metrics\_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458	
5	RandomForestRegressor	test	0.8717	0.8710	53038.5069	0.1212	
6	XGBRegressor	train	0.9825	0.9825	35584.0027	0.0818	
7	XGBRegressor	test	0.9260	0.9256	59532.0485	0.1184	
	RMSE			Comments			
0	40566.2303			Outliers flagging, no normalization.			
1	93599.4998			Outliers flagging, no normalization.			
2	37273.0054			Removing top 1%, no normalization.			
3	104088.2866			Removing top 1%, no normalization.			
4	29386.2798			Removing top 5%, no normalization.			
5	76308.9748			Removing top 5%, no normalization.			
6	48751.1351			Outliers flagging, no normalization.			
7	97798.9764			Outliers flagging, no normalization.			

### 1.5.5 XGBoost removing top 1%

```

[29]: # Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_99 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_99 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=seed)

```

[30]: X\_train

[30] :

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	\
16385	3	1.50	1290	8175	1.0	0	0	
16816	3	2.00	1830	10873	1.0	0	0	
6038	3	1.75	1840	8086	1.0	0	0	
16785	3	1.75	1540	10545	2.0	0	0	
15744	3	2.00	2090	15790	1.0	0	0	
...	...	...	...	...	...	...	...	
155	3	1.00	1180	7669	1.0	0	0	
878	3	1.00	1020	8100	1.0	0	0	
75	4	4.00	3430	35102	2.0	0	0	
14650	4	3.00	2530	5625	1.0	0	0	
345	4	1.00	1000	7134	1.0	0	0	
	condition	grade	sqft_above	...	yr_built	yr_renovated	zipcode	\
16385	4	7	820	...	1952	0	98004	
16816	3	8	1830	...	1989	0	98023	
6038	4	8	1840	...	1964	0	98052	
16785	4	6	1540	...	1978	0	98045	
15744	3	9	2090	...	1992	0	98034	
...	...	...	...	...	...	...	...	
155	4	7	1180	...	1967	0	98058	
878	3	7	1020	...	1954	0	98198	
75	4	10	2390	...	1986	0	98075	
14650	3	8	1470	...	1976	0	98034	
345	3	6	1000	...	1943	0	98178	
	lat	long	sqft_living15	sqft_lot15	year_sold	month_sold	\	
16385	47.6296	-122.205	2130	8577	2014	8		
16816	47.3066	-122.394	2490	8976	2014	7		
6038	47.6700	-122.155	1840	8060	2014	6		
16785	47.4451	-121.763	1540	10000	2015	3		
15744	47.7296	-122.199	1820	8770	2014	11		
...	...	...	...	...	...	...		
155	47.4479	-122.176	1190	7669	2014	7		
878	47.3586	-122.314	1020	8100	2014	12		
75	47.5822	-121.987	3240	35020	2014	11		
14650	47.7094	-122.233	1840	7070	2014	7		
345	47.4897	-122.240	1020	7138	2014	7		
	day_sold							
16385		20						
16816		31						
6038		23						
16785		11						
15744		20						
...	...							
155		28						

```
878      19
75       5
14650    11
345     23
```

```
[17116 rows x 21 columns]
```

```
[31]: xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[31]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

```
[32]: metrics_df = add_new_metrics(metrics_df,
                                    xgb_clf,
                                    X_train,
                                    y_train,
                                    split = "train",
                                    comments="Removing top 1%, no normalization.")
```

```
metrics_df = add_new_metrics(metrics_df,
                             xgb_clf,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 1%, no normalization.")
```

```
[33]: metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458	
5	RandomForestRegressor	test	0.8717	0.8710	53038.5069	0.1212	
6	XGBRegressor	train	0.9825	0.9825	35584.0027	0.0818	

```

7          XGBRegressor    test  0.9260      0.9256  59532.0485  0.1184
8          XGBRegressor    train  0.9679      0.9679  36410.9174  0.0830
9          XGBRegressor    test  0.8891      0.8885  61265.8381  0.1234

          RMSE           Comments
0  40566.2303  Outliers flagging, no normalization.
1  93599.4998  Outliers flagging, no normalization.
2  37273.0054  Removing top 1%, no normalization.
3  104088.2866  Removing top 1%, no normalization.
4  29386.2798  Removing top 5%, no normalization.
5  76308.9748  Removing top 5%, no normalization.
6  48751.1351  Outliers flagging, no normalization.
7  97798.9764  Outliers flagging, no normalization.
8  50584.7096  Removing top 1%, no normalization.
9  97618.3705  Removing top 1%, no normalization.

```

### 1.5.6 XGBoost removing top 5%

```
[34]: # Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_95 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_95 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=seed)
```

```
[35]: xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[35]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                  colsample_bylevel=None, colsample_bynode=None,
                  colsample_bytree=None, device=None, early_stopping_rounds=None,
                  enable_categorical=False, eval_metric=None, feature_types=None,
                  feature_weights=None, gamma=None, grow_policy=None,
                  importance_type=None, interaction_constraints=None,
                  learning_rate=None, max_bin=None, max_cat_threshold=None,
                  max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                  max_leaves=None, min_child_weight=None, missing=nan,
                  monotone_constraints=None, multi_strategy=None, n_estimators=None,
                  n_jobs=None, num_parallel_tree=None, ...)
```

```
[36]: metrics_df = add_new_metrics(metrics_df,
                                  xgb_clf,
                                  X_train,
```

```

        y_train,
        split = "train",
        comments="Removing top 5%, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             xgb_clf,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 5%, no normalization.")

```

[37] : metrics\_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458	
5	RandomForestRegressor	test	0.8717	0.8710	53038.5069	0.1212	
6	XGBRegressor	train	0.9825	0.9825	35584.0027	0.0818	
7	XGBRegressor	test	0.9260	0.9256	59532.0485	0.1184	
8	XGBRegressor	train	0.9679	0.9679	36410.9174	0.0830	
9	XGBRegressor	test	0.8891	0.8885	61265.8381	0.1234	
10	XGBRegressor	train	0.9569	0.9568	32103.1211	0.0783	
11	XGBRegressor	test	0.8806	0.8800	51602.5277	0.1189	

	RMSE	Comments
0	40566.2303	Outliers flagging, no normalization.
1	93599.4998	Outliers flagging, no normalization.
2	37273.0054	Removing top 1%, no normalization.
3	104088.2866	Removing top 1%, no normalization.
4	29386.2798	Removing top 5%, no normalization.
5	76308.9748	Removing top 5%, no normalization.
6	48751.1351	Outliers flagging, no normalization.
7	97798.9764	Outliers flagging, no normalization.
8	50584.7096	Removing top 1%, no normalization.
9	97618.3705	Removing top 1%, no normalization.
10	43810.9507	Removing top 5%, no normalization.
11	73599.3380	Removing top 5%, no normalization.

The feature flagging for both 99% and 95% worked better than removing the outliers, which means that the information about the top priced houses is still important to accurately predict the prices. In that sense, we will continue the analysis using the flagging of the columns instead of dropping them.

## 2 Data Leakage Analysis: Outlier Flagging

Our initial approach calculated outlier thresholds using the entire dataset, introducing data leakage by allowing test set information to influence training features. While this is common in exploratory phases, it risks inflating model performance. In this section, we apply hypothesis testing to determine if correcting this leakage results in a statistically significant difference in model predictions.

```
[38]: # Apply outlier flagging
def apply_flagging(X, y):
    # quantiles
    quantile_99 = y.quantile(0.99)
    quantile_95 = y.quantile(0.95)

    print("Quantile 99: ", quantile_99)
    print("Quantile 95: ", quantile_95)

    # Flag the datapoints inside each quantiles 99 and 95
    X["q_99"] = (y < quantile_99).astype(int)
    X["q_95"] = (y < quantile_95).astype(int)

    return X
```

### 2.1 Preparing Datasets

```
[39]: # Split into train and test
seed = 13
# The price is the target variable
y = df_clean["price"]

# All other variables are the features for the baseline model
X = df_clean.drop(["price"], axis=1)
```

#### 2.1.1 Dataset without Leakage

```
[40]: # Data without leakage
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=seed)

# Flagging
X_train = apply_flagging(X_train, y_train)
X_test = apply_flagging(X_test, y_test)
```

```
Quantile 99: 1980000.0
Quantile 95: 1150000.0
Quantile 99: 1945159.9999999944
Quantile 95: 1185000.0
```

## 2.1.2 Dataset with Leakage

```
[41]: # Data with leakage
```

```
X_leakage = apply_flagging(X, y)

X_train_leak, X_test_leak, y_train_leak, y_test_leak = train_test_split(X_leakage, y, test_size=0.2, random_state=seed)
```

Quantile 99: 1964400.0000000051

Quantile 95: 1156479.9999999974

## 2.1.3 Experiment metrics df

```
[42]: # New metrics dataframe
```

```
leak_df = create_metrics_df()
leak_df
```

[42]: Empty DataFrame

Columns: [Model, Split, R2, Adjusted\_R2, MAE, RMSE, MAPE, Comments]

Index: []

## 2.2 Statistical Significance Test: Random Forest Performance

### 2.2.1 Without leakage

```
[43]: # most common hyperparameters or the default ones
```

```
rf_regressor = RandomForestRegressor(random_state=seed)#default values +
random_state = 13
rf_regressor.fit(X_train, y_train)

leak_df = add_new_metrics(leak_df,
                          rf_regressor,
                          X_train,
                          y_train,
                          split = "train",
                          comments="No leakage.")

leak_df = add_new_metrics(leak_df,
                          rf_regressor,
                          X_test,
                          y_test,
                          split = "test",
                          comments="No leakage.")
```

```
leak_df
```

```
[43]:
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9880	0.9879	22693.7463	0.0455	
1	RandomForestRegressor	test	0.9331	0.9327	59616.8735	0.1197	

	RMSE	Comments
0	40491.8085	No leakage.
1	92989.3473	No leakage.

## 2.2.2 With leakage

```
[44]: # most common hyperparameters or the default ones
```

```
rf_regressor_leak = RandomForestRegressor(random_state=seed)#default values +  
↳random_state = 13  
rf_regressor_leak.fit(X_train_leak, y_train_leak)  
  
leak_df = add_new_metrics(leak_df,  
                           rf_regressor_leak,  
                           X_train_leak,  
                           y_train_leak,  
                           split = "train",  
                           comments="with leakage.")  
  
leak_df = add_new_metrics(leak_df,  
                           rf_regressor_leak,  
                           X_test_leak,  
                           y_test_leak,  
                           split = "test",  
                           comments="with leakage.")  
  
leak_df
```

```
[44]:
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9880	0.9879	22693.7463	0.0455	
1	RandomForestRegressor	test	0.9331	0.9327	59616.8735	0.1197	
2	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
3	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	

	RMSE	Comments
0	40491.8085	No leakage.
1	92989.3473	No leakage.
2	40566.2303	with leakage.
3	93599.4998	with leakage.

```
[45]: # Hypothesis testing for the test scores

# Ho: scores_leakage = scores_no_leakage
# Ha: scores_leakage != scores_no_leakage

# Method 1: With leakage
scores_leakage = cross_val_score(rf_regressor_leak, X_test_leak, y_test_leak,
                                 cv=5, scoring='r2')

# Method 2: Without leakage
scores_no_leakage = cross_val_score(rf_regressor, X_test, y_test,
                                      cv=5, scoring='r2')

# Paired t-test (same folds, so paired)
t_stat, p_value = stats.ttest_rel(scores_leakage, scores_no_leakage)

print(f"Mean R2 with leakage: {scores_leakage.mean():.4f} ± {scores_leakage.
    ~std():.4f}")
print(f"Mean R2 without leakage: {scores_no_leakage.mean():.4f} ±
    ~{scores_no_leakage.std():.4f}")
print(f"p-value: {p_value:.4f}")

if p_value < 0.05:
    print("Difference is statistically significant. We reject the null
          hypothesis.")
else:
    print("Difference is NOT statistically significant. We cannot reject the
          null hypothesis.")
```

Mean R<sup>2</sup> with leakage: 0.9049 ± 0.0178  
 Mean R<sup>2</sup> without leakage: 0.9061 ± 0.0180  
 p-value: 0.1119  
 Difference is NOT statistically significant. We cannot reject the null hypothesis.

```
[46]: # Hypothesis testing for the train scores

# Ho: scores_leakage = scores_no_leakage
# Ha: scores_leakage != scores_no_leakage

# Method 1: With leakage
scores_leakage = cross_val_score(rf_regressor_leak, X_train_leak, y_train_leak,
                                 cv=5, scoring='r2')

# Method 2: Without leakage
scores_no_leakage = cross_val_score(rf_regressor, X_train, y_train,
                                      cv=5, scoring='r2')
```

```

# Paired t-test (same folds, so paired)
t_stat, p_value = stats.ttest_rel(scores_leakage, scores_no_leakage)

print(f"Mean R2 with leakage: {scores_leakage.mean():.4f} ± {scores_leakage.
    ↴std():.4f}")
print(f"Mean R2 without leakage: {scores_no_leakage.mean():.4f} ±
    ↴{scores_no_leakage.std():.4f}")
print(f"p-value: {p_value:.4f}")

if p_value < 0.05:
    print("Difference is statistically significant. We reject the null
        ↴hypothesis.")
else:
    print("Difference is NOT statistically significant. We cannot reject the
        ↴null hypothesis.")

```

```

Mean R2 with leakage: 0.9054 ± 0.0179
Mean R2 without leakage: 0.9051 ± 0.0180
p-value: 0.6313
Difference is NOT statistically significant. We cannot reject the null
hypothesis.

```

## 2.3 Statistical Significance Test: XGBoost Performance

### 2.3.1 No leakage

```

[47]: xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)

leak_df = add_new_metrics(leak_df,
                          xgb_clf,
                          X_train,
                          y_train,
                          split = "train",
                          comments="No leakage.")

leak_df = add_new_metrics(leak_df,
                          xgb_clf,
                          X_test,
                          y_test,
                          split = "test",
                          comments="No leakage.")

leak_df

```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9880	0.9879	22693.7463	0.0455	

1	RandomForestRegressor	test	0.9331	0.9327	59616.8735	0.1197
2	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456
3	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203
4	XGBRegressor	train	0.9820	0.9820	35896.1715	0.0825
5	XGBRegressor	test	0.9240	0.9236	59986.6293	0.1188

	RMSE	Comments
0	40491.8085	No leakage.
1	92989.3473	No leakage.
2	40566.2303	with leakage.
3	93599.4998	with leakage.
4	49457.6309	No leakage.
5	99079.9228	No leakage.

```
[48]: xgb_clf_leak = xgb.XGBRegressor(seed = seed)
xgb_clf_leak.fit(X_train_leak, y_train_leak)
```

```
leak_df = add_new_metrics(leak_df,
                           xgb_clf_leak,
                           X_train_leak,
                           y_train_leak,
                           split = "train",
                           comments="with leakage.")

leak_df = add_new_metrics(leak_df,
                           xgb_clf_leak,
                           X_test_leak,
                           y_test_leak,
                           split = "test",
                           comments="with leakage.")
```

```
leak_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9880	0.9879	22693.7463	0.0455	
1	RandomForestRegressor	test	0.9331	0.9327	59616.8735	0.1197	
2	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
3	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
4	XGBRegressor	train	0.9820	0.9820	35896.1715	0.0825	
5	XGBRegressor	test	0.9240	0.9236	59986.6293	0.1188	
6	XGBRegressor	train	0.9825	0.9825	35584.0027	0.0818	
7	XGBRegressor	test	0.9260	0.9256	59532.0485	0.1184	

	RMSE	Comments
0	40491.8085	No leakage.
1	92989.3473	No leakage.
2	40566.2303	with leakage.

```

3 93599.4998 with leakage.
4 49457.6309    No leakage.
5 99079.9228    No leakage.
6 48751.1351    with leakage.
7 97798.9764    with leakage.

```

```
[49]: # Hypothesis testing for the test scores

# Ho: scores_leakage = scores_no_leakage
# Ha: scores_leakage != scores_no_leakage

# Method 1: With leakage
scores_leakage = cross_val_score(xgb_clf_leak, X_test_leak, y_test_leak,
                                 cv=5, scoring='r2')

# Method 2: Without leakage
scores_no_leakage = cross_val_score(xgb_clf, X_test, y_test,
                                     cv=5, scoring='r2')

# Paired t-test (same folds, so paired)
t_stat, p_value = stats.ttest_rel(scores_leakage, scores_no_leakage)

print(f"Mean R2 with leakage: {scores_leakage.mean():.4f} ± {scores_leakage.
    std():.4f}")
print(f"Mean R2 without leakage: {scores_no_leakage.mean():.4f} ±
    {scores_no_leakage.std():.4f}")
print(f"p-value: {p_value:.4f}")

if p_value < 0.05:
    print("Difference is statistically significant. We reject the null
        hypothesis.")
else:
    print("Difference is NOT statistically significant. We cannot reject the
        null hypothesis.")


```

```

Mean R2 with leakage: 0.9073 ± 0.0139
Mean R2 without leakage: 0.9064 ± 0.0166
p-value: 0.6127
Difference is NOT statistically significant. We cannot reject the null
hypothesis.

```

```
[50]: # Hypothesis testing for the train scores

# Ho: scores_leakage = scores_no_leakage
# Ha: scores_leakage != scores_no_leakage

# Method 1: With leakage
```

```

scores_leakage = cross_val_score(xgb_clf_leak, X_train_leak, y_train_leak,
                                 cv=5, scoring='r2')

# Method 2: Without leakage
scores_no_leakage = cross_val_score(xgb_clf, X_train, y_train,
                                    cv=5, scoring='r2')

# Paired t-test (same folds, so paired)
t_stat, p_value = stats.ttest_rel(scores_leakage, scores_no_leakage)

print(f"Mean R2 with leakage: {scores_leakage.mean():.4f} ± {scores_leakage.
    std():.4f}")
print(f"Mean R2 without leakage: {scores_no_leakage.mean():.4f} ±
    {scores_no_leakage.std():.4f}")
print(f"p-value: {p_value:.4f}")

if p_value < 0.05:
    print("Difference is statistically significant. We reject the null
        hypothesis.")
else:
    print("Difference is NOT statistically significant. We cannot reject the
        null hypothesis.")

```

```

Mean R2 with leakage: 0.9063 ± 0.0137
Mean R2 without leakage: 0.9095 ± 0.0149
p-value: 0.5156
Difference is NOT statistically significant. We cannot reject the null
hypothesis.

```

## 2.4 Conclusion on Leakage

Statistical testing showed no significant difference between the approaches. For the scope of this analysis, we will proceed with the current implementation for simplicity, while acknowledging that a production-grade deployment would require the leakage-free pipeline to ensure strict data isolation.

## 2.5 Export Dataset and Metrics

```
[51]: # filename_metrics = "Metrics/outlier_analysis_metrics.csv"
      # metrics_df.to_csv(filename_metrics, index = False)
```

```
[52]: # df_clean
```

```
[53]: # filename_data = "Data/v1_house_sales.csv"
      # df_clean.to_csv(filename_data, index = False)
```