

3_Outliers_Analysis

February 13, 2026

1 Analyse the Influence of Outliers

1.1 Imports

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Models & Normalization
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
```

1.2 Open the clean data

```
[2]: file_path = "Data/cleaned_house_sales.csv"

df_clean = pd.read_csv(file_path)
df_clean.head()
```

```
[2]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront \
0  221900.0        3       1.00      1180      5650     1.0          0
1  538000.0        3       2.25      2570      7242     2.0          0
2  180000.0        2       1.00       770     10000     1.0          0
3  604000.0        4       3.00      1960      5000     1.0          0
4  510000.0        3       2.00      1680      8080     1.0          0

      view  condition  grade  ...  yr_built  yr_renovated  zipcode      lat \
0      0         3      7  ...    1955                  0    98178  47.5112
1      0         3      7  ...    1951                 1991    98125  47.7210
2      0         3      6  ...    1933                  0    98028  47.7379
3      0         5      7  ...    1965                  0    98136  47.5208
4      0         3      8  ...    1987                  0    98074  47.6168

      long  sqft_living15  sqft_lot15  year_sold  month_sold  day_sold
0 -122.257           1340        5650     2014         10        13
1 -122.319           1690        7639     2014         12         9
```

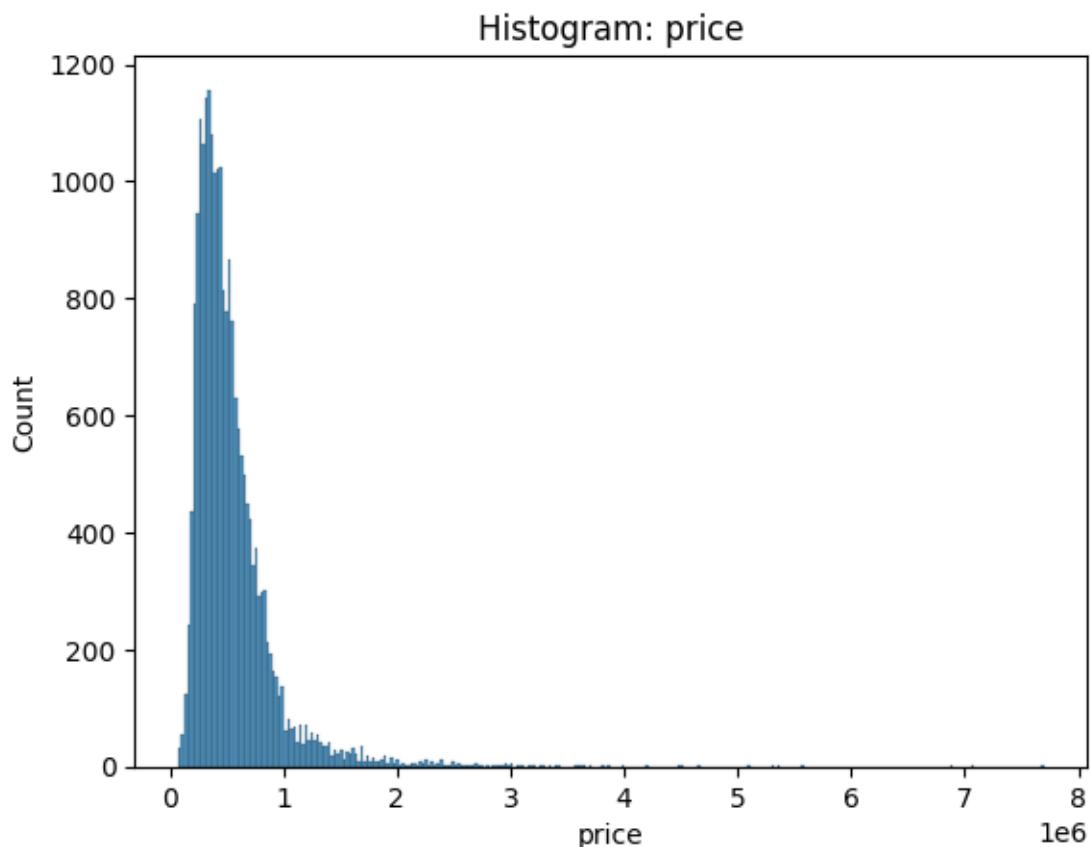
```
2 -122.233          2720        8062      2015         2       25
3 -122.393          1360        5000      2014        12        9
4 -122.045          1800        7503      2015         2       18
```

[5 rows x 22 columns]

1.3 Analysing Prices

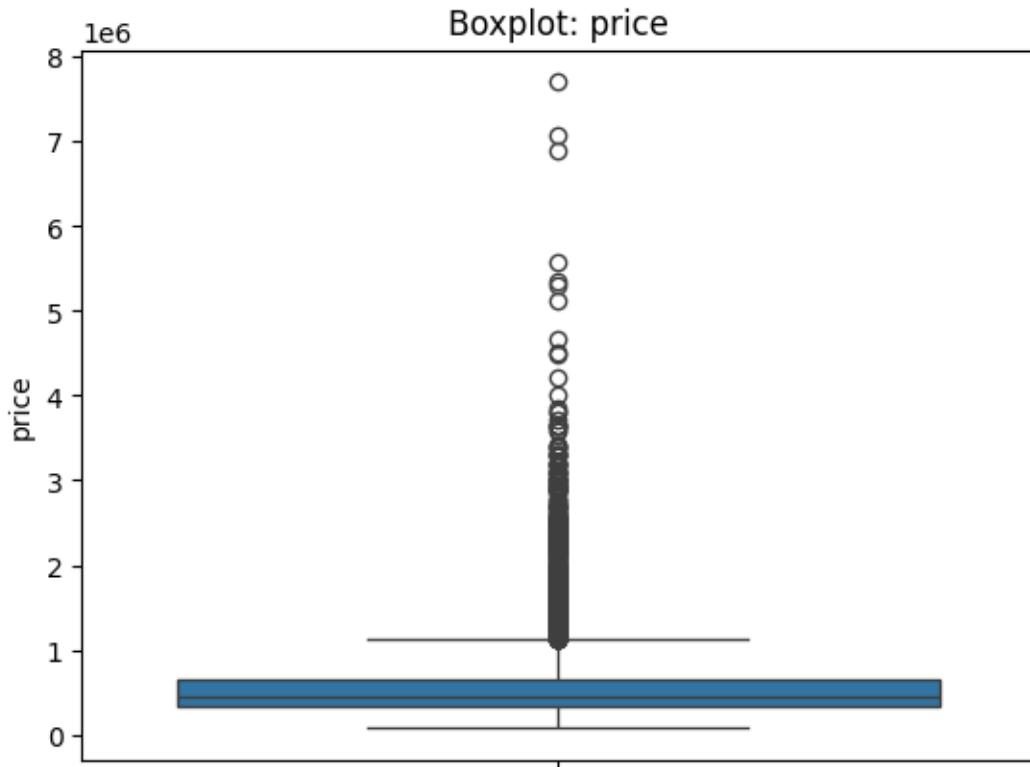
```
[3]: sns.histplot(df_clean.price)
plt.title("Histogram: price")
```

```
[3]: Text(0.5, 1.0, 'Histogram: price')
```



```
[4]: sns.boxplot(df_clean.price)
plt.title("Boxplot: price")
```

```
[4]: Text(0.5, 1.0, 'Boxplot: price')
```



```
[5]: df_clean["price"].describe()
```

```
[5]: count    2.161300e+04
      mean     5.400881e+05
      std      3.671272e+05
      min      7.500000e+04
      25%     3.219500e+05
      50%     4.500000e+05
      75%     6.450000e+05
      max      7.700000e+06
      Name: price, dtype: float64
```

```
[6]: # 11 houses over 4 million
      df_clean[df_clean.price > 4e6]
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
1164	5110800.0	5	5.25	8010	45517	2.0	
1315	5300000.0	6	6.00	7390	24829	2.0	
1448	5350000.0	5	5.00	8000	23985	2.0	
2626	4500000.0	5	5.50	6640	40014	2.0	
3914	7062500.0	5	4.50	10040	37325	2.0	
4411	5570000.0	5	5.75	9200	35069	2.0	

7252	7700000.0	6	8.00	12050	27600	2.5		
8092	4668000.0	5	6.75	9640	13068	1.0		
8638	4489000.0	4	3.00	6430	27517	2.0		
9254	6885000.0	6	7.75	9890	31374	2.0		
12370	4208000.0	5	6.00	7440	21540	2.0		
	waterfront	view	condition	grade	...	yr_built	yr_renovated	\
1164	1	4	3	12	...	1999	0	
1315	1	4	4	12	...	1991	0	
1448	0	4	3	12	...	2009	0	
2626	1	4	3	12	...	2004	0	
3914	1	2	3	11	...	1940	2001	
4411	0	0	3	13	...	2001	0	
7252	0	3	4	13	...	1910	1987	
8092	1	4	3	12	...	1983	2009	
8638	0	0	3	12	...	2001	0	
9254	0	4	3	13	...	2001	0	
12370	0	0	3	12	...	2003	0	
	zipcode	lat	long	sqft_living15	sqft_lot15	year_sold	\	
1164	98033	47.6767	-122.211	3430	26788	2014		
1315	98040	47.5631	-122.210	4320	24619	2015		
1448	98004	47.6232	-122.220	4600	21750	2015		
2626	98155	47.7493	-122.280	3030	23408	2014		
3914	98004	47.6500	-122.214	3930	25449	2014		
4411	98039	47.6289	-122.233	3560	24345	2014		
7252	98102	47.6298	-122.323	3940	8800	2014		
8092	98040	47.5570	-122.210	3270	10454	2014		
8638	98004	47.6208	-122.219	3720	14592	2014		
9254	98039	47.6305	-122.240	4540	42730	2014		
12370	98006	47.5692	-122.189	4740	19329	2015		
	month_sold	day_sold						
1164	10	20						
1315	4	13						
1448	4	13						
2626	8	15						
3914	6	11						
4411	8	4						
7252	10	13						
8092	6	17						
8638	6	18						
9254	9	19						
12370	5	6						

[11 rows x 22 columns]

1.4 How much of the data is represented by outliers?

```
[7]: df_sorted = df_clean["price"].sort_values()

# Cumulative Quatiles
fig = px.ecdf(df_sorted, x="price", title="Cumulative Distribution of House Prices")

# Add a marker line at the 95th and 99th percentiles
fig.add_hline(y=0.95, line_dash="dot", annotation_text="95% of data", annotation_position="bottom right")
fig.show()
```

```
[8]: # quantiles
quantile_99 = df_clean.price.quantile(0.99)
quantile_95 = df_clean.price.quantile(0.95)

print("Quantile 99: ", quantile_99)
print("Quantile 95: ", quantile_95)
```

Quantile 99: 1964400.000000051

Quantile 95: 1156479.999999974

```
[9]: # Flag the datapoints inside each quantiles 99 and 95
df_clean["q_99"] = (df_clean.price < quantile_99).astype(int)
df_clean["q_95"] = (df_clean.price < quantile_95).astype(int)

df_clean.head()
```

```
[9]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront \
0    221900.0        3       1.00      1180      5650     1.0          0
1    538000.0        3       2.25      2570      7242     2.0          0
2    180000.0        2       1.00       770     10000     1.0          0
3    604000.0        4       3.00      1960      5000     1.0          0
4    510000.0        3       2.00      1680      8080     1.0          0

      view  condition  grade  ...  zipcode      lat      long  sqft_living15 \
0      0         3      7  ...  98178  47.5112 -122.257        1340
1      0         3      7  ...  98125  47.7210 -122.319        1690
2      0         3      6  ...  98028  47.7379 -122.233        2720
3      0         5      7  ...  98136  47.5208 -122.393        1360
4      0         3      8  ...  98074  47.6168 -122.045        1800

      sqft_lot15  year_sold  month_sold  day_sold  q_99  q_95
0        5650    2014           10        13     1     1
1        7639    2014           12         9     1     1
2        8062    2015            2        25     1     1
3        5000    2014           12         9     1     1
```

```
4          7503      2015          2          18          1          1
```

```
[5 rows x 24 columns]
```

```
[10]: print(f"The 99th quantile exclude {df_clean.shape[0] - df_clean.q_99.sum()} datapoints")
print(f"The 95th quantile exclude {df_clean.shape[0] - df_clean.q_95.sum()} datapoints")
```

```
The 99th quantile exclude 217 datapoints
The 95th quantile exclude 1081 datapoints
```

1.5 Sensitivity to outliers (metrics)

```
[11]: from utils import *

metrics_df = create_metrics_df()
```

1.5.1 Random Forest flagging the outliers

```
[12]: # Split into train and test
seed = 13
# The price is the target variable
y = df_clean["price"]

# All other variables are the features for the baseline model
X = df_clean.drop(["price"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=seed)
```

```
[13]: X_train
```

```
[13]:    bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  view \
1571        4       1.50      2000     6778     1.0          0      0
16330       4       2.50      2630     48706     2.0          0      0
12786       4       2.50      2620     9525     2.5          0      0
12524       3       2.50      1610     6000     2.0          0      0
16179       3       1.00      880      18205     1.0          0      0
...
153         ...      ...      ...
866         3       2.50      3460     6590     2.0          0      0
74          3       1.75      1790     50529     1.0          0      0
14512       2       1.00      820      5040     1.0          0      0
338         3       1.75      1420     8250     1.0          0      0
```

	condition	grade	sqft_above	...	zipcode	lat	long	\
1571	4	7	1170	...	98198	47.3708	-122.311	
16330	3	8	2630	...	98072	47.7750	-122.125	
12786	4	9	2620	...	98040	47.5631	-122.219	
12524	4	7	1610	...	98038	47.3490	-122.036	
16179	4	6	880	...	98178	47.5013	-122.244	
...	
153	3	12	3540	...	98006	47.5620	-122.162	
866	3	7	3460	...	98056	47.4802	-122.188	
74	5	7	1090	...	98042	47.3511	-122.073	
14512	3	7	820	...	98199	47.6498	-122.388	
338	3	7	1420	...	98133	47.7535	-122.354	
	sqft_living15	sqft_lot15	year_sold	month_sold	day_sold	q_99	q_95	
1571	1940	7531	2015	3	23	1	1	
16330	2680	48706	2014	5	21	1	1	
12786	2580	9525	2014	8	5	1	1	
12524	1570	6000	2014	8	26	1	1	
16179	1110	16115	2014	6	24	1	1	
...	
153	3160	9750	2015	4	1	0	0	
866	2490	6312	2015	4	27	1	1	
74	1940	50529	2015	3	16	1	1	
14512	1730	5760	2014	8	20	1	1	
338	1740	8000	2014	8	26	1	1	

[17290 rows x 23 columns]

```
[14]: # most common hyperparameters or the default ones
from sklearn.ensemble import RandomForestRegressor

rf_regressor = RandomForestRegressor(random_state=seed) #default values + ↴random_state = 13
rf_regressor.fit(X_train, y_train)

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_train,
                             y_train,
                             split = "train",
                             comments="Outliers flagging, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_test,
                             y_test,
```

```
        split = "test",
        comments="Outliers flagging, no normalization.")
```

[15]: metrics_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
		RMSE		Comments			
0		40566.2303		Outliers flagging, no normalization.			
1		93599.4998		Outliers flagging, no normalization.			

1.5.2 Random Forest removing top 1% outliers

```
[16]: # Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_99 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_99 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=seed)
```

[17]: X_train

16385	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	\
16385	3	1.50	1290	8175	1.0	0	0	
16816	3	2.00	1830	10873	1.0	0	0	
6038	3	1.75	1840	8086	1.0	0	0	
16785	3	1.75	1540	10545	2.0	0	0	
15744	3	2.00	2090	15790	1.0	0	0	
...	
155	3	1.00	1180	7669	1.0	0	0	
878	3	1.00	1020	8100	1.0	0	0	
75	4	4.00	3430	35102	2.0	0	0	
14650	4	3.00	2530	5625	1.0	0	0	
345	4	1.00	1000	7134	1.0	0	0	
	condition	grade	sqft_above	...	yr_built	yr_renovated	zipcode	\
16385	4	7	820	...	1952	0	98004	
16816	3	8	1830	...	1989	0	98023	
6038	4	8	1840	...	1964	0	98052	
16785	4	6	1540	...	1978	0	98045	
15744	3	9	2090	...	1992	0	98034	

```

...
      ...   ...   ...   ...   ...   ...
155       4     7     1180   ...    1967   ...
878       3     7     1020   ...    1954   ...
75        4    10     2390   ...    1986   ...
14650     3     8     1470   ...    1976   ...
345       3     6     1000   ...    1943   ...
                                           ...
          lat    long  sqft_living15  sqft_lot15  year_sold  month_sold \
16385  47.6296 -122.205           2130        8577    2014         8
16816  47.3066 -122.394           2490        8976    2014         7
6038   47.6700 -122.155           1840        8060    2014         6
16785  47.4451 -121.763           1540       10000    2015         3
15744  47.7296 -122.199           1820        8770    2014        11
...
      ...   ...   ...   ...
155    47.4479 -122.176           1190        7669    2014         7
878    47.3586 -122.314           1020        8100    2014        12
75     47.5822 -121.987           3240       35020    2014        11
14650   47.7094 -122.233           1840        7070    2014         7
345    47.4897 -122.240           1020        7138    2014         7
                                           ...
          day_sold
16385      20
16816      31
6038       23
16785      11
15744      20
...
      ...
155       28
878       19
75        5
14650      11
345       23

```

[17116 rows x 21 columns]

```

[18]: # most common hyperparameters or the default ones
from sklearn.ensemble import RandomForestRegressor

rf_regressor = RandomForestRegressor(random_state=seed)#default values +
random_state = 13
rf_regressor.fit(X_train, y_train)

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_train,
                             y_train,

```

```

        split = "train",
        comments="Removing top 1%, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 1%, no normalization.")

```

[19]: metrics_df

```

[19]:          Model  Split      R2  Adjusted_R2      MAE    MAPE  \
0  RandomForestRegressor  train  0.9879      0.9879  22769.7215  0.0456
1  RandomForestRegressor  test   0.9322      0.9319  59863.8409  0.1203
2  RandomForestRegressor  train  0.9826      0.9826  23185.3777  0.0475
3  RandomForestRegressor  test   0.8739      0.8733  64327.4972  0.1290

      RMSE           Comments
0  40566.2303  Outliers flagging, no normalization.
1  93599.4998  Outliers flagging, no normalization.
2  37273.0054  Removing top 1%, no normalization.
3  104088.2866  Removing top 1%, no normalization.

```

1.5.3 Random Forest removing top 5% outliers

```

[20]: # Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_95 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_95 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=seed)

```

[21]: X_train

```

[21]:    bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  view  \
7523         3       1.00      1200     7810     1.0          0     0
20438        4       2.50      3250     4500     2.0          0     0
8571         4       2.50      2180     3893     2.0          0     0
15677        3       1.75      1990     5600     1.0          0     1
8032         3       1.75      2360     4063     1.0          0     0
...
159          ...       ...       ...       ...       ...       ...

```

907	3	1.75	1410	9315	1.0	0	0	
78	3	1.00	1410	5060	1.0	0	0	
15224	4	2.50	2570	8178	1.0	0	2	
353	2	1.00	990	3120	1.0	0	2	
	condition	grade	sqft_above	...	yr_built	yr_renovated	zipcode	\
7523	4	7	1200	...	1967		0	98038
20438	3	8	3250	...	2008		0	98059
8571	3	8	2180	...	1999		0	98117
15677	3	8	1330	...	1941		0	98199
8032	5	7	1180	...	1940		0	98117
...	
159	4	8	1760	...	1968		0	98059
907	5	7	1410	...	1960		0	98031
78	4	7	910	...	1956		0	98133
15224	3	8	1710	...	1961		0	98118
353	5	7	790	...	1907		0	98103
	lat	long	sqft_living15	sqft_lot15	year_sold	month_sold	\	
7523	47.3631	-122.050	1590	7800	2015		4	
20438	47.4944	-122.150	3030	4598	2014		11	
8571	47.6886	-122.388	1710	4550	2014		12	
15677	47.6500	-122.415	2630	6780	2014		9	
8032	47.6902	-122.382	1660	4063	2014		8	
...	
159	47.4715	-122.118	1730	11180	2014		10	
907	47.3969	-122.198	1630	8250	2014		10	
78	47.7073	-122.340	1130	5693	2014		6	
15224	47.5483	-122.261	2050	7500	2015		3	
353	47.6800	-122.353	1930	3120	2014		12	
	day_sold							
7523		25						
20438		4						
8571		8						
15677		3						
8032		18						
...	...							
159		3						
907		7						
78		9						
15224		9						
353		3						

[16425 rows x 21 columns]

```
[22]: # most common hyperparameters or the default ones
from sklearn.ensemble import RandomForestRegressor

rf_regressor = RandomForestRegressor(random_state=seed)#default values +
↳random_state = 13
rf_regressor.fit(X_train, y_train)

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_train,
                             y_train,
                             split = "train",
                             comments="Removing top 5%, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             rf_regressor,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 5%, no normalization.")
```

[23]: metrics_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458	
5	RandomForestRegressor	test	0.8717	0.8710	53038.5069	0.1212	
		RMSE		Comments			
0	40566.2303			Outliers flagging, no normalization.			
1	93599.4998			Outliers flagging, no normalization.			
2	37273.0054			Removing top 1%, no normalization.			
3	104088.2866			Removing top 1%, no normalization.			
4	29386.2798			Removing top 5%, no normalization.			
5	76308.9748			Removing top 5%, no normalization.			

1.5.4 XGBoost flagging the outliers

```
[24]: # Split into train and test
seed = 13
# The price is the target variable
y = df_clean["price"]
```

```

# All other variables are the features for the baseline model
X = df_clean.drop(["price"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=seed)

```

```
[25]: import xgboost as xgb

xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[25]: XGBRegressor(base_score=None, booster=None, callbacks=None,
       colsample_bylevel=None, colsample_bynode=None,
       colsample_bytree=None, device=None, early_stopping_rounds=None,
       enable_categorical=False, eval_metric=None, feature_types=None,
       feature_weights=None, gamma=None, grow_policy=None,
       importance_type=None, interaction_constraints=None,
       learning_rate=None, max_bin=None, max_cat_threshold=None,
       max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
       max_leaves=None, min_child_weight=None, missing=np.nan,
       monotone_constraints=None, multi_strategy=None, n_estimators=None,
       n_jobs=None, num_parallel_tree=None, ...)
```

```
[26]: metrics_df = add_new_metrics(metrics_df,
                                   xgb_clf,
                                   X_train,
                                   y_train,
                                   split = "train",
                                   comments="Outliers flagging, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             xgb_clf,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Outliers flagging, no normalization.")
```

```
[27]: metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458	

```

5 RandomForestRegressor    test  0.8717      0.8710  53038.5069  0.1212
6          XGBRegressor   train  0.9825      0.9825  35584.0027  0.0818
7          XGBRegressor   test   0.9260      0.9256  59532.0485  0.1184

              RMSE                      Comments
0  40566.2303  Outliers flagging, no normalization.
1  93599.4998  Outliers flagging, no normalization.
2  37273.0054  Removing top 1%, no normalization.
3  104088.2866  Removing top 1%, no normalization.
4  29386.2798  Removing top 5%, no normalization.
5  76308.9748  Removing top 5%, no normalization.
6  48751.1351  Outliers flagging, no normalization.
7  97798.9764  Outliers flagging, no normalization.

```

1.5.5 XGBoost removing top 1%

```
[28]: # Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_99 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_99 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=seed)
```

```
[29]: X_train
```

```
[29]:      bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  view \
16385        3       1.50      1290      8175     1.0        0        0
16816        3       2.00      1830     10873     1.0        0        0
6038         3       1.75      1840      8086     1.0        0        0
16785        3       1.75      1540     10545     2.0        0        0
15744        3       2.00      2090     15790     1.0        0        0
...
155          3       1.00      1180      7669     1.0        0        0
878          3       1.00      1020      8100     1.0        0        0
75           4       4.00      3430     35102     2.0        0        0
14650        4       3.00      2530      5625     1.0        0        0
345          4       1.00      1000      7134     1.0        0        0

      condition  grade  sqft_above  ...  yr_built  yr_renovated  zipcode \
16385        4      7       820  ...     1952            0    98004
16816        3      8       1830  ...     1989            0    98023
6038         4      8       1840  ...     1964            0    98052
16785        4      6       1540  ...     1978            0    98045
```

```

15744          3     9      2090 ...    1992          0   98034
...        ...
155          4     7      1180 ...    1967          0   98058
878          3     7      1020 ...    1954          0   98198
75           4    10      2390 ...    1986          0   98075
14650         3     8      1470 ...    1976          0   98034
345          3     6      1000 ...    1943          0   98178

      lat      long  sqft_living15  sqft_lot15  year_sold  month_sold \
16385  47.6296 -122.205       2130       8577    2014         8
16816  47.3066 -122.394       2490       8976    2014         7
6038   47.6700 -122.155       1840       8060    2014         6
16785  47.4451 -121.763       1540      10000    2015         3
15744  47.7296 -122.199       1820       8770    2014        11
...
155   47.4479 -122.176       1190       7669    2014         7
878   47.3586 -122.314       1020       8100    2014        12
75    47.5822 -121.987       3240      35020    2014        11
14650  47.7094 -122.233       1840       7070    2014         7
345   47.4897 -122.240       1020       7138    2014         7

      day_sold
16385      20
16816      31
6038      23
16785      11
15744      20
...
155      28
878      19
75       5
14650      11
345      23

```

[17116 rows x 21 columns]

[30]: `import xgboost as xgb`

```
xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

[30]: `XGBRegressor(base_score=None, booster=None, callbacks=None,
 colsample_bylevel=None, colsample_bynode=None,
 colsample_bytree=None, device=None, early_stopping_rounds=None,
 enable_categorical=False, eval_metric=None, feature_types=None,
 feature_weights=None, gamma=None, grow_policy=None,
 importance_type=None, interaction_constraints=None,`

```

learning_rate=None, max_bin=None, max_cat_threshold=None,
max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
max_leaves=None, min_child_weight=None, missing=nan,
monotone_constraints=None, multi_strategy=None, n_estimators=None,
n_jobs=None, num_parallel_tree=None, ...)

```

```
[31]: metrics_df = add_new_metrics(metrics_df,
                                 xgb_clf,
                                 X_train,
                                 y_train,
                                 split = "train",
                                 comments="Removing top 1%, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             xgb_clf,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 1%, no normalization.")
```

```
[32]: metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458	
5	RandomForestRegressor	test	0.8717	0.8710	53038.5069	0.1212	
6	XGBRegressor	train	0.9825	0.9825	35584.0027	0.0818	
7	XGBRegressor	test	0.9260	0.9256	59532.0485	0.1184	
8	XGBRegressor	train	0.9679	0.9679	36410.9174	0.0830	
9	XGBRegressor	test	0.8891	0.8885	61265.8381	0.1234	
					RMSE	Comments	
0	40566.2303				Outliers flagging, no normalization.		
1	93599.4998				Outliers flagging, no normalization.		
2	37273.0054				Removing top 1%, no normalization.		
3	104088.2866				Removing top 1%, no normalization.		
4	29386.2798				Removing top 5%, no normalization.		
5	76308.9748				Removing top 5%, no normalization.		
6	48751.1351				Outliers flagging, no normalization.		
7	97798.9764				Outliers flagging, no normalization.		
8	50584.7096				Removing top 1%, no normalization.		
9	97618.3705				Removing top 1%, no normalization.		

1.5.6 XGBoost removing top 5%

```
[33]: # Split into train and test
# The price is the target variable
y = df_clean[df_clean.q_95 == 1]["price"]

# All other variables are the features for the baseline model
X = df_clean[df_clean.q_95 == 1].drop(["price", "q_99", "q_95"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=seed)
```

```
[34]: import xgboost as xgb

xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[34]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

```
[35]: metrics_df = add_new_metrics(metrics_df,
                                   xgb_clf,
                                   X_train,
                                   y_train,
                                   split = "train",
                                   comments="Removing top 5%, no normalization.")

metrics_df = add_new_metrics(metrics_df,
                             xgb_clf,
                             X_test,
                             y_test,
                             split = "test",
                             comments="Removing top 5%, no normalization.")
```

```
[36]: metrics_df
```

```
[36]:
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9879	0.9879	22769.7215	0.0456	
1	RandomForestRegressor	test	0.9322	0.9319	59863.8409	0.1203	
2	RandomForestRegressor	train	0.9826	0.9826	23185.3777	0.0475	
3	RandomForestRegressor	test	0.8739	0.8733	64327.4972	0.1290	
4	RandomForestRegressor	train	0.9806	0.9806	19934.8930	0.0458	
5	RandomForestRegressor	test	0.8717	0.8710	53038.5069	0.1212	
6	XGBRegressor	train	0.9825	0.9825	35584.0027	0.0818	
7	XGBRegressor	test	0.9260	0.9256	59532.0485	0.1184	
8	XGBRegressor	train	0.9679	0.9679	36410.9174	0.0830	
9	XGBRegressor	test	0.8891	0.8885	61265.8381	0.1234	
10	XGBRegressor	train	0.9569	0.9568	32103.1211	0.0783	
11	XGBRegressor	test	0.8806	0.8800	51602.5277	0.1189	

	RMSE	Comments
0	40566.2303	Outliers flagging, no normalization.
1	93599.4998	Outliers flagging, no normalization.
2	37273.0054	Removing top 1%, no normalization.
3	104088.2866	Removing top 1%, no normalization.
4	29386.2798	Removing top 5%, no normalization.
5	76308.9748	Removing top 5%, no normalization.
6	48751.1351	Outliers flagging, no normalization.
7	97798.9764	Outliers flagging, no normalization.
8	50584.7096	Removing top 1%, no normalization.
9	97618.3705	Removing top 1%, no normalization.
10	43810.9507	Removing top 5%, no normalization.
11	73599.3380	Removing top 5%, no normalization.

The feature flagging for both 99% and 95% worked better than removing the outliers, which means that the information about the top priced houses is still important to accurately predict the prices. In that sense, we will continue the analysis using the flagging of the columns instead of dropping them.

1.5.7 Export dataset and metrics

```
[41]: filename_metrics = "Metrics/outlier_analysis_metrics.csv"

metrics_df.to_csv(filename_metrics, index = False)
```

```
[38]: df_clean
```

```
[38]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	\
0	221900.0	3	1.00	1180	5650	1.0	
1	538000.0	3	2.25	2570	7242	2.0	
2	180000.0	2	1.00	770	10000	1.0	
3	604000.0	4	3.00	1960	5000	1.0	
4	510000.0	3	2.00	1680	8080	1.0	
...	

21608	360000.0	3	2.50	1530	1131	3.0			
21609	400000.0	4	2.50	2310	5813	2.0			
21610	402101.0	2	0.75	1020	1350	2.0			
21611	400000.0	3	2.50	1600	2388	2.0			
21612	325000.0	2	0.75	1020	1076	2.0			
	waterfront	view	condition	grade	...	zipcode	lat	long	\
0	0	0	3	7	...	98178	47.5112	-122.257	
1	0	0	3	7	...	98125	47.7210	-122.319	
2	0	0	3	6	...	98028	47.7379	-122.233	
3	0	0	5	7	...	98136	47.5208	-122.393	
4	0	0	3	8	...	98074	47.6168	-122.045	
...	
21608	0	0	3	8	...	98103	47.6993	-122.346	
21609	0	0	3	8	...	98146	47.5107	-122.362	
21610	0	0	3	7	...	98144	47.5944	-122.299	
21611	0	0	3	8	...	98027	47.5345	-122.069	
21612	0	0	3	7	...	98144	47.5941	-122.299	
	sqft_living15	sqft_lot15	year_sold	month_sold	day_sold	q_99	q_95		
0	1340	5650	2014	10	13	1	1		
1	1690	7639	2014	12	9	1	1		
2	2720	8062	2015	2	25	1	1		
3	1360	5000	2014	12	9	1	1		
4	1800	7503	2015	2	18	1	1		
...	
21608	1530	1509	2014	5	21	1	1		
21609	1830	7200	2015	2	23	1	1		
21610	1020	2007	2014	6	23	1	1		
21611	1410	1287	2015	1	16	1	1		
21612	1020	1357	2014	10	15	1	1		

[21613 rows x 24 columns]

```
[42]: filename_data = "Data/v1_house_sales.csv"
df_clean.to_csv(filename_data, index = False)
```

[]: