

4_Feature_Selection

February 13, 2026

```
[2]: # Import libraries
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.model_selection import train_test_split
from utils import *
```

```
[3]: # Load king_country_dataset
df = pd.read_csv('Data/cleaned_house_sales.csv')
df.head()
```

```
[3]:      price  bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront \
0  221900.0         3       1.00      1180      5650     1.0          0
1  538000.0         3       2.25      2570      7242     2.0          0
2  180000.0         2       1.00       770     10000     1.0          0
3  604000.0         4       3.00      1960      5000     1.0          0
4  510000.0         3       2.00      1680      8080     1.0          0

      view  condition  grade  ...  yr_built  yr_renovated  zipcode        lat \
0      0           3     7  ...    1955                  0   98178  47.5112
1      0           3     7  ...    1951                 1991  98125  47.7210
2      0           3     6  ...    1933                  0   98028  47.7379
3      0           5     7  ...    1965                  0   98136  47.5208
4      0           3     8  ...    1987                  0   98074  47.6168

      long  sqft_living15  sqft_lot15  year_sold  month_sold  day_sold
0 -122.257        1340        5650    2014        10        13
1 -122.319        1690        7639    2014        12         9
2 -122.233        2720        8062    2015         2        25
3 -122.393        1360        5000    2014        12         9
4 -122.045        1800        7503    2015         2        18

[5 rows x 22 columns]
```

```
[4]: # Load king_country_dataset
base_metrics = pd.read_csv('Metrics/baseline_metrics.csv')
base_metrics = base_metrics.drop([2, 3], axis=0)
base_metrics
```

```
[4]:
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	
1	LinearRegression	test	0.7162	0.7148	125985.6747	
4	RandomForestRegressor	train	0.9820	0.9819	25948.4444	
5	RandomForestRegressor	test	0.8959	0.8954	67269.8770	
6	XGBRegressor	train	0.9780	0.9780	39126.7558	
7	XGBRegressor	test	0.9015	0.9010	65712.7448	

	RMSE	MAPE	Comments
0	202864.5703	0.2561	Baseline model
1	191531.3335	0.2596	Baseline model
4	49547.1197	0.0486	Baseline, no normalization, random_state 13, d...
5	115994.2051	0.1271	Baseline, no normalization, random_state 13, d...
6	54676.7727	0.0872	Baseline, no normalization, default values.
7	112860.4995	0.1246	Baseline, no normalization, default values.

In the OLS we see that the p-values of some columns are high. We would like to test if without those columns the models work better or change. The columns are:

- sqft_lot - floors

The Ridge and Lasso Coefficients were also low in some columns and we would like to analize the same on: - month_sold - day_sold - yr_renovated

0.1 Drop 'sqft_lot'

Defining target and features

```
[ ]: X = df.drop(['price', 'sqft_lot'], axis=1) # Features
y = df['price'] # Target
```

Split the data

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, ↴random_state=13)
```

0.1.1 Linear Regression

```
[ ]: from sklearn.linear_model import LinearRegression
# Create the Linear Regression estimator
lm = LinearRegression()

# Perform the fitting
lm.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = create_metrics_df()
metrics_df = add_new_metrics(
    metrics_df,
    lm,
```

```
X_train,  
y_train,  
"train",  
"Whithout sqft_lot"  
)  
  
metrics_df
```

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    lm,  
    X_test,  
    y_test,  
    "test",  
    "Whithout sqft_lot"  
)  
  
metrics_df
```

0.1.2 Random Forest

```
[ ]: from sklearn.ensemble import RandomForestRegressor  
rf_regressor = RandomForestRegressor(random_state=13)  
  
rf_regressor.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    rf_regressor,  
    X_train,  
    y_train,  
    "train",  
    "Whithout sqft_lot"  
)
```

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    rf_regressor,  
    X_test,  
    y_test,  
    "test",  
    "Whithout sqft_lot"  
)  
metrics_df
```

0.1.3 XGBoost

```
[ ]: import xgboost as xgb  
  
xgb_reg = xgb.XGBRegressor()  
xgb_reg.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    xgb_reg,  
    X_train,  
    y_train,  
    "train",  
    "Whithout sqft_lot"  
)  
metrics_df
```

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    xgb_reg,  
    X_test,  
    y_test,  
    "test",  
    "Whithout sqft_lot"  
)  
metrics_df
```

0.1.4 Conclusion

```
[ ]: base_metrics
```

Without this feature the modules don't change significantly so we decided to drop it

0.2 Drop 'floors'

Defining target and features

```
[ ]: X = df.drop(['price','floors'], axis=1) # Features  
y = df['price'] # Target  
X.head()
```

Split the data

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2,  
    ↪random_state=13)
```

0.2.1 Linear Regression

```
[ ]: from sklearn.linear_model import LinearRegression
      # Create the Linear Regression estimator
lm = LinearRegression()

      # Perform the fitting
lm.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = create_metrics_df()
metrics_df = add_new_metrics(
    metrics_df,
    lm,
    X_train,
    y_train,
    "train",
    "Whithout floors"
)

metrics_df
```

```
[ ]: metrics_df = add_new_metrics(
    metrics_df,
    lm,
    X_test,
    y_test,
    "test",
    "Whithout floors"
)

metrics_df
```

0.2.2 Random Forest

```
[ ]: from sklearn.ensemble import RandomForestRegressor
rf_regressor = RandomForestRegressor(random_state=13)

rf_regressor.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = add_new_metrics(
    metrics_df,
    rf_regressor,
    X_train,
    y_train,
    "train",
```

```
        "Whithout floors"
)
metrics_df
```

```
[ ]: metrics_df = add_new_metrics(
    metrics_df,
    rf_regressor,
    X_test,
    y_test,
    "test",
    "Whithout floors"
)
metrics_df
```

0.2.3 XGBoost

```
[ ]: import xgboost as xgb

xgb_reg = xgb.XGBRegressor()
xgb_reg.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_train,
    y_train,
    "train",
    "Whithout floors"
)
metrics_df
```

```
[ ]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_test,
    y_test,
    "test",
    "Whithout floors"
)
metrics_df
```

0.2.4 Conclusion

```
[ ]: metrics_df.to_csv('Metrics/drop_floors.csv')
```

```
[ ]: base_metrics
```

Without this feature the modules don't change significantly so we decided to drop it

0.3 Drop 'month_sold'

Defining target and features

```
[ ]: X = df.drop(['price','month_sold'], axis=1) # Features  
y = df['price'] # Target  
X.head()
```

Split the data

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2,  
random_state=13)
```

0.3.1 Linear Regression

```
[ ]: from sklearn.linear_model import LinearRegression  
# Create the Linear Regression estimator  
lm = LinearRegression()  
  
# Perform the fitting  
lm.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = create_metrics_df()  
metrics_df = add_new_metrics(  
    metrics_df,  
    lm,  
    X_train,  
    y_train,  
    "train",  
    "Whithout month_sold"  
)  
  
metrics_df
```

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    lm,  
    X_test,  
    y_test,  
    "test",  
    "Whithout month_sold"  
)  
  
metrics_df
```

0.3.2 Random Forest

```
[ ]: from sklearn.ensemble import RandomForestRegressor  
rf_regressor = RandomForestRegressor(random_state=13)  
  
rf_regressor.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    rf_regressor,  
    X_train,  
    y_train,  
    "train",  
    "Whithout month_sold"  
)  
metrics_df
```

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    rf_regressor,  
    X_test,  
    y_test,  
    "test",  
    "Whithout month_sold"  
)  
metrics_df
```

0.3.3 XGBoost

```
[ ]: import xgboost as xgb  
  
xgb_reg = xgb.XGBRegressor()  
xgb_reg.fit(X_train, y_train)
```

Evaluation

```
[ ]: metrics_df = add_new_metrics(  
    metrics_df,  
    xgb_reg,  
    X_train,  
    y_train,  
    "train",  
    "Whithout month_sold"  
)  
metrics_df
```

```
[ ]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_test,
    y_test,
    "test",
    "Whithout month_sold"
)
metrics_df
```

0.3.4 Conclusion

```
[ ]: metrics_df.to_csv('Metrics/drop_month_sold.csv')
```

```
[ ]: base_metrics
```

Without this feature the modules don't change significantly so we decided to drop it

0.4 Drop 'day_sold'

Defining target and features

```
[128]: X = df.drop(['price','day_sold'], axis=1) # Features
y = df['price'] # Target
X.head()
```

```
[128]:    bedrooms  bathrooms  sqft_living  sqft_lot  floors  waterfront  view \
0          3        1.00      1180      5650     1.0          0         0
1          3        2.25      2570      7242     2.0          0         0
2          2        1.00       770     10000     1.0          0         0
3          4        3.00      1960      5000     1.0          0         0
4          3        2.00      1680      8080     1.0          0         0

    condition  grade  sqft_above  sqft_basement  yr_built  yr_renovated \
0            3      7        1180                  0        1955                  0
1            3      7        2170                  400       1951        1991
2            3      6        770                  0        1933                  0
3            5      7        1050                 910        1965                  0
4            3      8        1680                  0        1987                  0

    zipcode      lat      long  sqft_living15  sqft_lot15  year_sold  month_sold
0    98178  47.5112 -122.257        1340      5650      2014        10
1    98125  47.7210 -122.319        1690      7639      2014        12
2    98028  47.7379 -122.233        2720      8062      2015         2
3    98136  47.5208 -122.393        1360      5000      2014        12
4    98074  47.6168 -122.045        1800      7503      2015         2
```

Split the data

```
[129]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random_state=13)
```

0.4.1 Linear Regression

```
[130]: from sklearn.linear_model import LinearRegression
# Create the Linear Regression estimator
lm = LinearRegression()

# Perform the fitting
lm.fit(X_train, y_train)
```

```
[130]: LinearRegression()
```

Evaluation

```
[131]: metrics_df = create_metrics_df()
metrics_df = add_new_metrics(
    metrics_df,
    lm,
    X_train,
    y_train,
    "train",
    "Whithout day_sold"
)

metrics_df
```

g:\Mi unidad\IronHack\Project IronKaggel\utils.py:115: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
updated_df = pd.concat([metrics_df, new_row_df], ignore_index=True)
```

```
[131]:      Model  Split      R2  Adjusted_R2          MAE          RMSE  \
0  LinearRegression  train  0.6977      0.6973  125937.4648  202884.7569

      MAPE          Comments
0  0.256  Whithout day_sold
```

```
[132]: metrics_df = add_new_metrics(
    metrics_df,
    lm,
    X_test,
    y_test,
    "test",
    "Whithout day_sold"
```

```
)  
metrics_df
```

```
[132]:      Model  Split      R2  Adjusted_R2      MAE      RMSE \
0  LinearRegression  train  0.6977      0.6973  125937.4648  202884.7569
1  LinearRegression  test   0.7161      0.7148  125973.0357  191559.9896  
  
      MAPE      Comments
0  0.2560  Whithout day_sold
1  0.2595  Whithout day_sold
```

0.4.2 Random Forest

```
[133]: from sklearn.ensemble import RandomForestRegressor
rf_regressor = RandomForestRegressor(random_state=13)  
  
rf_regressor.fit(X_train, y_train)
```

```
[133]: RandomForestRegressor(random_state=13)
```

Evaluation

```
[134]: metrics_df = add_new_metrics(
    metrics_df,
    rf_regressor,
    X_train,
    y_train,
    "train",
    "Whithout day_sold"
)
metrics_df
```

```
[134]:      Model  Split      R2  Adjusted_R2      MAE      RMSE \
0  LinearRegression  train  0.6977      0.6973  125937.4648
1  LinearRegression  test   0.7161      0.7148  125973.0357
2  RandomForestRegressor  train  0.9824      0.9824  25800.6256  
  
      RMSE      MAPE      Comments
0  202884.7569  0.2560  Whithout day_sold
1  191559.9896  0.2595  Whithout day_sold
2  48907.9896  0.0483  Whithout day_sold
```

```
[135]: metrics_df = add_new_metrics(
    metrics_df,
    rf_regressor,
    X_test,
    y_test,
```

```

        "test",
        "Whithout day_sold"
)
metrics_df
```

[135]:

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6973	125937.4648	
1	LinearRegression	test	0.7161	0.7148	125973.0357	
2	RandomForestRegressor	train	0.9824	0.9824	25800.6256	
3	RandomForestRegressor	test	0.8959	0.8954	67286.1944	

	RMSE	MAPE	Comments
0	202884.7569	0.2560	Whithout day_sold
1	191559.9896	0.2595	Whithout day_sold
2	48907.9896	0.0483	Whithout day_sold
3	115992.8695	0.1269	Whithout day_sold

0.4.3 XGBoost

[136]:

```

import xgboost as xgb

xgb_reg = xgb.XGBRegressor()
xgb_reg.fit(X_train, y_train)
```

[136]:

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=None, device=None, early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None, feature_types=None,
             feature_weights=None, gamma=None, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=None, max_bin=None, max_cat_threshold=None,
             max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
             max_leaves=None, min_child_weight=None, missing=nan,
             monotone_constraints=None, multi_strategy=None, n_estimators=None,
             n_jobs=None, num_parallel_tree=None, ...)
```

Evaluation

[137]:

```

metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_train,
    y_train,
    "train",
    "Whithout day_sold"
)
metrics_df
```

```
[137]:
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6973	125937.4648	
1	LinearRegression	test	0.7161	0.7148	125973.0357	
2	RandomForestRegressor	train	0.9824	0.9824	25800.6256	
3	RandomForestRegressor	test	0.8959	0.8954	67286.1944	
4	XGBRegressor	train	0.9783	0.9782	39015.3232	

	RMSE	MAPE	Comments
0	202884.7569	0.2560	Whithout day_sold
1	191559.9896	0.2595	Whithout day_sold
2	48907.9896	0.0483	Whithout day_sold
3	115992.8695	0.1269	Whithout day_sold
4	54398.0531	0.0871	Whithout day_sold

```
[138]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_test,
    y_test,
    "test",
    "Whithout day_sold"
)
metrics_df
```

```
[138]:
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6973	125937.4648	
1	LinearRegression	test	0.7161	0.7148	125973.0357	
2	RandomForestRegressor	train	0.9824	0.9824	25800.6256	
3	RandomForestRegressor	test	0.8959	0.8954	67286.1944	
4	XGBRegressor	train	0.9783	0.9782	39015.3232	
5	XGBRegressor	test	0.9060	0.9056	64958.8204	

	RMSE	MAPE	Comments
0	202884.7569	0.2560	Whithout day_sold
1	191559.9896	0.2595	Whithout day_sold
2	48907.9896	0.0483	Whithout day_sold
3	115992.8695	0.1269	Whithout day_sold
4	54398.0531	0.0871	Whithout day_sold
5	110209.5275	0.1238	Whithout day_sold

0.4.4 Conclusion

```
[139]: metrics_df.to_csv('Metrics/drop_day_sold.csv')
```

```
[144]: base_metrics
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	
1	LinearRegression	test	0.7162	0.7148	125985.6747	
4	RandomForestRegressor	train	0.9820	0.9819	25948.4444	
5	RandomForestRegressor	test	0.8959	0.8954	67269.8770	
6	XGBRegressor	train	0.9780	0.9780	39126.7558	
7	XGBRegressor	test	0.9015	0.9010	65712.7448	
	RMSE	MAPE			Comments	
0	202864.5703	0.2561			Baseline model	
1	191531.3335	0.2596			Baseline model	
4	49547.1197	0.0486	Baseline, no normalization, random_state 13, d...			
5	115994.2051	0.1271	Baseline, no normalization, random_state 13, d...			
6	54676.7727	0.0872	Baseline, no normalization, default values.			
7	112860.4995	0.1246	Baseline, no normalization, default values.			

0.5 Drop 'yr_renovated'

Defining target and features

```
[145]: X = df.drop(['price', 'yr_renovated'], axis=1) # Features
y = df['price'] # Target
X.head()
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	\
0	3	1.00	1180	5650	1.0	0	0	
1	3	2.25	2570	7242	2.0	0	0	
2	2	1.00	770	10000	1.0	0	0	
3	4	3.00	1960	5000	1.0	0	0	
4	3	2.00	1680	8080	1.0	0	0	
	condition	grade	sqft_above	sqft_basement	yr_built	zipcode	lat	\
0	3	7	1180	0	1955	98178	47.5112	
1	3	7	2170	400	1951	98125	47.7210	
2	3	6	770	0	1933	98028	47.7379	
3	5	7	1050	910	1965	98136	47.5208	
4	3	8	1680	0	1987	98074	47.6168	
	long	sqft_living15	sqft_lot15	year_sold	month_sold	day_sold		
0	-122.257	1340	5650	2014	10	13		
1	-122.319	1690	7639	2014	12	9		
2	-122.233	2720	8062	2015	2	25		
3	-122.393	1360	5000	2014	12	9		
4	-122.045	1800	7503	2015	2	18		

Split the data

```
[146]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random_state=13)
```

0.5.1 Linear Regression

```
[147]: from sklearn.linear_model import LinearRegression
# Create the Linear Regression estimator
lm = LinearRegression()

# Perform the fitting
lm.fit(X_train, y_train)
```

```
[147]: LinearRegression()
```

Evaluation

```
[148]: metrics_df = create_metrics_df()
metrics_df = add_new_metrics(
    metrics_df,
    lm,
    X_train,
    y_train,
    "train",
    "Whithout yr_renovated"
)

metrics_df
```

g:\Mi unidad\IronHack\Project IronKaggel\utils.py:115: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
updated_df = pd.concat([metrics_df, new_row_df], ignore_index=True)
```

```
[148]:      Model  Split      R2  Adjusted_R2        MAE        RMSE  \
0  LinearRegression  train  0.6973      0.6969  126097.584  203019.8769

      MAPE           Comments
0  0.2563  Whithout yr_renovated
```

```
[149]: metrics_df = add_new_metrics(
    metrics_df,
    lm,
    X_test,
    y_test,
    "test",
    "Whithout yr_renovated"
```

```
)  
metrics_df
```

```
[149]:      Model  Split      R2  Adjusted_R2      MAE      RMSE  \  
0  LinearRegression  train  0.6973      0.6969  126097.584  203019.8769  
1  LinearRegression  test   0.7158      0.7144  126123.658  191673.0385  
  
      MAPE          Comments  
0  0.2563  Whithout yr_renovated  
1  0.2600  Whithout yr_renovated
```

0.5.2 Random Forest

```
[150]: from sklearn.ensemble import RandomForestRegressor  
rf_regressor = RandomForestRegressor(random_state=13)  
  
rf_regressor.fit(X_train, y_train)
```

```
[150]: RandomForestRegressor(random_state=13)
```

Evaluation

```
[151]: metrics_df = add_new_metrics(  
        metrics_df,  
        rf_regressor,  
        X_train,  
        y_train,  
        "train",  
        "Whithout yr_renovated"  
)  
metrics_df
```

```
[151]:      Model  Split      R2  Adjusted_R2      MAE      RMSE  \  
0  LinearRegression  train  0.6973      0.6969  126097.584  203019.8769  
1  LinearRegression  test   0.7158      0.7144  126123.658  191673.0385  
2  RandomForestRegressor  train  0.9819      0.9819  25962.452   49633.6552  
  
      MAPE          Comments  
0  0.2563  Whithout yr_renovated  
1  0.2600  Whithout yr_renovated  
2  0.0487  Whithout yr_renovated
```

```
[152]: metrics_df = add_new_metrics(  
        metrics_df,  
        rf_regressor,  
        X_test,  
        y_test,
```

```

        "test",
        "Whithout yr_renovated"
)
metrics_df

```

[152]:

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6973	0.6969	126097.5840	
1	LinearRegression	test	0.7158	0.7144	126123.6580	
2	RandomForestRegressor	train	0.9819	0.9819	25962.4520	
3	RandomForestRegressor	test	0.8958	0.8953	67255.3388	

	RMSE	MAPE	Comments
0	203019.8769	0.2563	Whithout yr_renovated
1	191673.0385	0.2600	Whithout yr_renovated
2	49633.6552	0.0487	Whithout yr_renovated
3	116066.5787	0.1272	Whithout yr_renovated

0.5.3 XGBoost

[153]:

```

import xgboost as xgb

xgb_reg = xgb.XGBRegressor()
xgb_reg.fit(X_train, y_train)

```

[153]:

```

XGBRegressor(base_score=None, booster=None, callbacks=None,
            colsample_bylevel=None, colsample_bynode=None,
            colsample_bytree=None, device=None, early_stopping_rounds=None,
            enable_categorical=False, eval_metric=None, feature_types=None,
            feature_weights=None, gamma=None, grow_policy=None,
            importance_type=None, interaction_constraints=None,
            learning_rate=None, max_bin=None, max_cat_threshold=None,
            max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
            max_leaves=None, min_child_weight=None, missing=nan,
            monotone_constraints=None, multi_strategy=None, n_estimators=None,
            n_jobs=None, num_parallel_tree=None, ...)

```

Evaluation

[154]:

```

metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_train,
    y_train,
    "train",
    "Whithout yr_renovated"
)
metrics_df

```

```
[154]:
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6973	0.6969	126097.5840	
1	LinearRegression	test	0.7158	0.7144	126123.6580	
2	RandomForestRegressor	train	0.9819	0.9819	25962.4520	
3	RandomForestRegressor	test	0.8958	0.8953	67255.3388	
4	XGBRegressor	train	0.9774	0.9774	39451.3132	

	RMSE	MAPE	Comments
0	203019.8769	0.2563	Whithout yr_renovated
1	191673.0385	0.2600	Whithout yr_renovated
2	49633.6552	0.0487	Whithout yr_renovated
3	116066.5787	0.1272	Whithout yr_renovated
4	55430.4617	0.0869	Whithout yr_renovated

```
[155]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_test,
    y_test,
    "test",
    "Whithout yr_renovated"
)
metrics_df
```

```
[155]:
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6973	0.6969	126097.5840	
1	LinearRegression	test	0.7158	0.7144	126123.6580	
2	RandomForestRegressor	train	0.9819	0.9819	25962.4520	
3	RandomForestRegressor	test	0.8958	0.8953	67255.3388	
4	XGBRegressor	train	0.9774	0.9774	39451.3132	
5	XGBRegressor	test	0.8991	0.8986	66615.4646	

	RMSE	MAPE	Comments
0	203019.8769	0.2563	Whithout yr_renovated
1	191673.0385	0.2600	Whithout yr_renovated
2	49633.6552	0.0487	Whithout yr_renovated
3	116066.5787	0.1272	Whithout yr_renovated
4	55430.4617	0.0869	Whithout yr_renovated
5	114219.5770	0.1262	Whithout yr_renovated

0.5.4 Conclusion

```
[ ]: metrics_df.to_csv('Metrics/drop_day_sold.csv')

[ ]: base_metrics
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	

```

1      LinearRegression    test  0.7162      0.7148  125985.6747
4  RandomForestRegressor train  0.9820      0.9819  25948.4444
5  RandomForestRegressor  test  0.8959      0.8954  67269.8770
6      XGBRegressor       train  0.9780      0.9780  39126.7558
7      XGBRegressor       test   0.9015      0.9010  65712.7448

          RMSE      MAPE                               Comments
0  202864.5703  0.2561                           Baseline model
1  191531.3335  0.2596                           Baseline model
4  49547.1197  0.0486 Baseline, no normalization, random_state 13, d...
5  115994.2051  0.1271 Baseline, no normalization, random_state 13, d...
6  54676.7727  0.0872 Baseline, no normalization, default values.
7  112860.4995  0.1246 Baseline, no normalization, default values.

```

It makes it slightly worse

0.6 Drop ‘bedrooms’

Defining target and features

```
[5]: X = df.drop(['price', 'bedrooms'], axis=1) # Features
y = df['price'] # Target
X.head()
```

```
[5]:    bathrooms  sqft_living  sqft_lot  floors  waterfront  view  condition \
0        1.00     1180      5650      1.0          0      0      3
1        2.25     2570      7242      2.0          0      0      3
2        1.00      770     10000      1.0          0      0      3
3        3.00     1960      5000      1.0          0      0      5
4        2.00     1680      8080      1.0          0      0      3

    grade  sqft_above  sqft_basement  yr_built  yr_renovated  zipcode      lat \
0      7        1180                  0     1955                  0    98178  47.5112
1      7        2170                  400    1951                1991    98125  47.7210
2      6        770                  0     1933                  0    98028  47.7379
3      7        1050                 910    1965                  0    98136  47.5208
4      8        1680                  0     1987                  0    98074  47.6168

    long  sqft_living15  sqft_lot15  year_sold  month_sold  day_sold
0 -122.257        1340      5650      2014        10       13
1 -122.319        1690      7639      2014        12       9
2 -122.233        2720      8062      2015         2      25
3 -122.393        1360      5000      2014        12       9
4 -122.045        1800      7503      2015         2      18
```

Split the data

```
[6]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2, random_state=13)
```

0.6.1 Random Forest

```
[8]: from sklearn.ensemble import RandomForestRegressor  
rf_regressor = RandomForestRegressor(random_state=13)  
  
rf_regressor.fit(X_train, y_train)
```

```
[8]: RandomForestRegressor(random_state=13)
```

Evaluation

```
[9]: metrics_df = create_metrics_df()  
metrics_df = add_new_metrics(  
    metrics_df,  
    rf_regressor,  
    X_train,  
    y_train,  
    "train",  
    "Whithout bedrooms"  
)  
metrics_df
```

g:\Mi unidad\IronHack\Project IronKaggel\utils.py:115: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
updated_df = pd.concat([metrics_df, new_row_df], ignore_index=True)
```

```
[9]:          Model  Split      R2  Adjusted_R2        MAE        RMSE  \  
0  RandomForestRegressor  train  0.9823      0.9822  25888.0036  49132.8562  
  
      MAPE           Comments  
0  0.0486  Whithout bedrooms
```

```
[10]: metrics_df = add_new_metrics(  
    metrics_df,  
    rf_regressor,  
    X_test,  
    y_test,  
    "test",  
    "Whithout yr_renovated"  
)  
metrics_df
```

```
[10]:          Model  Split      R2  Adjusted_R2        MAE        RMSE  \  
0  RandomForestRegressor  train  0.9823      0.9822  25888.0036  49132.8562  
1  RandomForestRegressor  test   0.8969      0.8964  67042.6859  115463.7677
```

	MAPE	Comments
0	0.0486	Whithout bedrooms
1	0.1270	Whithout yr_renovated

0.6.2 XGBoost

```
[11]: import xgboost as xgb

xgb_reg = xgb.XGBRegressor()
xgb_reg.fit(X_train, y_train)
```

```
[11]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

Evaluation

```
[12]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_train,
    y_train,
    "train",
    "Whithout yr_renovated"
)
metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	RMSE	\
0	RandomForestRegressor	train	0.9823	0.9822	25888.0036	49132.8562	
1	RandomForestRegressor	test	0.8969	0.8964	67042.6859	115463.7677	
2	XGBRegressor	train	0.9786	0.9786	38678.5863	54008.8395	

	MAPE	Comments
0	0.0486	Whithout bedrooms
1	0.1270	Whithout yr_renovated
2	0.0860	Whithout yr_renovated

```
[13]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
```

```

        X_test,
        y_test,
        "test",
        "Whithout yr_renovated"
    )
metrics_df

```

[13]:

	Model	Split	R2	Adjusted_R2	MAE	RMSE	\
0	RandomForestRegressor	train	0.9823	0.9822	25888.0036	49132.8562	
1	RandomForestRegressor	test	0.8969	0.8964	67042.6859	115463.7677	
2	XGBRegressor	train	0.9786	0.9786	38678.5863	54008.8395	
3	XGBRegressor	test	0.9003	0.8999	65793.1286	113502.9973	

	MAPE	Comments
0	0.0486	Whithout bedrooms
1	0.1270	Whithout yr_renovated
2	0.0860	Whithout yr_renovated
3	0.1241	Whithout yr_renovated

0.6.3 Conclusion

[156]: metrics_df.to_csv('Metrics/drop_yr_renovated.csv')

[]:

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	
1	LinearRegression	test	0.7162	0.7148	125985.6747	
4	RandomForestRegressor	train	0.9820	0.9819	25948.4444	
5	RandomForestRegressor	test	0.8959	0.8954	67269.8770	
6	XGBRegressor	train	0.9780	0.9780	39126.7558	
7	XGBRegressor	test	0.9015	0.9010	65712.7448	

	RMSE	MAPE	Comments
0	202864.5703	0.2561	Baseline model
1	191531.3335	0.2596	Baseline model
4	49547.1197	0.0486	Baseline, no normalization, random_state 13, d...
5	115994.2051	0.1271	Baseline, no normalization, random_state 13, d...
6	54676.7727	0.0872	Baseline, no normalization, default values.
7	112860.4995	0.1246	Baseline, no normalization, default values.

It improves a lot

0.7 Drop ‘sqft_lot’, ‘floors’, ‘month_sold’ and ‘day_sold’ toguether

Defining target and features

```
[157]: X = df.drop(['price', 'sqft_lot', 'floors', 'month_sold', 'day_sold'], axis=1) #  
      ↪Features  
y = df['price'] # Target  
X.head()
```

```
[157]:    bedrooms  bathrooms  sqft_living  waterfront  view  condition  grade  \  
0            3        1.00       1180          0       0         3         7  
1            3        2.25       2570          0       0         3         7  
2            2        1.00       770           0       0         3         6  
3            4        3.00      1960           0       0         5         7  
4            3        2.00      1680           0       0         3         8  
  
    sqft_above  sqft_basement  yr_built  yr_renovated  zipcode  lat  \  
0       1180              0       1955             0     98178  47.5112  
1       2170              400      1951            1991     98125  47.7210  
2        770              0       1933             0     98028  47.7379  
3      1050              910      1965             0     98136  47.5208  
4       1680              0       1987             0     98074  47.6168  
  
    long  sqft_living15  sqft_lot15  year_sold  
0 -122.257        1340       5650    2014  
1 -122.319        1690       7639    2014  
2 -122.233        2720       8062    2015  
3 -122.393        1360       5000    2014  
4 -122.045        1800       7503    2015
```

Split the data

```
[158]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.2,  
      ↪random_state=13)
```

0.7.1 Linear Regression

```
[159]: from sklearn.linear_model import LinearRegression  
      # Create the Linear Regression estimator  
lm = LinearRegression()  
  
      # Perform the fitting  
lm.fit(X_train, y_train)
```

```
[159]: LinearRegression()
```

Evaluation

```
[160]: metrics_df = create_metrics_df()  
metrics_df = add_new_metrics(  
    metrics_df,  
    lm,
```

```

        X_train,
        y_train,
        "train",
        "Whithout 4 columns"
    )

metrics_df

```

g:\Mi unidad\IronHack\Project IronKaggel\utils.py:115: FutureWarning: The behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In a future version, this will no longer exclude empty or all-NA columns when determining the result dtypes. To retain the old behavior, exclude the relevant entries before the concat operation.

```
updated_df = pd.concat([metrics_df, new_row_df], ignore_index=True)
```

[160]:

	Model	Split	R2	Adjusted_R2	MAE	RMSE	\
0	LinearRegression	train	0.6976	0.6973	126045.8997	202924.743	
			MAPE	Comments			
0			0.2561	Whithout 4 columns			

[161]:

```

metrics_df = add_new_metrics(
    metrics_df,
    lm,
    X_test,
    y_test,
    "test",
    "Whithout 4 columns"
)

metrics_df

```

[161]:

	Model	Split	R2	Adjusted_R2	MAE	RMSE	\
0	LinearRegression	train	0.6976	0.6973	126045.8997	202924.7430	
1	LinearRegression	test	0.7157	0.7145	126103.8682	191706.7474	
			MAPE	Comments			
0			0.2561	Whithout 4 columns			
1			0.2597	Whithout 4 columns			

0.7.2 Random Forest

[162]:

```

from sklearn.ensemble import RandomForestRegressor
rf_regressor = RandomForestRegressor(random_state=13)

rf_regressor.fit(X_train, y_train)

```

[162]: RandomForestRegressor(random_state=13)

Evaluation

```
[163]: metrics_df = add_new_metrics(  
        metrics_df,  
        rf_regressor,  
        X_train,  
        y_train,  
        "train",  
        "Whithout 4 columns"  
)  
metrics_df
```

```
[163]:  
          Model  Split      R2  Adjusted_R2      MAE  \  
0  LinearRegression  train  0.6976      0.6973  126045.8997  
1  LinearRegression  test   0.7157      0.7145  126103.8682  
2 RandomForestRegressor  train  0.9823      0.9823  25887.6388  
  
      RMSE      MAPE      Comments  
0  202924.7430  0.2561  Whithout 4 columns  
1  191706.7474  0.2597  Whithout 4 columns  
2   49105.0748  0.0486  Whithout 4 columns
```

```
[164]: metrics_df = add_new_metrics(  
        metrics_df,  
        rf_regressor,  
        X_test,  
        y_test,  
        "test",  
        "Whithout 4 columns"  
)  
metrics_df
```

```
[164]:  
          Model  Split      R2  Adjusted_R2      MAE  \  
0  LinearRegression  train  0.6976      0.6973  126045.8997  
1  LinearRegression  test   0.7157      0.7145  126103.8682  
2 RandomForestRegressor  train  0.9823      0.9823  25887.6388  
3 RandomForestRegressor  test   0.8975      0.8971  67388.4881  
  
      RMSE      MAPE      Comments  
0  202924.7430  0.2561  Whithout 4 columns  
1  191706.7474  0.2597  Whithout 4 columns  
2   49105.0748  0.0486  Whithout 4 columns  
3  115083.0470  0.1276  Whithout 4 columns
```

0.7.3 XGBoost

```
[165]: import xgboost as xgb

xgb_reg = xgb.XGBRegressor()
xgb_reg.fit(X_train, y_train)
```

```
[165]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

Evaluation

```
[166]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
    X_train,
    y_train,
    "train",
    "Whithout 4 columns"
)
metrics_df
```

```
[166]:
```

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6976	0.6973	126045.8997	
1	LinearRegression	test	0.7157	0.7145	126103.8682	
2	RandomForestRegressor	train	0.9823	0.9823	25887.6388	
3	RandomForestRegressor	test	0.8975	0.8971	67388.4881	
4	XGBRegressor	train	0.9756	0.9756	40967.3963	

	RMSE	MAPE	Comments
0	202924.7430	0.2561	Whithout 4 columns
1	191706.7474	0.2597	Whithout 4 columns
2	49105.0748	0.0486	Whithout 4 columns
3	115083.0470	0.1276	Whithout 4 columns
4	57584.1626	0.0907	Whithout 4 columns

```
[167]: metrics_df = add_new_metrics(
    metrics_df,
    xgb_reg,
```

```

        X_test,
        y_test,
        "test",
        "Whithout 4 columns"
)
metrics_df

```

[167]:

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6976	0.6973	126045.8997	
1	LinearRegression	test	0.7157	0.7145	126103.8682	
2	RandomForestRegressor	train	0.9823	0.9823	25887.6388	
3	RandomForestRegressor	test	0.8975	0.8971	67388.4881	
4	XGBRegressor	train	0.9756	0.9756	40967.3963	
5	XGBRegressor	test	0.9008	0.9004	67442.8650	

	RMSE	MAPE	Comments
0	202924.7430	0.2561	Whithout 4 columns
1	191706.7474	0.2597	Whithout 4 columns
2	49105.0748	0.0486	Whithout 4 columns
3	115083.0470	0.1276	Whithout 4 columns
4	57584.1626	0.0907	Whithout 4 columns
5	113232.2624	0.1288	Whithout 4 columns

0.7.4 Conclusion

[168]:

```
metrics_df.to_csv('Metrics/drop_4_columns.csv')
```

[169]:

```
base_metrics
```

[169]:

	Model	Split	R2	Adjusted_R2	MAE	\
0	LinearRegression	train	0.6977	0.6974	125948.1118	
1	LinearRegression	test	0.7162	0.7148	125985.6747	
4	RandomForestRegressor	train	0.9820	0.9819	25948.4444	
5	RandomForestRegressor	test	0.8959	0.8954	67269.8770	
6	XGBRegressor	train	0.9780	0.9780	39126.7558	
7	XGBRegressor	test	0.9015	0.9010	65712.7448	

	RMSE	MAPE	Comments
0	202864.5703	0.2561	Baseline model
1	191531.3335	0.2596	Baseline model
4	49547.1197	0.0486	Baseline, no normalization, random_state 13, d...
5	115994.2051	0.1271	Baseline, no normalization, random_state 13, d...
6	54676.7727	0.0872	Baseline, no normalization, default values.
7	112860.4995	0.1246	Baseline, no normalization, default values.

Dropping this columns improves slightly in general the scores of Linear Regression and XGB but not Random Forest