

# 5\_FeatureEngineering

February 13, 2026

## 1 Ironkaggle

```
[31]: import pandas as pd
import numpy as np
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
import seaborn as sns
import matplotlib.pyplot as plt
import xgboost as xgb
import plotly.express as px
from sklearn.linear_model import LinearRegression, Ridge, Lasso
import statsmodels.api as sm
```

### 1.0.1 Load Data - Show basics

```
[2]: dataset_csv= pd.read_csv("king_ country_ houses_aa.csv")
dataset = dataset_csv.copy()
```

## 2 Functions used for Analysis

```
[3]: ### Functions - Taken from Anne's Utils.py

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.metrics import (r2_score,
                             mean_absolute_error,
                             mean_absolute_percentage_error,
                             mean_squared_error,
```

```

        root_mean_squared_error)

from sklearn.base import BaseEstimator
from typing import SupportsFloat, Union, Dict

# Define types for clearer code
ArrayLike = Union[pd.DataFrame, pd.Series, np.ndarray]

def adjusted_r2(y_true: ArrayLike, y_pred: ArrayLike, X: ArrayLike) -> float:
    """
    Calculate Adjusted R2, handling both Arrays and DataFrames.
    """
    r2 = r2_score(y_true, y_pred)
    n = len(y_true)

    # Handle X shape safely (whether it's a DataFrame or Numpy array)
    if hasattr(X, 'shape'):
        p = X.shape[1] if len(X.shape) > 1 else 1
    else:
        p = 1 # Fallback for 1D lists

    return 1 - (1 - r2) * (n - 1) / (n - p - 1)

def create_metrics_df():
    """
    Create the dataframe to store the metrics
    """

    Returns:
    pd.DataFrame: The empty metrics DataFrame.
    """
    columns = ["Model", "Split", "R2", "Adjusted_R2", "MAE", "RMSE", ↴
    ↵"MAPE", "Comments"]
    metrics_df = pd.DataFrame(columns=columns)
    return metrics_df

def _get_metrics(
    trained_model: BaseEstimator,
    X: ArrayLike,
    y: ArrayLike,
    split: str = "train",
    comments: str = "Baseline model"
) -> Dict[str, Union[str, float]]:
    """
    Internal function to calculate metrics for a single split.

    Args:
    """

```

```

trained_model: A fitted sklearn model.
X: Feature matrix (DataFrame or Numpy Array).
y: Target vector (Series or Numpy Array).
split: Label for the data split (e.g., 'Train', 'Test').
comments: User notes.

>Returns:
dict: A dictionary containing all calculated metrics.
"""

# Generate predictions
y_pred = trained_model.predict(X)

# Calculate metrics
new_row = {
    "Model": trained_model.__class__.__name__,
    "Split": split,
    "R2": np.round(r2_score(y, y_pred), 4),
    "Adjusted_R2": np.round(adjusted_r2(y, y_pred, X), 4),
    "MAE": np.round(mean_absolute_error(y, y_pred), 4),
    "MAPE": np.round(mean_absolute_percentage_error(y, y_pred), 4),
    "RMSE": np.round(root_mean_squared_error(y, y_pred), 4),
    "Comments": comments
}

return new_row

def add_new_metrics(
    metrics_df: pd.DataFrame,
    trained_model: BaseEstimator,
    X: ArrayLike,
    y: ArrayLike,
    split: str = "train",
    comments: str = "Baseline model"
) -> pd.DataFrame:
"""
Calculates metrics and appends them to the tracking DataFrame.

>Args:
metrics_df: The existing DataFrame to update.
trained_model: A fitted sklearn model.
X: Feature matrix.
y: Target vector.
split: "Train" or "Test".
comments: Notes about this run.

>Returns:
pd.DataFrame: The updated DataFrame with the new row.

```

```

"""
# Get the metrics dictionary
new_row_dict = _get_metrics(trained_model, X, y, split, comments)

# Create a DataFrame from the new row
new_row_df = pd.DataFrame([new_row_dict])

# Concatenate and RETURN the result
if metrics_df.empty:
    return new_row_df
else:
    updated_df = pd.concat([metrics_df, new_row_df], ignore_index=True)

return updated_df

def generate_heatmap(X):
    # 1. Calculate correlation
    corr = np.round(np.abs(X.corr()), 2)

    # 2. Create the mask (True for upper triangle and diagonal)
    mask = np.zeros_like(corr, dtype=bool)
    mask[np.triu_indices_from(mask)] = True

    # 3. Slice the matrix and mask to remove the completely empty row/column
    #     Row 0 is fully masked (hidden), Column -1 is fully masked (hidden)
    corr_sliced = corr.iloc[1:, :-1]
    mask_sliced = mask[1:, :-1]

    f, ax = plt.subplots(figsize=(16, 16))

    # 4. Plot using the sliced data
    #     Note: annot=True is safer than annot=corr when slicing
    sns.heatmap(
        corr_sliced,
        mask=mask_sliced,
        annot=True,           # Use True to automatically label values
        square=True,
        linewidths=.5,
        vmax=1
    )
    plt.show()

from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error

```

```

def get_r_squared(y_train, y_test, y_pred_train, y_pred_test):
    """Get the r2 for train and test and print the values.

    Args:
        y_train (_type_): Ground Truth target values for the train.
        y_test (_type_): Ground Truth target values for the test.
        y_pred_train (_type_): Predicted values for the train.
        y_pred_test (_type_): Predicted values for the test.

    Returns:
        tuple: tuple containing the r2 for the train and the test in that order.
    """
    r2_train = r2_score(y_train, y_pred_train)
    r2_test = r2_score(y_test, y_pred_test)

    print_metrics("R2", r2_train, r2_test)
    return (r2_train, r2_test)

def get_mse(y_train, y_test, y_pred_train, y_pred_test):
    """Get the mse for train and test and print the values.

    Args:
        y_train (_type_): Ground Truth target values for the train.
        y_test (_type_): Ground Truth target values for the test.
        y_pred_train (_type_): Predicted values for the train.
        y_pred_test (_type_): Predicted values for the test.

    Returns:
        tuple: tuple containing the mse for the train and the test in that order.
    """
    mse_train = mean_squared_error(y_train, y_pred_train)
    mse_test = mean_squared_error(y_test, y_pred_test)

    print_metrics("MSE", mse_train, mse_test)
    return (mse_train, mse_test)

def get_mae(y_train, y_test, y_pred_train, y_pred_test):
    """Get the mae for train and test and print the values.

    Args:
        y_train (_type_): Ground Truth target values for the train.
        y_test (_type_): Ground Truth target values for the test.
        y_pred_train (_type_): Predicted values for the train.
        y_pred_test (_type_): Predicted values for the test.

    Returns:

```

```

        tuple: tuple containing the mae for the train and the test in that_
        ↵order.

    """
    mae_train = mean_absolute_error(y_train, y_pred_train)
    mae_test = mean_absolute_error(y_test, y_pred_test)

    print_metrics("MSE", mae_train, mae_test)
    return (mae_train, mae_test)

def print_metrics(metric_name:str, train_score:SupportsFloat, test_score:
    ↵SupportsFloat):
    string = f"""
    {metric_name} score:
    train | {train_score}
    test  | {test_score}
    """
    print(string)

```

### 3 Basic Data Inspection

```
[4]: # Basic Data Inspection

# Check Shape of the Dataset
print(f"Dataset Shape : {dataset.shape[0]} rows x {dataset.shape[1]}")

# Check Column Names of the Dataset
print(f"\nColumn Names : {dataset.columns.tolist()}")

# Temp Dataset
dataset_alt = dataset.copy()

# Missing Values Check
print("\nCheck for possible Null Values:")
missingVal = dataset.isnull().sum()
print(missingVal)
if missingVal.sum() == 0:
    print("NO Missing Values found")

# Check for duplicates
print(f"\nChecking for Duplicates :")
duplicatesVal = dataset.duplicated().sum()
if duplicatesVal == 0:
    print(f"No Duplicates found")
```

Dataset Shape : 21613 rows x 21

```
Column Names : ['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',  
'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade', 'sqft_above',  
'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode', 'lat', 'long',  
'sqft_living15', 'sqft_lot15']
```

Check for possible Null Values:

```
id          0  
date        0  
price       0  
bedrooms    0  
bathrooms   0  
sqft_living 0  
sqft_lot    0  
floors      0  
waterfront  0  
view        0  
condition   0  
grade       0  
sqft_above  0  
sqft_basement 0  
yr_built    0  
yr_renovated 0  
zipcode     0  
lat         0  
long        0  
sqft_living15 0  
sqft_lot15  0  
dtype: int64  
NO Missing Values found
```

Checking for Duplicates :

```
No Duplicates found
```

## 4 Feature Engineering

### Proposal

- Convert Date Columnn to Pandas Data-Time series and extract 3 new columns : day,month and year.
- Create 2 new columns called : Year\_Since\_Renovation and Was\_Renovated(boolean)

```
[5]: # Convert Date Column to Day,month and Year  
dataset_alt.date = pd.to_datetime(dataset_alt.date)  
dataset_alt["year_sold"] = dataset_alt.date.dt.year  
dataset_alt["month_sold"] = dataset_alt.date.dt.month  
dataset_alt["day_sold"] = dataset_alt.date.dt.day
```

```

# Selected features for FEATURE ENGINEERING
dataset_alt['yr_built'] = pd.to_numeric(dataset_alt['yr_built'],  

    ↪errors='coerce')
dataset_alt['yr_renovated'] = pd.to_numeric(dataset_alt['yr_renovated'],  

    ↪errors='coerce')
dataset_alt['year_sold'] = pd.to_numeric(dataset_alt['year_sold'],  

    ↪errors='coerce')

# Add Columns : Age at Sale, Year since renovation, Was Renovated(bool) to the  

    ↪Dataset
dataset_alt["Age_at_Sale"] = dataset_alt["year_sold"] - dataset_alt["yr_built"]

dataset_alt["Year_since_Renovation"] = dataset_alt["year_sold"] -  

    ↪dataset_alt["yr_renovated"]

dataset_alt.loc[dataset_alt["yr_renovated"] == 0, "Year_since_Renovation"] =  

    ↪dataset_alt["year_sold"] - dataset_alt["yr_built"]

print(f"\nColumn Names : {dataset_alt.columns.tolist()}")

dataset_alt["was_renovated"] = (dataset_alt["yr_renovated"] > 0) &  

    ↪(dataset_alt["yr_renovated"] > dataset_alt["yr_built"])

print(f"\nData Type of columns :\n{dataset_alt.dtypes}")

Cleaned_Dataset = dataset_alt.drop(columns = ["date", "id", "yr_built",  

    ↪"year_sold", "month_sold" ]) # yr_built, year_sold
Cleaned_Dataset.info()
print(f"\nColumn Names : {Cleaned_Dataset.columns.tolist()}")
Cleaned_Dataset.to_csv("king_country_houses_fe.csv", index=False)

```

Column Names : ['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft\_living',  
 'sqft\_lot', 'floors', 'waterfront', 'view', 'condition', 'grade', 'sqft\_above',  
 'sqft\_basement', 'yr\_built', 'yr\_renovated', 'zipcode', 'lat', 'long',  
 'sqft\_living15', 'sqft\_lot15', 'year\_sold', 'month\_sold', 'day\_sold',  
 'Age\_at\_Sale', 'Year\_since\_Renovation']

Data Type of columns :

id	int64
date	datetime64[ns]
price	float64
bedrooms	int64
bathrooms	float64
sqft_living	int64
sqft_lot	int64
floors	float64

```

waterfront           int64
view                int64
condition           int64
grade               int64
sqft_above          int64
sqft_basement       int64
yr_built            int64
yr_renovated        int64
zipcode             int64
lat                 float64
long                float64
sqft_living15       int64
sqft_lot15          int64
year_sold           int32
month_sold          int32
day_sold            int32
Age_at_Sale         int64
Year_since_Renovation int64
was_renovated       bool
dtype: object
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 22 columns):
 #   Column           Non-Null Count Dtype  
 --- 
 0   price            21613 non-null  float64 
 1   bedrooms          21613 non-null  int64   
 2   bathrooms         21613 non-null  float64 
 3   sqft_living       21613 non-null  int64   
 4   sqft_lot          21613 non-null  int64   
 5   floors            21613 non-null  float64 
 6   waterfront        21613 non-null  int64   
 7   view              21613 non-null  int64   
 8   condition          21613 non-null  int64   
 9   grade              21613 non-null  int64   
 10  sqft_above         21613 non-null  int64   
 11  sqft_basement      21613 non-null  int64   
 12  yr_renovated       21613 non-null  int64   
 13  zipcode            21613 non-null  int64   
 14  lat                21613 non-null  float64 
 15  long               21613 non-null  float64 
 16  sqft_living15      21613 non-null  int64   
 17  sqft_lot15          21613 non-null  int64   
 18  day_sold           21613 non-null  int32   
 19  Age_at_Sale         21613 non-null  int64   
 20  Year_since_Renovation 21613 non-null  int64   
 21  was_renovated       21613 non-null  bool    
dtypes: bool(1), float64(5), int32(1), int64(15)

```

```
memory usage: 3.4 MB
```

```
Column Names : ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot',
'floors', 'waterfront', 'view', 'condition', 'grade', 'sqft_above',
'sqft_basement', 'yr_renovated', 'zipcode', 'lat', 'long', 'sqft_living15',
'sqft_lot15', 'day_sold', 'Age_at_Sale', 'Year_since_Renovation',
'was_renovated']
```

## 5 Split Data for train\_test\_split

- Here we split the Data :
  - X - Holds the Cleaned Dataset without PRICE
  - Y - Holds the Cleaned Dataset with only PRICE

```
[6]: X = Cleaned_Dataset.drop(columns = ["price"])
y = Cleaned_Dataset["price"]
seed = 13

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,random_state=seed)

print(f"X_train Training Samples: {X_train.shape[0]}")
print(f"X_test Samples: {X_test.shape[0]}")
print(f"Y_train Training Samples: {y_train.shape[0]}")
print(f"Y_test Training Samples: {y_test.shape[0]}")

print(f"Split: {len(X_train)} train ({len(X_train)/len(Cleaned_Dataset):.1%}) | {len(X_test)} test ({len(X_test)/len(Cleaned_Dataset):.1%})")
```

```
X_train Training Samples: 17290
X_test Samples: 4323
Y_train Training Samples: 17290
Y_test Training Samples: 4323
Split: 17290 train (80.0%) | 4323 test (20.0%)
```

## 6 Correlation

- Here a simple correlation without controlling any parameters is done.
- And consequently a Heatmap

```
[7]: check_correlation = Cleaned_Dataset.corr()
print("\nCorrelation Analysis :")
print(check_correlation.round(3))

# Plot correlation Heatmap
```

```

plt.figure(figsize=(15,15))
sns.heatmap(check_correlation, annot=True, cmap='coolwarm', center=0, fmt='.
˓→2f', square=True, linewidths=1)
plt.title("Housing Dataset Correlation Heatmap", fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
plt.figure()

```

Correlation Analysis :

	price	bedrooms	bathrooms	sqft_living	sqft_lot	\
price	1.000	0.308	0.525	0.702	0.090	
bedrooms	0.308	1.000	0.516	0.577	0.032	
bathrooms	0.525	0.516	1.000	0.755	0.088	
sqft_living	0.702	0.577	0.755	1.000	0.173	
sqft_lot	0.090	0.032	0.088	0.173	1.000	
floors	0.257	0.175	0.501	0.354	-0.005	
waterfront	0.266	-0.007	0.064	0.104	0.022	
view	0.397	0.080	0.188	0.285	0.075	
condition	0.036	0.028	-0.125	-0.059	-0.009	
grade	0.667	0.357	0.665	0.763	0.114	
sqft_above	0.606	0.478	0.685	0.877	0.184	
sqft_basement	0.324	0.303	0.284	0.435	0.015	
yr_renovated	0.126	0.019	0.051	0.055	0.008	
zipcode	-0.053	-0.153	-0.204	-0.199	-0.130	
lat	0.307	-0.009	0.025	0.053	-0.086	
long	0.022	0.129	0.223	0.240	0.230	
sqft_living15	0.585	0.392	0.569	0.756	0.145	
sqft_lot15	0.082	0.029	0.087	0.183	0.719	
day_sold	-0.015	-0.008	-0.005	-0.007	0.001	
Age_at_Sale	-0.054	-0.154	-0.506	-0.318	-0.053	
Year_since_Renovation	-0.106	-0.166	-0.537	-0.344	-0.053	
was_renovated	0.126	0.019	0.050	0.055	0.008	

	floors	waterfront	view	condition	grade	...	\
price	0.257	0.266	0.397	0.036	0.667	...	
bedrooms	0.175	-0.007	0.080	0.028	0.357	...	
bathrooms	0.501	0.064	0.188	-0.125	0.665	...	
sqft_living	0.354	0.104	0.285	-0.059	0.763	...	
sqft_lot	-0.005	0.022	0.075	-0.009	0.114	...	
floors	1.000	0.024	0.029	-0.264	0.458	...	
waterfront	0.024	1.000	0.402	0.017	0.083	...	
view	0.029	0.402	1.000	0.046	0.251	...	
condition	-0.264	0.017	0.046	1.000	-0.145	...	
grade	0.458	0.083	0.251	-0.145	1.000	...	
sqft_above	0.524	0.072	0.168	-0.158	0.756	...	
sqft_basement	-0.246	0.081	0.277	0.174	0.168	...	

yr_renovated	0.006	0.093	0.104	-0.061	0.014	...
zipcode	-0.059	0.030	0.085	0.003	-0.185	...
lat	0.050	-0.014	0.006	-0.015	0.114	...
long	0.125	-0.042	-0.078	-0.107	0.198	...
sqft_living15	0.280	0.086	0.280	-0.093	0.713	...
sqft_lot15	-0.011	0.031	0.073	-0.003	0.119	...
day_sold	-0.007	0.011	0.011	-0.005	-0.012	...
Age_at_Sale	-0.490	0.026	0.053	0.361	-0.447	...
Year_since_Renovation	-0.506	0.000	0.018	0.396	-0.461	...
was_renovated	0.006	0.093	0.104	-0.060	0.014	...

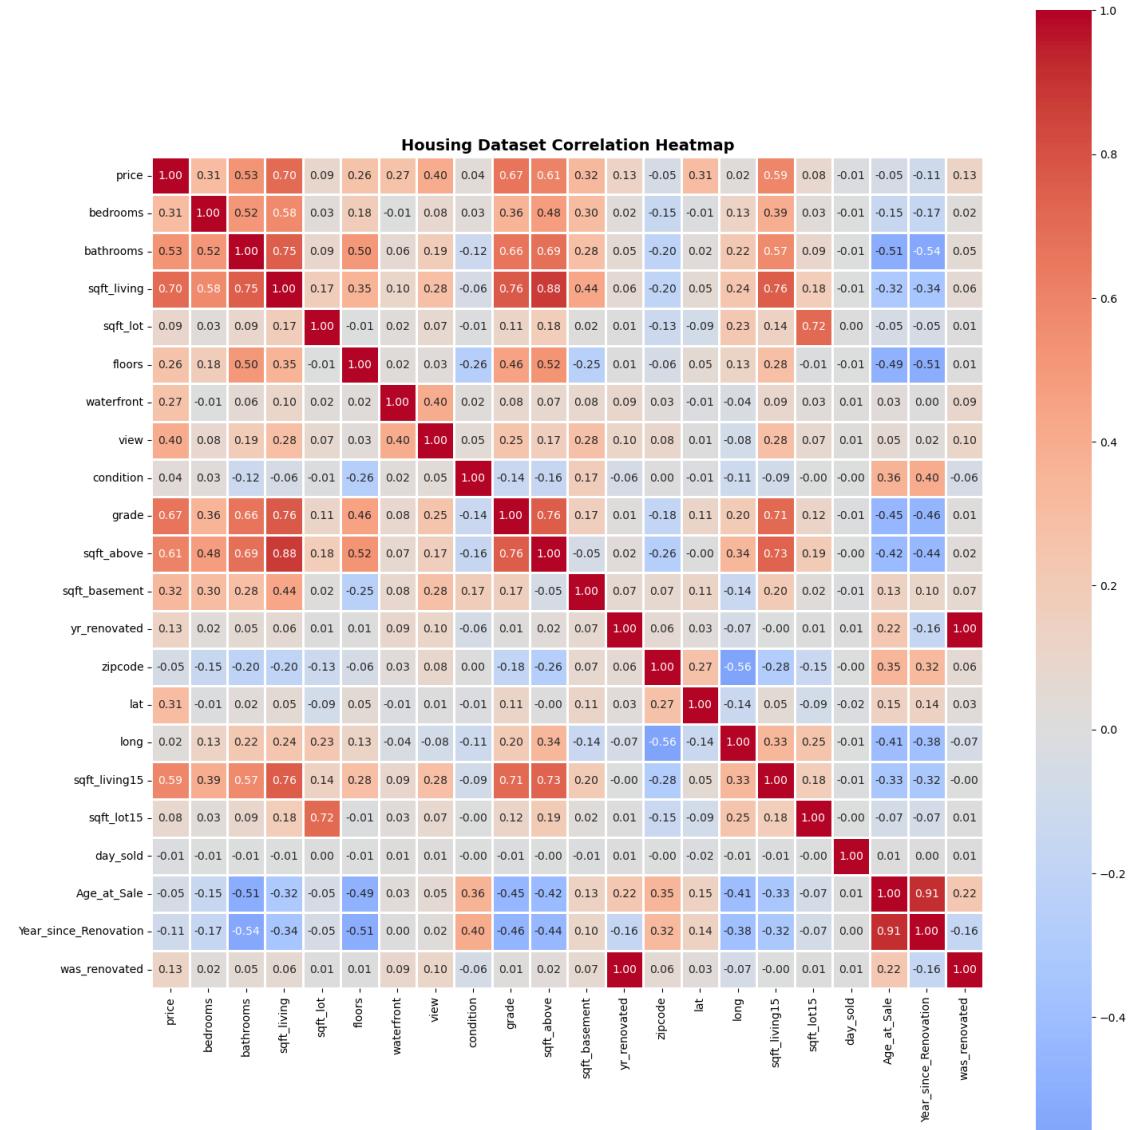
	yr_renovated	zipcode	lat	long	sqft_living15	\
price	0.126	-0.053	0.307	0.022		0.585
bedrooms	0.019	-0.153	-0.009	0.129		0.392
bathrooms	0.051	-0.204	0.025	0.223		0.569
sqft_living	0.055	-0.199	0.053	0.240		0.756
sqft_lot	0.008	-0.130	-0.086	0.230		0.145
floors	0.006	-0.059	0.050	0.125		0.280
waterfront	0.093	0.030	-0.014	-0.042		0.086
view	0.104	0.085	0.006	-0.078		0.280
condition	-0.061	0.003	-0.015	-0.107		-0.093
grade	0.014	-0.185	0.114	0.198		0.713
sqft_above	0.023	-0.261	-0.001	0.344		0.732
sqft_basement	0.071	0.075	0.111	-0.145		0.200
yr_renovated	1.000	0.064	0.029	-0.068		-0.003
zipcode	0.064	1.000	0.267	-0.564		-0.279
lat	0.029	0.267	1.000	-0.136		0.049
long	-0.068	-0.564	-0.136	1.000		0.335
sqft_living15	-0.003	-0.279	0.049	0.335		1.000
sqft_lot15	0.008	-0.147	-0.086	0.254		0.183
day_sold	0.008	-0.003	-0.016	-0.007		-0.009
Age_at_Sale	0.224	0.347	0.148	-0.409		-0.327
Year_since_Renovation	-0.165	0.321	0.135	-0.383		-0.325
was_renovated	1.000	0.064	0.029	-0.068		-0.003

	sqft_lot15	day_sold	Age_at_Sale	\
price	0.082	-0.015	-0.054	
bedrooms	0.029	-0.008	-0.154	
bathrooms	0.087	-0.005	-0.506	
sqft_living	0.183	-0.007	-0.318	
sqft_lot	0.719	0.001	-0.053	
floors	-0.011	-0.007	-0.490	
waterfront	0.031	0.011	0.026	
view	0.073	0.011	0.053	
condition	-0.003	-0.005	0.361	
grade	0.119	-0.012	-0.447	
sqft_above	0.194	-0.002	-0.424	
sqft_basement	0.017	-0.010	0.133	

yr_renovated	0.008	0.008	0.224
zipcode	-0.147	-0.003	0.347
lat	-0.086	-0.016	0.148
long	0.254	-0.007	-0.409
sqft_living15	0.183	-0.009	-0.327
sqft_lot15	1.000	-0.003	-0.071
day_sold	-0.003	1.000	0.006
Age_at_Sale	-0.071	0.006	1.000
Year_since_Renovation	-0.070	0.003	0.910
was_renovated	0.008	0.008	0.225

	Year_since_Renovation	was_renovated
price	-0.106	0.126
bedrooms	-0.166	0.019
bathrooms	-0.537	0.050
sqft_living	-0.344	0.055
sqft_lot	-0.053	0.008
floors	-0.506	0.006
waterfront	0.000	0.093
view	0.018	0.104
condition	0.396	-0.060
grade	-0.461	0.014
sqft_above	-0.436	0.023
sqft_basement	0.102	0.071
yr_renovated	-0.165	1.000
zipcode	0.321	0.064
lat	0.135	0.029
long	-0.383	-0.068
sqft_living15	-0.325	-0.003
sqft_lot15	-0.070	0.008
day_sold	0.003	0.008
Age_at_Sale	0.910	0.225
Year_since_Renovation	1.000	-0.164
was_renovated	-0.164	1.000

[22 rows x 22 columns]

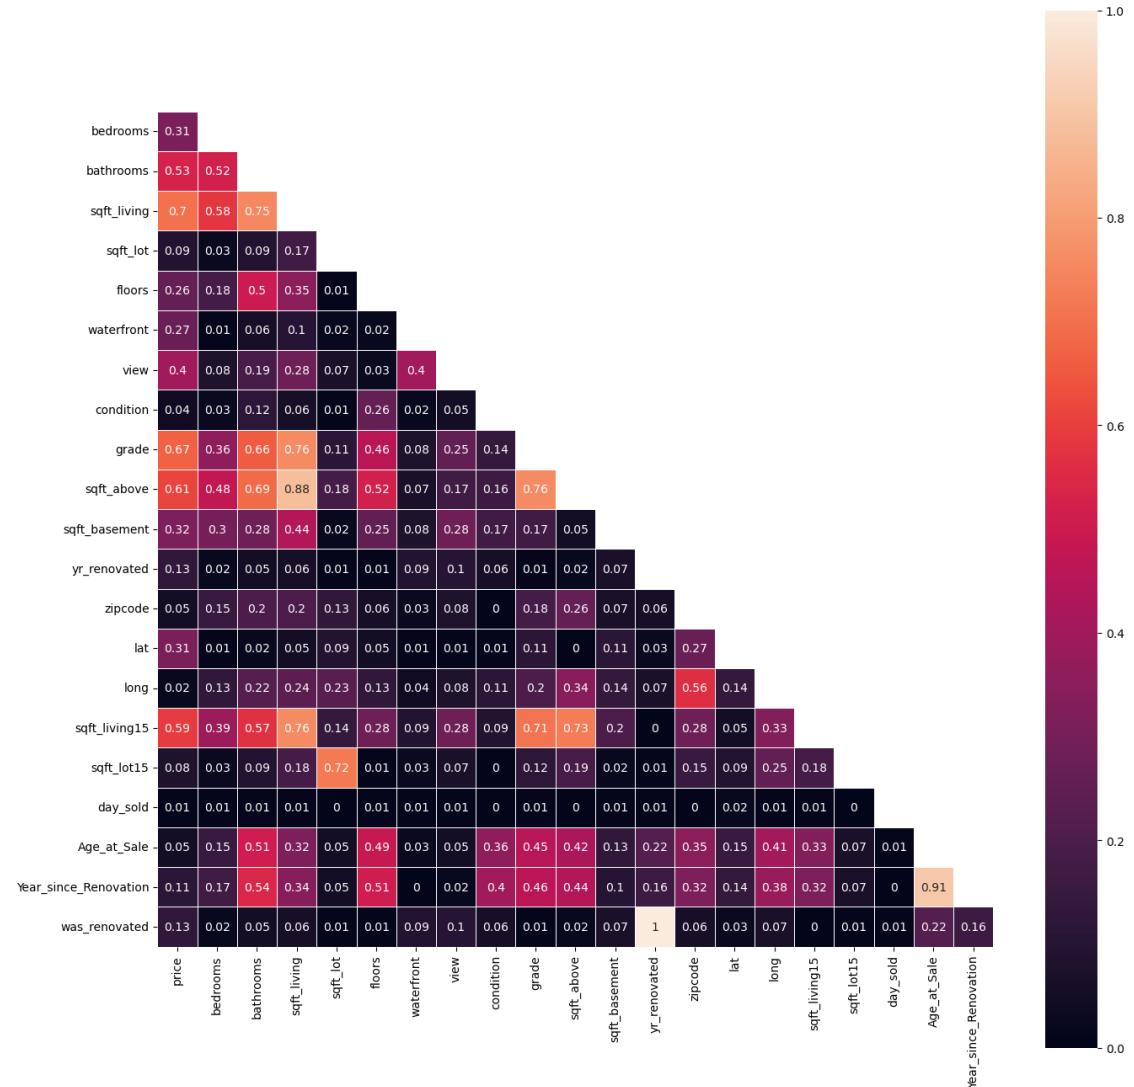


[7]: <Figure size 640x480 with 0 Axes>

<Figure size 640x480 with 0 Axes>

## 7 Parameterized Heatmap

[8]: generate\_heatmap(Cleaned\_Dataset)



## 8 Feature Importance

- Lasso Regularization

```
[9]: std_scaler = StandardScaler()

X_train_standard = std_scaler.fit_transform(X_train)
X_test_standard = std_scaler.fit_transform(X_test)

print(X_train_standard.min(), X_train_standard.max())
```

-3.956684808851382 39.35388067380475

```
[10]: features = X_train.columns  
features
```

```
[10]: Index(['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',  
           'waterfront', 'view', 'condition', 'grade', 'sqft_above',  
           'sqft_basement', 'yr_renovated', 'zipcode', 'lat', 'long',  
           'sqft_living15', 'sqft_lot15', 'day_sold', 'Age_at_Sale',  
           'Year_since_Renovation', 'was_renovated'],  
           dtype='object')
```

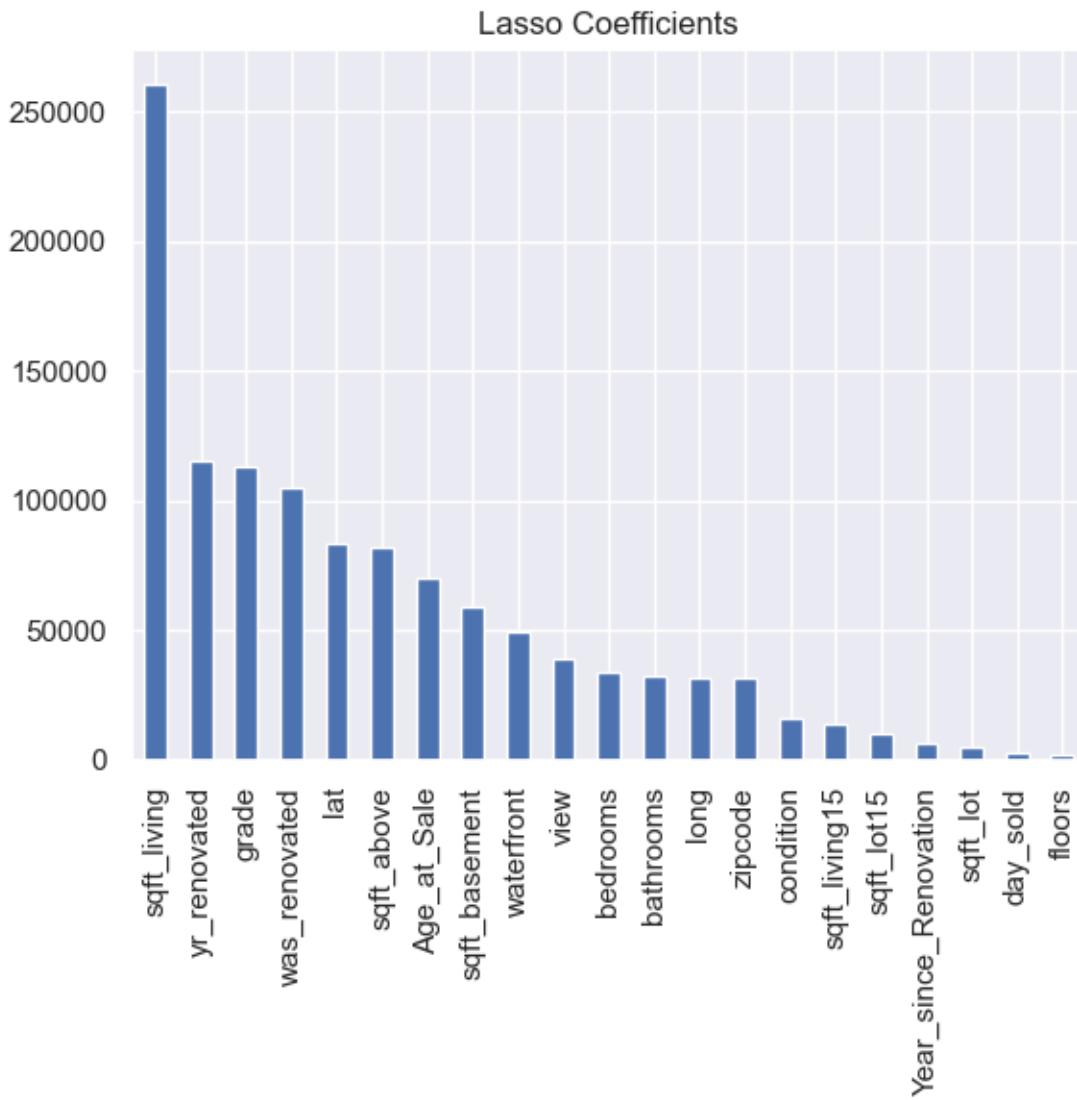
```
[11]: from IPython.display import clear_output  
  
lasso_regressor = Lasso(random_state=seed)  
  
lasso_regressor.fit(X_train_standard, y_train)  
  
# Remove warning  
clear_output()
```

```
[12]: # Regression coefficients  
  
coefs_lasso = pd.Series(np.abs(lasso_regressor.coef_), features).  
    ↪sort_values(ascending=False)  
  
coefs_lasso
```

```
[12]: sqft_living          260667.308014  
yr_renovated        115539.976929  
grade              113464.284916  
was_renovated       105055.743718  
lat                 83899.034194  
sqft_above          81996.071094  
Age_at_Sale         70493.906203  
sqft_basement        59474.251516  
waterfront          49347.059323  
view                39356.052812  
bedrooms            33745.986217  
bathrooms           32306.386441  
long                31980.674767  
zipcode             31606.104587  
condition            16509.503473  
sqft_living15        14330.914670  
sqft_lot15           10419.853206  
Year_since_Renovation  6825.039435  
sqft_lot              4926.447704  
day_sold             3153.359687  
floors               2301.221136
```

```
dtype: float64
```

```
[13]: sns.set_theme()  
coefs_lasso.plot(kind='bar', title='Lasso Coefficients')  
fig1 = plt.savefig("Graphs/Lasso_Coefficients.png")
```



## 9 Ridge (L2 Regularization)

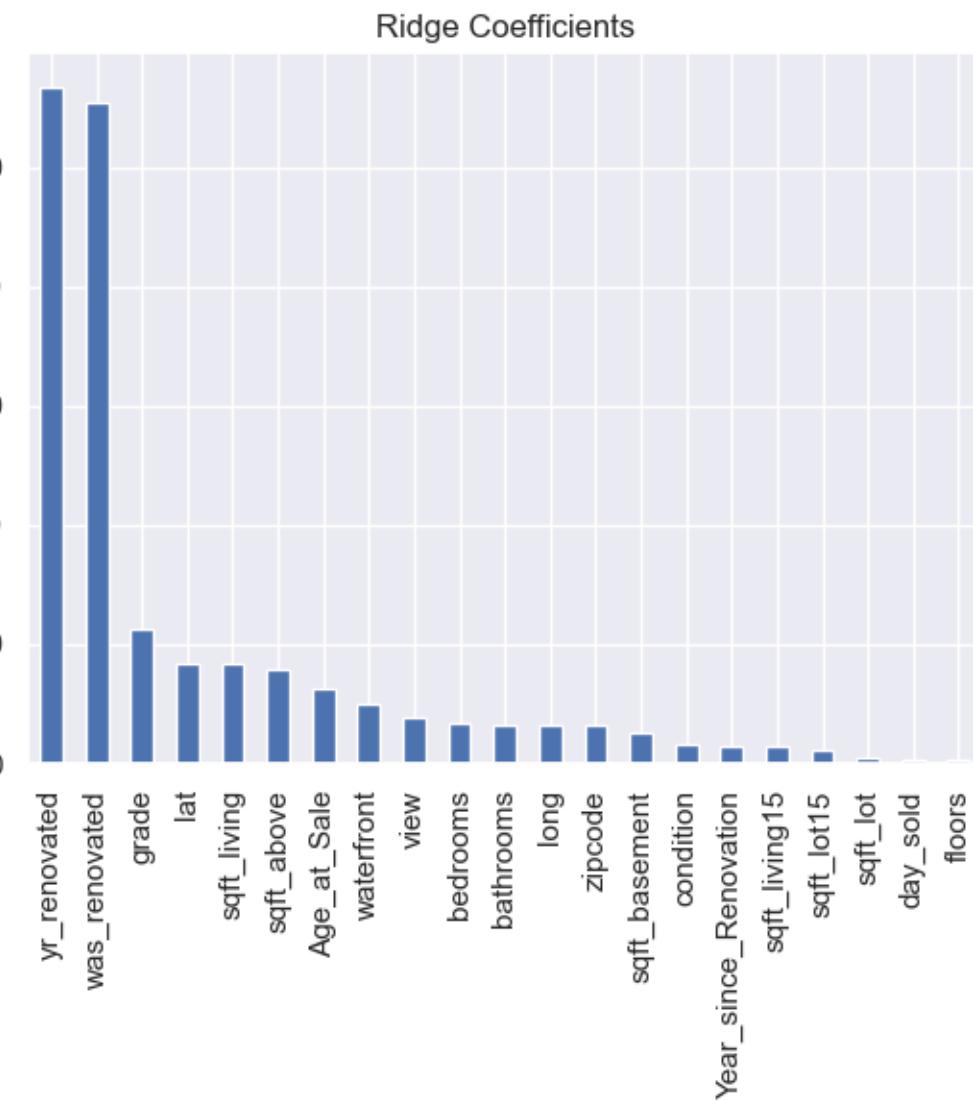
- Reduces coefficients magnitudes for correlated features. Should give us an idea of

```
[14]: ridge_regressor = Ridge(random_state=seed)  
ridge_regressor.fit(X_train_standard, y_train)
```

```
# Regression coefficients

coefs_ridge = pd.Series(np.abs(ridge_regressor.coef_), features).
    ↪sort_values(ascending=False)
```

```
[15]: sns.set_theme()
coefs_ridge.plot(kind='bar', title='Ridge Coefficients')
fig1 = plt.savefig("Graphs/Ridge_Coefficients.png")
```



## 10 Lasso vs Ridge

```
[16]: # combined

feature_importance = pd.DataFrame({"lasso": coefs_lasso, "ridge": coefs_ridge})
```

```
[17]: feature_importance = feature_importance.reset_index()
feature_importance
```

```
[17]:      index      lasso      ridge
0   Age_at_Sale  70493.906203  62758.474492
1 Year_since_Renovation  6825.039435  14832.195034
2   bathrooms  32306.386441  32258.641807
3   bedrooms  33745.986217  33822.267903
4   condition  16509.503473  16613.195141
5   day_sold  3153.359687  3174.078409
6     floors  2301.221136  2587.054029
7     grade  113464.284916  113284.414558
8       lat  83899.034194  83883.757531
9       long  31980.674767  31943.753755
10  sqft_above  81996.071094  78480.622929
11  sqft_basement  59474.251516  25424.127466
12  sqft_living  260667.308014  83030.526816
13  sqft_living15  14330.914670  14424.743626
14  sqft_lot  4926.447704  4972.626029
15  sqft_lot15  10419.853206  10457.284205
16     view  39356.052812  39373.650718
17  was_renovated  105055.743718  553980.792177
18    waterfront  49347.059323  49480.985886
19  yr_renovated  115539.976929  567539.527477
20     zipcode  31606.104587  31661.929745
```

```
[18]: import pandas as pd
import plotly.express as px

# Fix formating of the dataframe for plotting
df_melted = feature_importance.melt(id_vars='index',
                                      value_vars=['lasso', 'ridge'],
                                      var_name='Model',
                                      value_name='Coefficient')

# Create a interactive bar plot
fig = px.bar(df_melted.sort_values("Coefficient", ascending=False),
              x='index',
              y='Coefficient',
              color='Model',
              barmode='group',           # side-by-side, not stacked
```

```

color_discrete_map={
    'lasso': '#4c72b0',
    'ridge': '#dd8452'},
title='Feature Importance: Lasso vs Ridge')

# 3. Rotate x-axis labels
fig.update_layout(
    xaxis_tickangle=-45,
    xaxis_title='Features',
    yaxis_title='Absolute Coefficient Value'
)

fig.show()
fig.write_image("Graphs/Lasso_Ridge_Coefficients.png")

```

## 11 Linear Regression

- Given Numeriacal Dataset we perform a Linear Regression and start creating metrics

```
[19]: from sklearn.linear_model import LinearRegression

# Create Estimator
LinReg = LinearRegression()

# Perform Fitting
LinReg.fit(X_train, y_train)

# Make Predictions
y_train_predict = LinReg.predict(X_train)
y_test_predict = LinReg.predict(X_test)

print("Model trained successfully!")
print(f"Training predictions shape: {y_train_predict.shape}")
print(f"Testing predictions shape: {y_test_predict.shape}")
print(f"\nFirst 5 training predictions: {y_train_predict[:5]}")
print(f"First 5 actual training values: {y_train.iloc[:5].values}")

print(f"Intercept: {LinReg.intercept_:.2f}")
print(f"Number of coefficients: {len(LinReg.coef_)}")
print(f"\nTop3 Influential features:")
feature_importance = sorted(zip(X_train.columns,LinReg.coef_),key=lambda x:
    -abs(x[1]), reverse=True)[:3]
for feat, coef in feature_importance:
    print(f"{feat}: {coef:.4f}")

metrics_df = create_metrics_df()
```

```

metrics_df = add_new_metrics(metrics_df, LinReg, X_train, y_train, split="train", comments="Linear Regression FE Metrics - train - dropped features: date, id, yr_built, year_sold, month_sold")
metrics_df = add_new_metrics(metrics_df, LinReg, X_test, y_test, split="test", comments="Linear Regression FE Metrics - test - dropped features: date, id, yr_built, year_sold, month_sold")

metrics_df

```

Model trained successfully!

Training predictions shape: (17290,)

Testing predictions shape: (4323,)

First 5 training predictions: [247834.7561139 728804.13253798 810711.97302496  
216848.11429534  
90564.97897474]

First 5 actual training values: [284950. 625000. 838400. 282000. 218000.]

Intercept: 1355520.65

Number of coefficients: 21

Top3 Influential features:

was\_renovated: -9156171.3570

lat: 604623.1733

waterfront: 583972.7739

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6973	125914.8975	0.2557	
1	LinearRegression	test	0.7165	0.7151	125994.7484	0.2590	

	RMSE	Comments
0	202891.3616	Linear Regression FE Metrics - train - dropped...
1	191439.9182	Linear Regression FE Metrics - test - dropped ...

## 12 Mean Sq Error, Mean Absolute Error, Root Mean Squared Error

```

[20]: get_mse(y_train, y_test, y_train_predict, y_test_predict)

get_r_squared(y_train, y_test, y_train_predict, y_test_predict)

get_mae(y_train, y_test, y_train_predict, y_test_predict)

```

MSE score:

train | 41164904608.63687

test | 36649242289.43816

```
R2 score:  
train | 0.697664220549328  
test | 0.716453297645963
```

```
MSE score:  
train | 125914.89749025984  
test | 125994.74837969066
```

[20]: (125914.89749025984, 125994.74837969066)

## 13 Random Forest - Regression

- Random Forrest Regressor is applied, again without controlling any parameters.

```
[21]: import pandas as pd  
import numpy as np  
import statsmodels.api as sm  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
rf_regressor = RandomForestRegressor(n_estimators=110, random_state=seed)  
  
rf_regressor.fit(X_train, y_train)
```

[21]: RandomForestRegressor(n\_estimators=110, random\_state=13)

### 13.0.1 Predictions

```
[22]: y_pred = rf_regressor.predict(X_test)  
  
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)  
  
print(f"Mean Squared Error: {mse:.2f}")  
print(f"R-squared Score: {r2:.2f}")  
  
single_data = X_test.iloc[0].values.reshape(1, -1)  
predicted_value = rf_regressor.predict(single_data)  
print(f"Predicted Value: {predicted_value[0]:.2f}")  
print(f"Actual Value: {y_test.iloc[0]:.2f}")
```

```

metrics_df = add_new_metrics(metrics_df, rf_regressor, X_train, y_train,□
    ↪split="train", comments="Random Forest Train FE Metrics - train - dropped□
    ↪features: date, id, yr_built, year_sold, month_sold")
metrics_df = add_new_metrics(metrics_df, rf_regressor, X_test, y_test,□
    ↪split="test", comments="Random Forest Test FE Metrics - test - dropped□
    ↪features: date, id, yr_built, year_sold, month_sold")

metrics_df

```

Mean Squared Error: 13500260163.85

R-squared Score: 0.90

Predicted Value: 290463.59

Actual Value: 324900.00

c:\Users\KaoticCharma\anaconda3\Lib\site-packages\sklearn\utils\validation.py:2691: UserWarning:

X does not have valid feature names, but RandomForestRegressor was fitted with feature names

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6973	125914.8975	0.2557	
1	LinearRegression	test	0.7165	0.7151	125994.7484	0.2590	
2	RandomForestRegressor	train	0.9820	0.9820	25980.2304	0.0486	
3	RandomForestRegressor	test	0.8956	0.8950	67606.7385	0.1280	
		RMSE			Comments		
0	202891.3616	Linear Regression FE Metrics - train - dropped...					
1	191439.9182	Linear Regression FE Metrics - test - dropped ...					
2	49538.9365	Random Forest Train FE Metrics - train - dropp...					
3	116190.6199	Random Forest Test FE Metrics - test - dropped...					

## 14 Random Forest - Regression with Feature Importance

- Random Forrest Regressor is applied, again without controlling any parameters.

```

[23]: from sklearn.ensemble import RandomForestRegressor

rf_regressor = RandomForestRegressor(random_state=seed)#default values +
    ↪random_state = 13
rf_regressor.fit(X_train, y_train)

```

```
[23]: RandomForestRegressor(random_state=13)
```

```

[24]: y_pred = rf_regressor.predict(X_test)

mse = mean_squared_error(y_test, y_pred)

```

```

r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared Score: {r2:.2f}")

single_data = X_test.iloc[0].values.reshape(1, -1)
predicted_value = rf_regressor.predict(single_data)
print(f"Predicted Value: {predicted_value[0]:.2f}")
print(f"Actual Value: {y_test.iloc[0]:.2f}")

```

Mean Squared Error: 13619426178.18

R-squared Score: 0.89

Predicted Value: 289293.45

Actual Value: 324900.00

c:\Users\KaoticCharma\anaconda3\Lib\site-packages\sklearn\utils\validation.py:2691: UserWarning:

X does not have valid feature names, but RandomForestRegressor was fitted with feature names

```

[25]: features = X_train.columns
coefs_rf = pd.Series(np.abs(rf_regressor.feature_importances_), features).
    ↪sort_values(ascending=False)

metrics_df = add_new_metrics(metrics_df, rf_regressor, X_train, y_train, ↪
    ↪split="train", comments="Random Forest Train FE + FI Metrics - train -"
    ↪dropped features: date, id, yr_built, year_sold, month_sold")
metrics_df = add_new_metrics(metrics_df, rf_regressor, X_test, y_test, ↪
    ↪split="test", comments="Random Forest Test FE + FI Metrics - test - dropped"
    ↪features: date, id, yr_built, year_sold, month_sold")

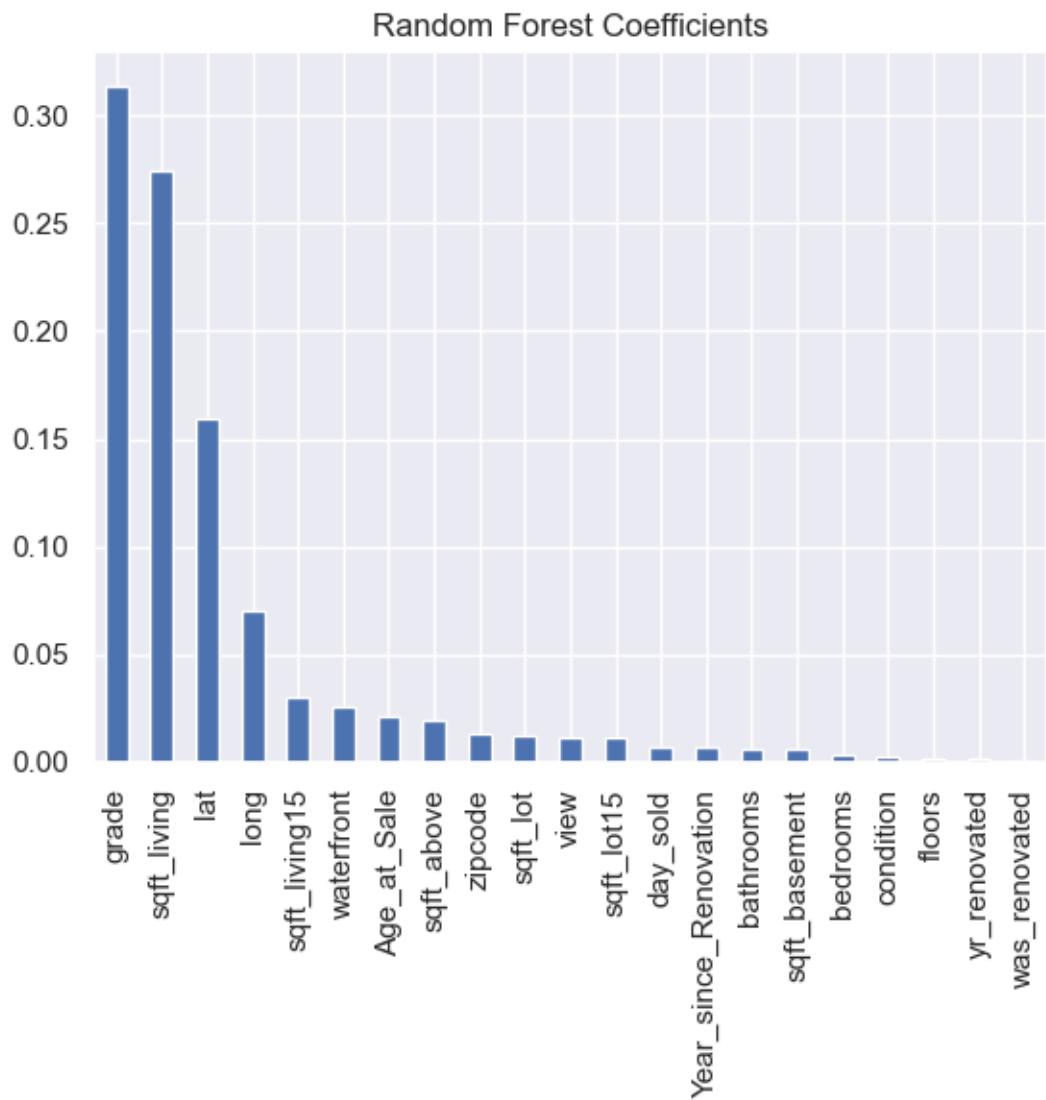
coefs_rf

```

grade	0.313535
sqft_living	0.274619
lat	0.159416
long	0.070064
sqft_living15	0.030581
waterfront	0.025367
Age_at_Sale	0.020887
sqft_above	0.019248
zipcode	0.012981
sqft_lot	0.012404
view	0.011900
sqft_lot15	0.011859
day_sold	0.007335

```
Year_since_Renovation      0.007207
bathrooms                  0.006546
sqft_basement               0.006109
bedrooms                   0.003223
condition                  0.002707
floors                     0.001969
yr_renovated                0.001513
was_renovated               0.000530
dtype: float64
```

```
[26]: coefs_rf.plot(kind='bar', title='Random Forest Coefficients')
fig1 = plt.savefig("Graphs/RandomForest_Coefficients.png")
```



## 15 XGBoost

```
[27]: xgb_regressor = xgb.XGBRegressor(n_estimators=100, random_state=seed)

xgb_regressor.fit(X_train, y_train)

metrics_df = add_new_metrics(metrics_df, xgb_regressor, X_train, y_train, □
    ↵split="train", comments="XGBoost Train FE Metrics - train - dropped features" □
    ↵- date, id, yr_built, year_sold, month_sold")

metrics_df = add_new_metrics(metrics_df, xgb_regressor, X_test, y_test, □
    ↵split="test", comments="XGBoost Test FE Metrics - test - dropped features -" □
    ↵date, id, yr_built, year_sold, month_sold")

metrics_df

metrics_df.to_csv("ModelMetrics_FE.csv", index=False)

metrics_df
```

```
[27]:
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6973	125914.8975	0.2557	
1	LinearRegression	test	0.7165	0.7151	125994.7484	0.2590	
2	RandomForestRegressor	train	0.9820	0.9820	25980.2304	0.0486	
3	RandomForestRegressor	test	0.8956	0.8950	67606.7385	0.1280	
4	RandomForestRegressor	train	0.9819	0.9819	26031.4643	0.0487	
5	RandomForestRegressor	test	0.8946	0.8941	67781.7366	0.1283	
6	XGBRegressor	train	0.9761	0.9761	40791.8177	0.0903	
7	XGBRegressor	test	0.9038	0.9033	66773.7091	0.1271	

	RMSE	Comments
0	202891.3616	Linear Regression FE Metrics - train - dropped...
1	191439.9182	Linear Regression FE Metrics - test - dropped ...
2	49538.9365	Random Forest Train FE Metrics - train - dropp...
3	116190.6199	Random Forest Test FE Metrics - test - dropped...
4	49598.7780	Random Forest Train FE + FI Metrics - train -...
5	116702.2972	Random Forest Test FE + FI Metrics - test - dr...
6	57067.4036	XGBoost Train FE Metrics - train - dropped fea...
7	111527.5852	XGBoost Test FE Metrics - test - dropped featu...

## 16 XGBoost with Feature Importance

```
[28]: import xgboost as xgb

xgb_clf = xgb.XGBRegressor(seed = seed)
xgb_clf.fit(X_train, y_train)
```

```
[28]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, device=None, early_stopping_rounds=None,
                   enable_categorical=False, eval_metric=None, feature_types=None,
                   feature_weights=None, gamma=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
                   max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                   max_leaves=None, min_child_weight=None, missing=nan,
                   monotone_constraints=None, multi_strategy=None, n_estimators=None,
                   n_jobs=None, num_parallel_tree=None, ...)
```

```
[29]: features = X_train.columns
coefs_xgb = pd.Series(np.abs(xgb_clf.feature_importances_), features).
    ↪sort_values(ascending=False)
metrics_df = add_new_metrics(metrics_df, xgb_clf, X_train, y_train, □
    ↪split="train", comments="XGBoost Train FE + FI Metrics - train - dropped"
    ↪features: date, id, yr_builtin, year_sold, month_sold")
metrics_df = add_new_metrics(metrics_df, xgb_clf, X_test, y_test, split="test", □
    ↪comments="XGBoost Test FE + FI Metrics - test - dropped features: date, id, □
    ↪yr_builtin, year_sold, month_sold")

coefs_xgb
```

```
[29]: grade                  0.382075
sqft_living                0.169310
waterfront                  0.144378
lat                         0.081549
long                        0.042398
view                        0.034126
Age_at_Sale                 0.024606
sqft_living15               0.022307
zipcode                      0.019849
yr_renovated                 0.013571
sqft_above                   0.012126
bathrooms                    0.010843
sqft_lot                     0.008284
condition                    0.007685
sqft_basement                 0.007208
Year_since_Renovation        0.005419
sqft_lot15                   0.005056
floors                       0.004667
day_sold                     0.003014
bedrooms                      0.001529
was_renovated                 0.000000
dtype: float32
```

```
[30]: coefs_xgb.plot(kind='bar', title='XGBoost Coefficients')
fig1 = plt.savefig("Graphs/XGBoost_Coefficients.png")
metrics_df
metrics_df.to_csv("ModelMetrics_FE.csv", index=False)
metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	LinearRegression	train	0.6977	0.6973	125914.8975	0.2557	
1	LinearRegression	test	0.7165	0.7151	125994.7484	0.2590	
2	RandomForestRegressor	train	0.9820	0.9820	25980.2304	0.0486	
3	RandomForestRegressor	test	0.8956	0.8950	67606.7385	0.1280	
4	RandomForestRegressor	train	0.9819	0.9819	26031.4643	0.0487	
5	RandomForestRegressor	test	0.8946	0.8941	67781.7366	0.1283	
6	XGBRegressor	train	0.9761	0.9761	40791.8177	0.0903	
7	XGBRegressor	test	0.9038	0.9033	66773.7091	0.1271	
8	XGBRegressor	train	0.9761	0.9761	40791.8177	0.0903	
9	XGBRegressor	test	0.9038	0.9033	66773.7091	0.1271	

	RMSE	Comments
0	202891.3616	Linear Regression FE Metrics - train - dropped...
1	191439.9182	Linear Regression FE Metrics - test - dropped ...
2	49538.9365	Random Forest Train FE Metrics - train - dropp...
3	116190.6199	Random Forest Test FE Metrics - test - dropped...
4	49598.7780	Random Forest Train FE + FI Metrics - train -...
5	116702.2972	Random Forest Test FE + FI Metrics - test - dr...
6	57067.4036	XGBoost Train FE Metrics - train - dropped fea...
7	111527.5852	XGBoost Test FE Metrics - test - dropped featu...
8	57067.4036	XGBoost Train FE + FI Metrics - train - droppe...
9	111527.5852	XGBoost Test FE + FI Metrics - test - dropped ...

