

6_Model_Optimization

February 13, 2026

1 Model Hyperparameters Optimization

From the tested models, Random Forest and XGBoost seem to be the best ones, consistently getting high metrics on the test split. So we are going to run a gridsearch to find the best parameters for these two models on the v1 of te data.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

# Models & Normalization
from sklearn.model_selection import train_test_split, RandomizedSearchCV

# Evaluation
import statsmodels.api as sm

# Extra
from utils import *
```

```
[2]: # filename = "Data/v1_house_sales.csv"
# filename = "Data/v2_house_sales.csv"
filename = "Data/v3_house_sales.csv"

df = pd.read_csv(filename)
```

```
[3]: df.head()
```

```
[3]:   Unnamed: 0      price  bathrooms  sqft_living  waterfront  view  condition \
0          0  221900.0       1.00      1180          0      0      0      3
1          1  538000.0       2.25      2570          0      0      0      3
2          2  180000.0       1.00       770          0      0      0      3
3          3  604000.0       3.00      1960          0      0      0      5
4          4  510000.0       2.00      1680          0      0      0      3

      grade  sqft_above  sqft_basement  zipcode        lat        long  sqft_living15 \
0         7        1180                 0    98178  47.5112 -122.257           1340
```

1	7	2170	400	98125	47.7210	-122.319	1690
2	6	770	0	98028	47.7379	-122.233	2720
3	7	1050	910	98136	47.5208	-122.393	1360
4	8	1680	0	98074	47.6168	-122.045	1800
0	5650	1	1	59	59	False	
1	7639	1	1	63	23	True	
2	8062	1	1	82	82	False	
3	5000	1	1	49	49	False	
4	7503	1	1	28	28	False	

```
[4]: # Split into train and test
seed = 13
# The price is the target variable
y = df["price"]

# All other variables are the features for the baseline model
X = df.drop(["price"], axis=1)

# Train Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=seed)
```

```
[5]: X_train.head()
```

1571	1571	1.5	2000	0	0	4	7
16330	16330	2.5	2630	0	0	3	8
12786	12786	2.5	2620	0	0	4	9
12524	12524	2.5	1610	0	0	4	7
16179	16179	1.0	880	0	0	4	6
1571	1170	830	98198	47.3708	-122.311	1940	
16330	2630	0	98072	47.7750	-122.125	2680	
12786	2620	0	98040	47.5631	-122.219	2580	
12524	1610	0	98038	47.3490	-122.036	1570	
16179	880	0	98178	47.5013	-122.244	1110	
1571	7531	1	1	53	53		
16330	48706	1	1	28	28		
12786	9525	1	1	40	40		
12524	6000	1	1	21	21		
16179	16115	1	1	69	69		

```

was_renovated
1571      False
16330     False
12786     False
12524     False
16179     False

```

1.1 Metrics dataframe

```
[6]: from utils import *

metrics_df = create_metrics_df()
```

1.2 Random Forest

1.2.1 Grid Search

```
[7]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV

# Define the model with a fixed, high number of trees
# n_estimators=300 is a good balance of speed vs stability
rf_regressor = RandomForestRegressor(n_estimators=300, random_state=seed)

# Improved Parameter Grid
param_dist = {
    # Tree Depth: Control complexity
    'max_depth': [None, 10, 20, 30, 40, 50],

    # Split Criteria: Higher values prevent overfitting
    'min_samples_split': [2, 5, 10, 15, 20],

    # Leaf Size: Critical for regression smoothness
    'min_samples_leaf': [1, 2, 4, 8, 12],

    # Feature Selection: 'sqrt' is standard, but try fractions (0.3, 0.5)
    # Using floats (0.3) means "use 30% of features"
    'max_features': ['sqrt', 'log2', 0.3, 0.5, None],
}

# Bootstrapping: Usually True is best, but False can work for small data
'bootstrap': [True, False]
}

# Randomized Search
random_search = RandomizedSearchCV(
    estimator=rf_regressor,
    param_distributions=param_dist,
    n_iter=50,                      # 50 iterations is plenty for random search
```

```

        cv=3,
        scoring='neg_root_mean_squared_error', # Use RMSE (easier to interpret)
        n_jobs=-1,
        verbose=1,
        random_state=seed
    )

# Fit
random_search.fit(X_train, y_train)

# Results
print(f"Best parameters: {random_search.best_params_}")
best_rf = random_search.best_estimator_

```

Fitting 3 folds for each of 50 candidates, totalling 150 fits

```

/Users/anne/Documents/_IronHack/Projects/King-County-Housing-
Analysis/.venv/lib/python3.13/site-
packages/joblib/externals/loky/process_executor.py:782: UserWarning: A worker
stopped while some jobs were given to the executor. This can be caused by a too
short worker timeout or by a memory leak.

```

```
    warnings.warn(
```

```
Best parameters: {'min_samples_split': 5, 'min_samples_leaf': 1, 'max_features':
'log2', 'max_depth': None, 'bootstrap': True}
```

[8]: # Evaluate on train and test

```

best_params_str = f"Best parameters: {random_search.best_params_}"

metrics_df = add_new_metrics(metrics_df,
                             best_rf,
                             X_train,
                             y_train,
                             split = "train",
                             comments=best_params_str)

metrics_df = add_new_metrics(metrics_df,
                             best_rf,
                             X_test,
                             y_test,
                             split = "test",
                             comments=best_params_str)

```

[9]: metrics_df

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9792	0.9792	29830.8495	0.061	
1	RandomForestRegressor	test	0.9318	0.9315	61500.1162	0.126	

	RMSE	Comments
0	53176.4381	Best parameters: {'min_samples_split': 5, 'min...
1	93872.8212	Best parameters: {'min_samples_split': 5, 'min...

1.3 XGBoost

1.3.1 Grid Search

```
[10]: import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV

# Define the model
xgb_clf = xgb.XGBRegressor(seed=seed, objective='reg:squarederror')

param_dist = {
    'n_estimators': [100, 200, 300, 500, 1000], # More trees is usually okay
    ↪with early stopping
    'max_depth': [3, 5, 7, 10], # XGBoost prefers shallower
    ↪trees than RF
    'learning_rate': [0.01, 0.05, 0.1, 0.3], # Critical for XGBoost

    # "min_samples_leaf" equivalent:
    'min_child_weight': [1, 3, 5, 10],

    # "max_features" equivalent:
    'colsample_bytree': [0.5, 0.7, 1.0],

    # "min_samples_split" equivalent (approximate):
    'gamma': [0, 0.1, 0.5, 1],

    # Regularization (optional but good)
    'subsample': [0.6, 0.8, 1.0]
}

# Randomized search
random_search = RandomizedSearchCV(
    estimator=xgb_clf,
    param_distributions=param_dist,
    n_iter=50, # 50 is usually enough for random search
    cv=3,
    scoring='neg_root_mean_squared_error', # Better metric for regression
    n_jobs=-1,
    verbose=1,
    random_state=seed
)

# Fit
```

```

random_search.fit(X_train, y_train)

# Results
print(f"Best parameters: {random_search.best_params_}")
best_xgb = random_search.best_estimator_

```

Fitting 3 folds for each of 50 candidates, totalling 150 fits
Best parameters: {'subsample': 0.6, 'n_estimators': 1000, 'min_child_weight': 3, 'max_depth': 5, 'learning_rate': 0.05, 'gamma': 0, 'colsample_bytree': 0.5}

```
[11]: # Evaluate on train and test
best_params_str = f"Best parameters: {random_search.best_params_}"

metrics_df = add_new_metrics(metrics_df,
                             best_xgb,
                             X_train,
                             y_train,
                             split = "train",
                             comments=best_params_str)

metrics_df = add_new_metrics(metrics_df,
                             best_xgb,
                             X_test,
                             y_test,
                             split = "test",
                             comments=best_params_str)
```

```
[12]: metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9792	0.9792	29830.8495	0.0610	
1	RandomForestRegressor	test	0.9318	0.9315	61500.1162	0.1260	
2	XGBRegressor	train	0.9753	0.9752	41858.7250	0.0938	
3	XGBRegressor	test	0.9358	0.9355	58910.2638	0.1190	

	RMSE	Comments
0	53176.4381	Best parameters: {'min_samples_split': 5, 'min...
1	93872.8212	Best parameters: {'min_samples_split': 5, 'min...
2	58040.5476	Best parameters: {'subsample': 0.6, 'n_estimat...
3	91103.1330	Best parameters: {'subsample': 0.6, 'n_estimat...

1.4 Gradient Boosting

```
[13]: from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import RandomizedSearchCV

# Define the base model
gb_regressor = GradientBoostingRegressor(random_state=seed)
```

```

# Parameter Grid optimized for Gradient Boosting
param_dist = {
    # 1. Boosting Parameters (Critical Pair: Rate vs Trees)
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'n_estimators': [100, 200, 300, 500],

    # 2. Tree Structure
    # GB trees are usually shallow (depth 3-5 works best)
    'max_depth': [3, 4, 5, 6, 8],

    # 3. Regularization (Prevents overfitting)
    'min_samples_split': [2, 5, 10, 20],
    'min_samples_leaf': [1, 2, 4, 10],

    # 4. Stochastic Boosting (Using < 1.0 helps reduce variance)
    'subsample': [0.7, 0.8, 0.9, 1.0],

    # 5. Feature Randomness (Like RF, helps if features are correlated)
    'max_features': ['sqrt', 'log2', 0.5, None]
}

# Randomized Search
random_search = RandomizedSearchCV(
    estimator=gb_regressor,
    param_distributions=param_dist,
    n_iter=50,                      # 50 iterations is plenty
    cv=3,
    scoring='neg_root_mean_squared_error',
    n_jobs=-1,
    verbose=1,
    random_state=seed
)

# Fit
random_search.fit(X_train, y_train)

# Results
print(f"Best parameters: {random_search.best_params_}")
best_gb = random_search.best_estimator_

```

Fitting 3 folds for each of 50 candidates, totalling 150 fits
 Best parameters: {'subsample': 0.7, 'n_estimators': 200, 'min_samples_split': 20, 'min_samples_leaf': 10, 'max_features': 'sqrt', 'max_depth': 6, 'learning_rate': 0.1}

```
[14]: # Evaluate on train and test
best_params_str = f"Best parameters: {random_search.best_params_}"

metrics_df = add_new_metrics(metrics_df,
                             best_gb,
                             X_train,
                             y_train,
                             split = "train",
                             comments=best_params_str)

metrics_df = add_new_metrics(metrics_df,
                             best_gb,
                             X_test,
                             y_test,
                             split = "test",
                             comments=best_params_str)
```

```
[15]: metrics_df
```

	Model	Split	R2	Adjusted_R2	MAE	MAPE	\
0	RandomForestRegressor	train	0.9792	0.9792	29830.8495	0.0610	
1	RandomForestRegressor	test	0.9318	0.9315	61500.1162	0.1260	
2	XGBRegressor	train	0.9753	0.9752	41858.7250	0.0938	
3	XGBRegressor	test	0.9358	0.9355	58910.2638	0.1190	
4	GradientBoostingRegressor	train	0.9582	0.9581	52551.4223	0.1126	
5	GradientBoostingRegressor	test	0.9347	0.9345	61386.0436	0.1239	

	RMSE	Comments
0	53176.4381	Best parameters: {'min_samples_split': 5, 'min...
1	93872.8212	Best parameters: {'min_samples_split': 5, 'min...
2	58040.5476	Best parameters: {'subsample': 0.6, 'n_estimat...
3	91103.1330	Best parameters: {'subsample': 0.6, 'n_estimat...
4	75480.4486	Best parameters: {'subsample': 0.7, 'n_estimat...
5	91838.4174	Best parameters: {'subsample': 0.7, 'n_estimat...

1.5 Saving Best Parameters

```
[16]: # filename = "Metrics/v1_best_model.csv"
# filename = "Metrics/v2_best_model.csv"
filename = "Metrics/v3_best_model.csv"

metrics_df.to_csv(filename, index=False)
```

```
[ ]:
```